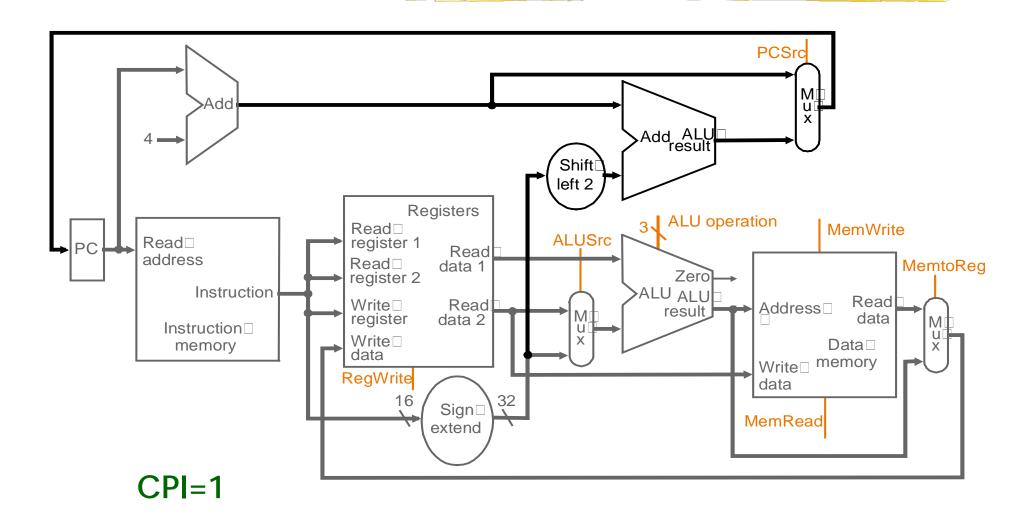
CS2006: 計算機組織

Designing a Multicycle Processor

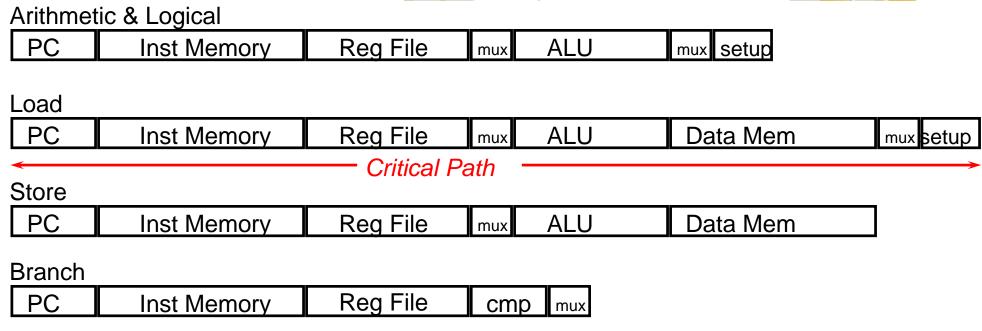
Outline

- Designing a processor
- Building the datapath
- A single-cycle implementation
- A multicycle implementation
- Microprogramming: simplifying control (Appendix C.4)
- Exceptions

Recap: A Single-Cycle Processor



What's Wrong with Single-cycle?



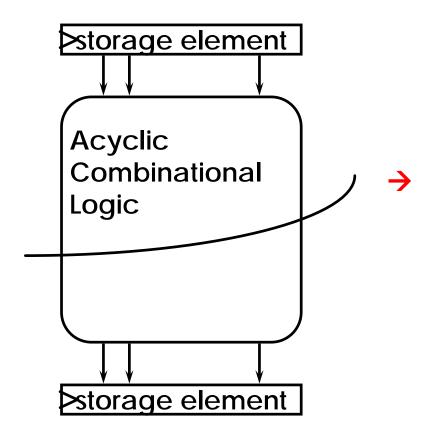
- Long cycle time
- All instructions take same time as the slowest
- Real memory is not so ideal
 - cannot always get job done in one (short) cycle
- ◆ A FU can only be used once => higher cost

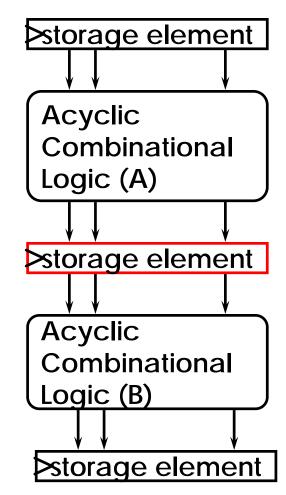
Outline

- Designing a processor
- Building the datapath
- A single-cycle implementation
- A multicycle implementation
 - Multicycle datapath
 - Multicycle execution steps
 - Multicycle control (Appendix C.3)
- Microprogramming: simplifying control (Appendix C.4)
- Exceptions

Multicycle Implementation

- Reduce cycle time
- Diff. Inst. take diff. cycles
- Share functional units

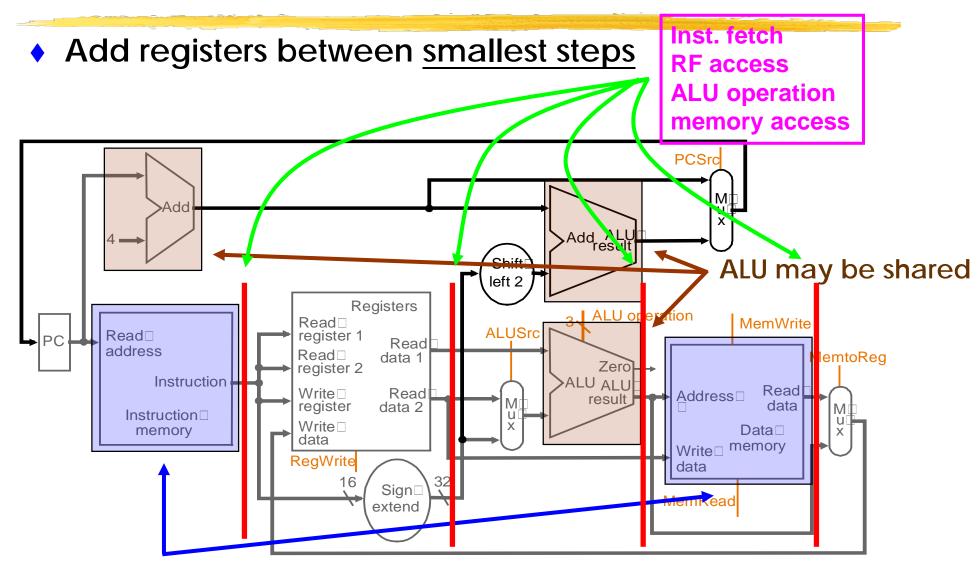




Multicycle Approach

- Break up the instructions into steps, each step takes a cycle
 - balance the amount of work to be done
 - restrict each cycle to use only one major functional unit
- At the end of a cycle
 - store values for use in later cycles (easiest thing to do)
 - introduce additional internal registers

Partition Single-Cycle Datapath



Multicycle Datapath

- 1 memory (inst. & data), 1 ALU (addr, PC+4, add,...), registers (IR, MDR, A, B, ALUOut)
 - Storage for subsequent inst. (arch.-visible) vs. storage for same inst. but in a subsequent cycle

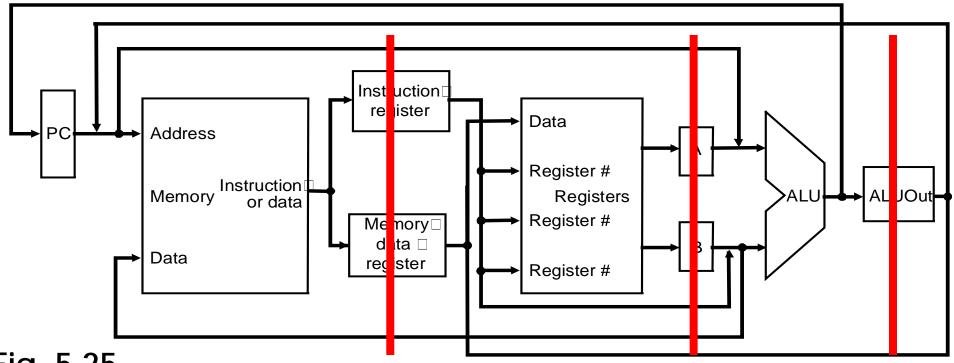
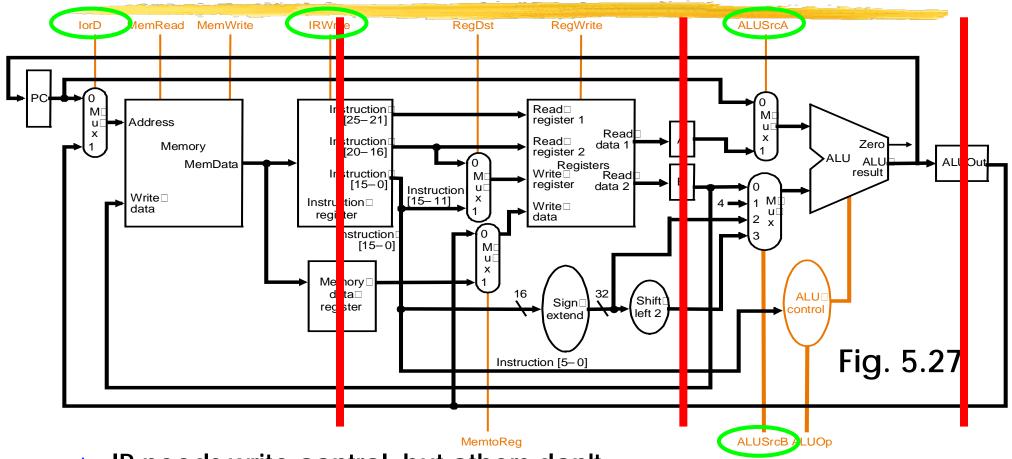


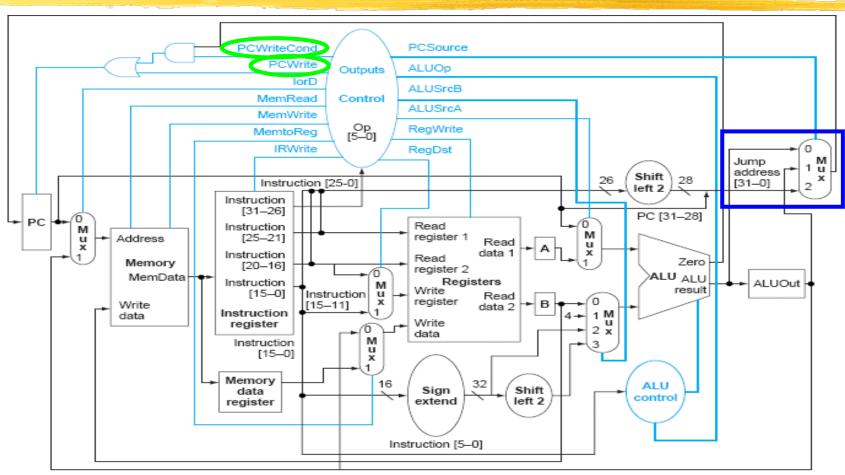
Fig. 5.25

Multicycle Datapath for Basic Inst.



- IR needs write control, but others don't
- MUX to select 2 sources to memory; memory needs read signal
- PC and A to one ALU input; four sources to another input

Adding Branch/Jump



- Three sources to PC
- Two PC write signals

Fig. 5.28

Outline

- Designing a processor
- Building the datapath
- A single-cycle implementation
- A multicycle implementation
 - Multicycle datapath
 - Multicycle execution steps
 - Multicycle control (Appendix C.3)
- Microprogramming: simplifying control (Appendix C.4)
- Exceptions

Five Execution Steps

- Instruction Fetch
- Instruction Decode and Register Fetch
- Execution, Memory Address Computation, or Branch Completion
- Memory Access or R-type Instruction Completion
- Memory Read Completion (Write-back)

INSTRUCTIONS TAKE FROM 3 - 5 CYCLES!

Step 1: Instruction Fetch

- Use PC to get instruction and put it in the Instruction Register (IR)
- Increment the PC by 4 and put the result back in the PC
- Can be described succinctly using RTL (Register-Transfer Language)

```
IR = Memory[PC];
PC = PC + 4;
```

Can we figure out the values of the control signals?

What is the advantage of updating the PC now?

Step 2: Instruction Decode and Register Fetch

- Read registers rs and rt (not harmful if not needed)
- Compute the branch address (not harmful if not needed)
- RTL:

```
A = Reg[IR[25-21]];
B = Reg[IR[20-16]];
ALUOut=PC+(sign-ext(IR[15-0])<<2);</pre>
```

We aren't setting any control lines based on the instruction type yet

(we are busy "decoding" it in control logic)

Step 3: Execution

ALU is performing one of three functions, based on instruction type:

Memory Reference:
ALUOut = A + sign-extend(IR[15-0]);

- R-type:
 ALUOut = A op B;
- Branch:
 if (A==B) PC = ALUOut;

Step 4: R-type or Memory-access

Loads and stores access memory

```
MDR = Memory[ALUOut];
    or
Memory[ALUOut] = B;
```

R-type instructions finish

```
Reg[IR[15-11]] = ALUOut;
```

The write actually takes place at the end of the cycle on the edge

Step 5: Write-back

Loads write to register

Reg[IR[
$$20-16$$
]] = MDR;

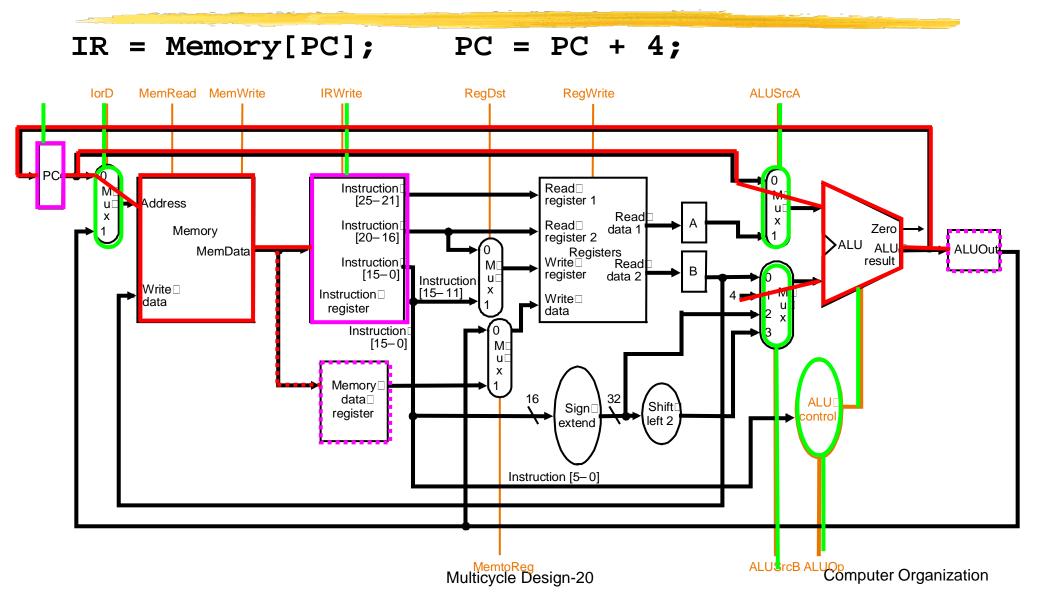
What about all the other instructions?

Summary of the Steps

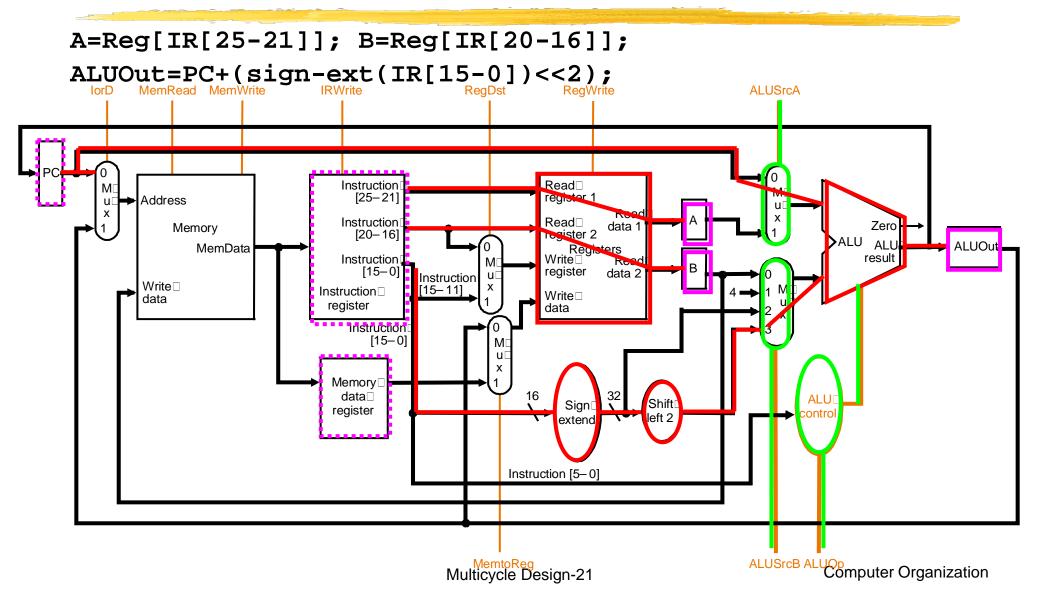
Step name	Action for R-type instructions	Action for memory-reference instructions	Action for branches	Action for jumps		
Instruction fetch	IR = Memory[PC] PC = PC + 4					
Instruction decode/register fetch	A = Reg [IR[25-21]] B = Reg [IR[20-16]] ALUOut = PC + (sign-extend (IR[15-0]) << 2)					
Execution, address computation, branch/jump completion	ALUOut = A op B	ALUOut = A + sign-extend (IR[15-0])	if (A ==B) then PC = ALUOut	PC = PC [31-28] II (IR[25-0]<<2)		
Memory access or R-type completion	Reg [IR[15-11]] = ALUOut	Load: MDR = Memory[ALUOut] or Store: Memory [ALUOut] = B				
Memory read completion		Load: Reg[IR[20-16]] = MDR				

Fig. 5.30

Cycle 1 of add

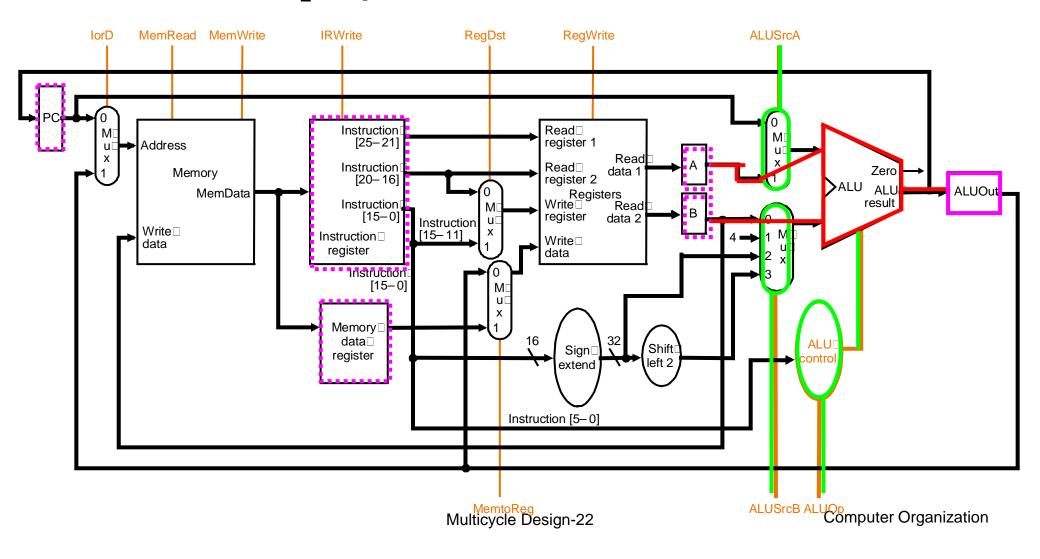


Cycle 2 of add



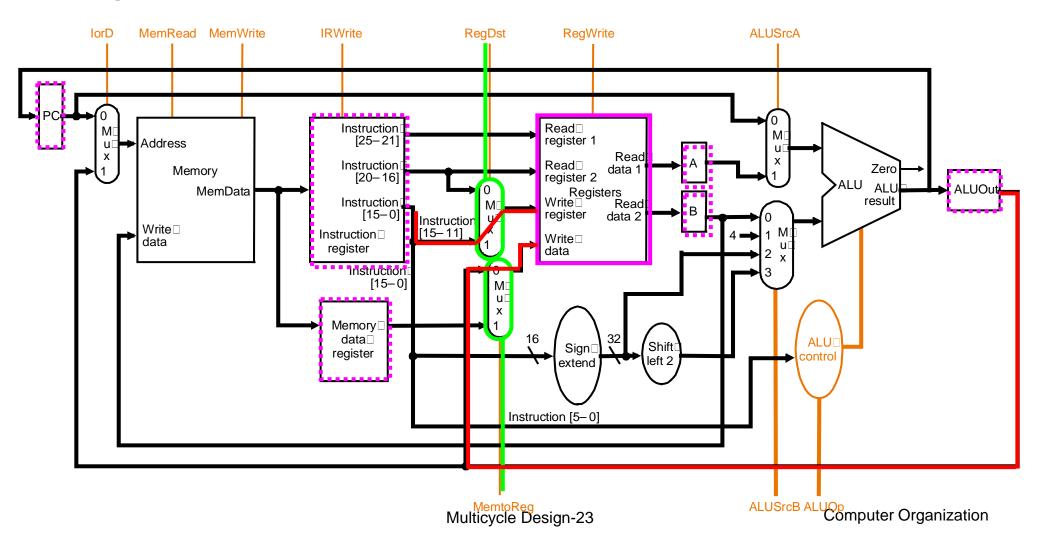
Cycle 3 of add

ALUOut = A op B;



Cycle 4 of add

Reg[IR[15-11]] = ALUOut;



Simple Questions

How many cycles will it take to execute this code?

```
lw $t2, 0($t3)
lw $t3, 4($t3)
beq $t2, $t3, Label
add $t5, $t2, $t3
sw $t5, 8($t3)
Label:
```

- What is going on during the 8th cycle of execution?
- In what cycle does the actual addition of \$t2 and \$t3 takes place?

Outline

- Designing a processor
- Building the datapath
- A single-cycle implementation
- A multicycle implementation
 - Multicycle datapath
 - Multicycle execution steps
 - Multicycle control (Appendix C.3)
- Microprogramming: simplifying control (Appendix C.4)
- Exceptions

Implementing the Control

- Value of control signals is dependent upon:
 - what instruction is being executed
 - which step is being performed
 - Control must specify both the signals to be set in any step and the next step in the sequence
- Control specification
 - Use a finite state machine (graphically)
 - Use microprogramming
- Implementation can be derived from the specification and use gates, ROM, or PLA

Controller Design: An Overview

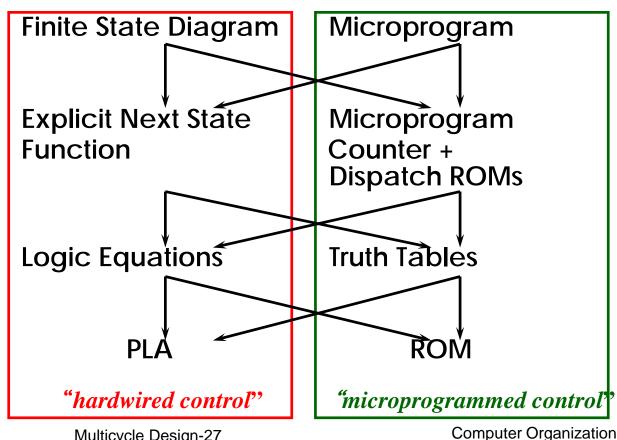
 Several possible initial representations, sequence control and logic representation, and control implementation => all may be determined indep.

Initial Rep.

Sequencing Control

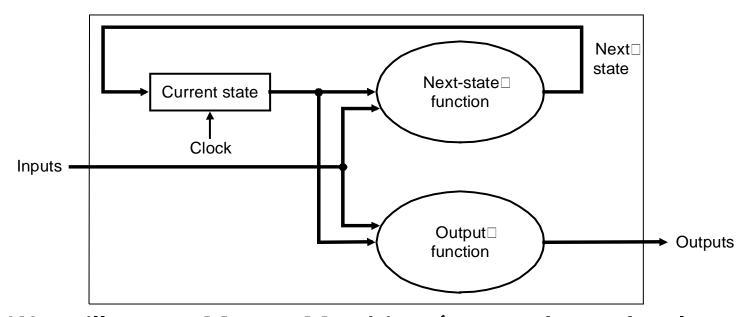
Logic Rep.

Implementation



Review: Finite State Machines

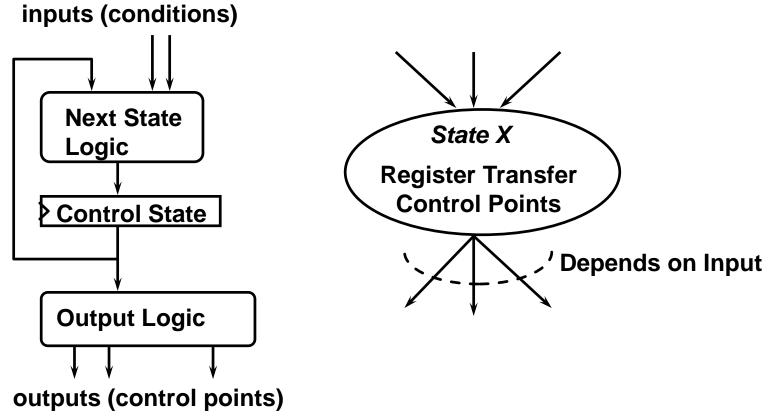
- Finite state machines:
 - a set of states and
 - next state (set by current state and input)
 - output (set by current state and possibly input)



 We will use a Moore Machine (output based only on the current state)

Our Control Model

- State specifies control points for RT
- Transfer at exiting state (same falling edge)
- One state takes one cycle

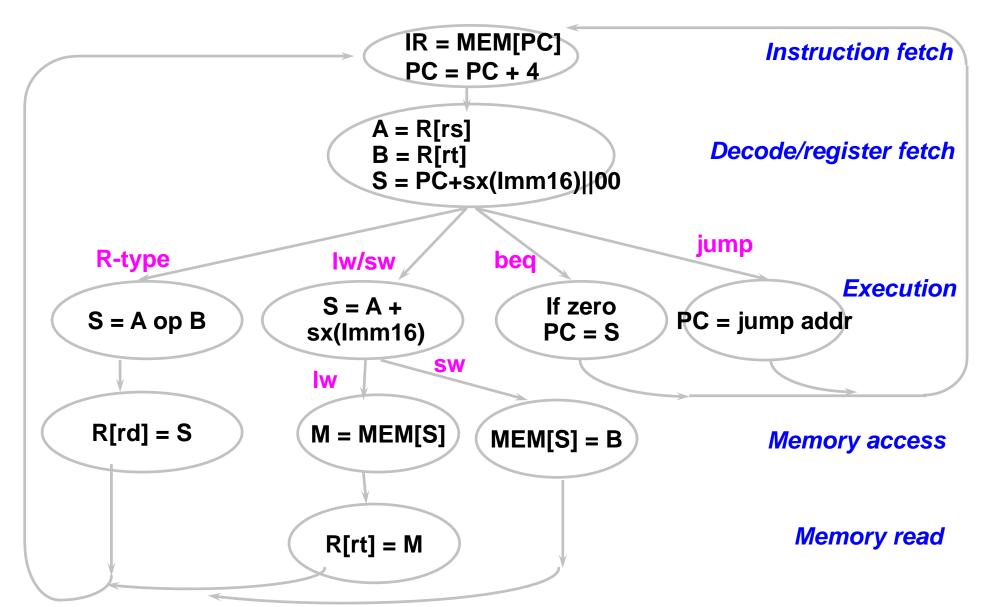


Summary of the Steps

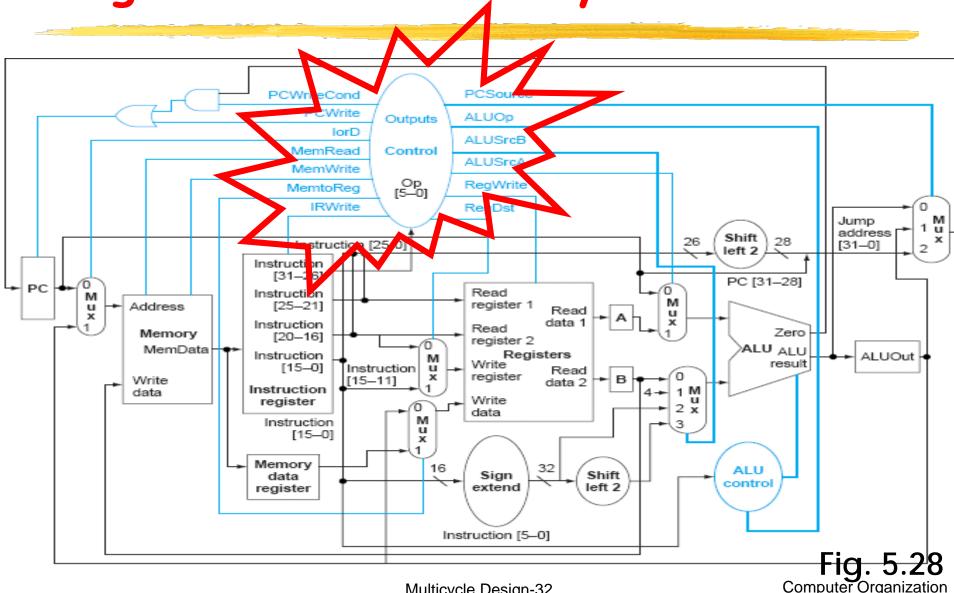
Step name	Action for R-type instructions	Action for memory-reference instructions	Action for branches	Action for jumps		
Instruction fetch	IR = Memory[PC] PC = PC + 4					
Instruction decode/register fetch	A = Reg [IR[25-21]] B = Reg [IR[20-16]] ALUOut = PC + (sign-extend (IR[15-0]) << 2)					
Execution, address computation, branch/jump completion	ALUOut = A op B	ALUOut = A + sign-extend (IR[15-0])	if (A ==B) then PC = ALUOut	PC = PC [31-28] II (IR[25-0]<<2)		
Memory access or R-type completion	Reg [IR[15-11]] = ALUOut	Load: MDR = Memory[ALUOut] or Store: Memory [ALUOut] = B				
Memory read completion		Load: Reg[IR[20-16]] = MDR				

Fig. 5.30

Control Specification for Multicycle



Organization of Multicycle Processor



Control Signals

Signal name ALUSrcA RegWrite MemtoReg RegDst MemRead MemWrite IorD IRWrite PCWrite PCWriteCond	1st ALU None Reg. da Reg. wr None None Memory None None	yhen deasserted J operand = PC Ita input = ALU Tite dest. no. = rt y address = PC	Ist ALU operand = Reg[rs] Reg file is written Reg. write data input = MDR Reg. write dest. no. = rd Memory at address is read Memory at address is written Memory address = ALUout IR = Memory PC = PCSource If zero then PC = PCSource	
	None		ii zero tileli PC = PC30urce	
Signal name	Value	<u>Effect</u>		
ALUOp	00	ALU adds	Fig. 5.29	9
AL LICHAD	01 10		ccording to func code	
ALUSrcB	00	2nd ALU input =	: B	
	01	2nd ALU input =	: 4	
	10	2nd ALU input =	sign extended IR[15-0]	
D00	11		sign ext., shift left 2 IR[15-0]	
PCSource	00	PC = ALU (PC +		
	01	PC = ALUOUT (b)	ranch target address)	
	10	PC = PC + 4[31 - 28]	8] : IR[25-0] << 2	

Mapping RT to Control Signals

High-level view of the finite state machine control

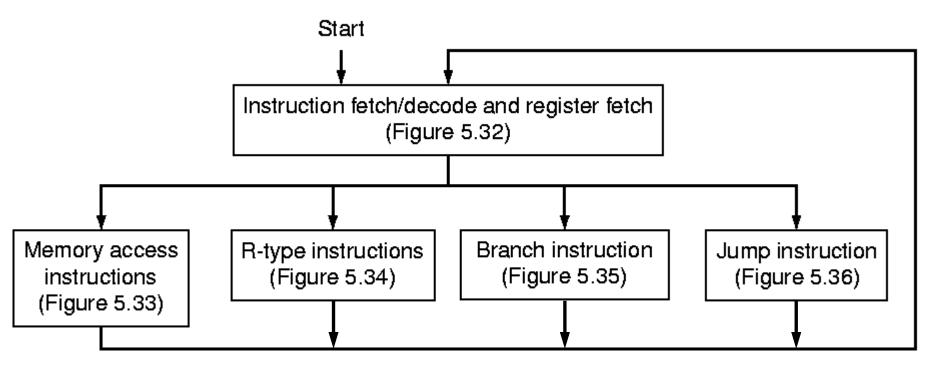
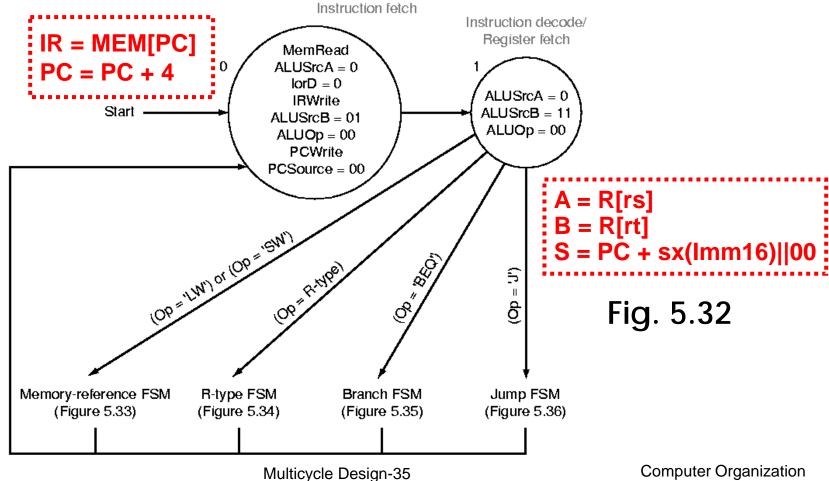


Fig. 5.31

Mapping RT to Control Signals

 Instruction fetch and decode portion of every instruction is identical:

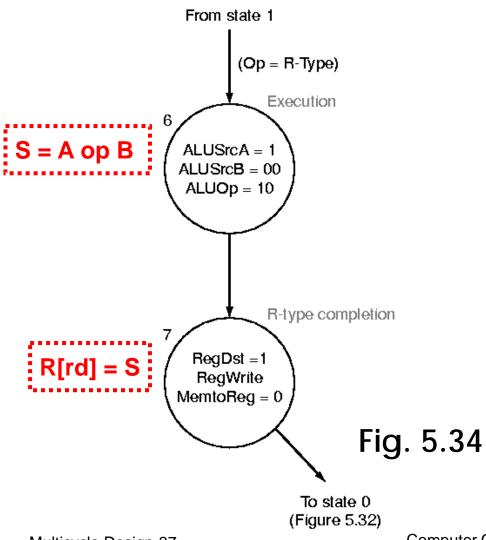


Mapping RT to Control Signals

 FSM for controlling From state 1 memory reference (Op = 'LW') or (Op = 'SW')instructions: Memory address computation ALUSrcA = 1S = A + sx(Imm16)ALUSrcB = 10 ALUOp = 00<u>`</u> Memory Memory access access 5 M = MEM[S] MEM[S] = BMemWrite MemRead IorD = 1IorD = 1Memory read completion step Fig. 5.33 RegWrite To state 0 MemtoReg =1 (Figure 5.32) RegDst = 0Computer Organization Multicycle Design-36

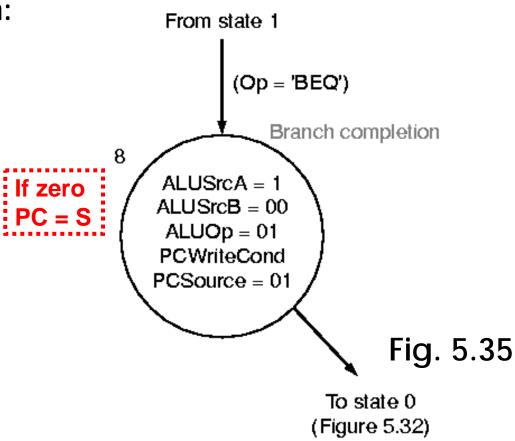
Mapping RT to Control Signals

 FSM for controlling R-type instructions:



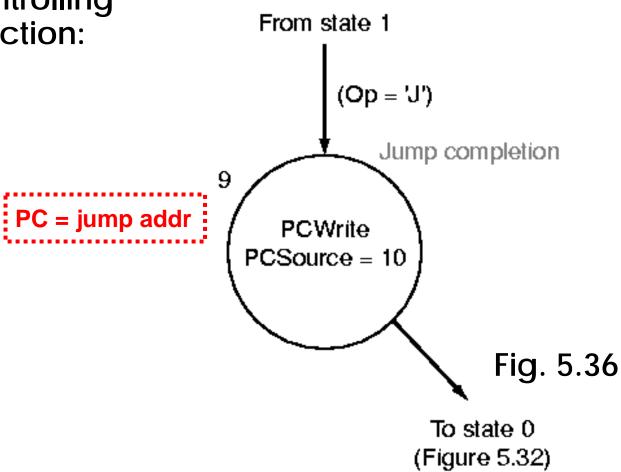
Mapping RT to Control Signals

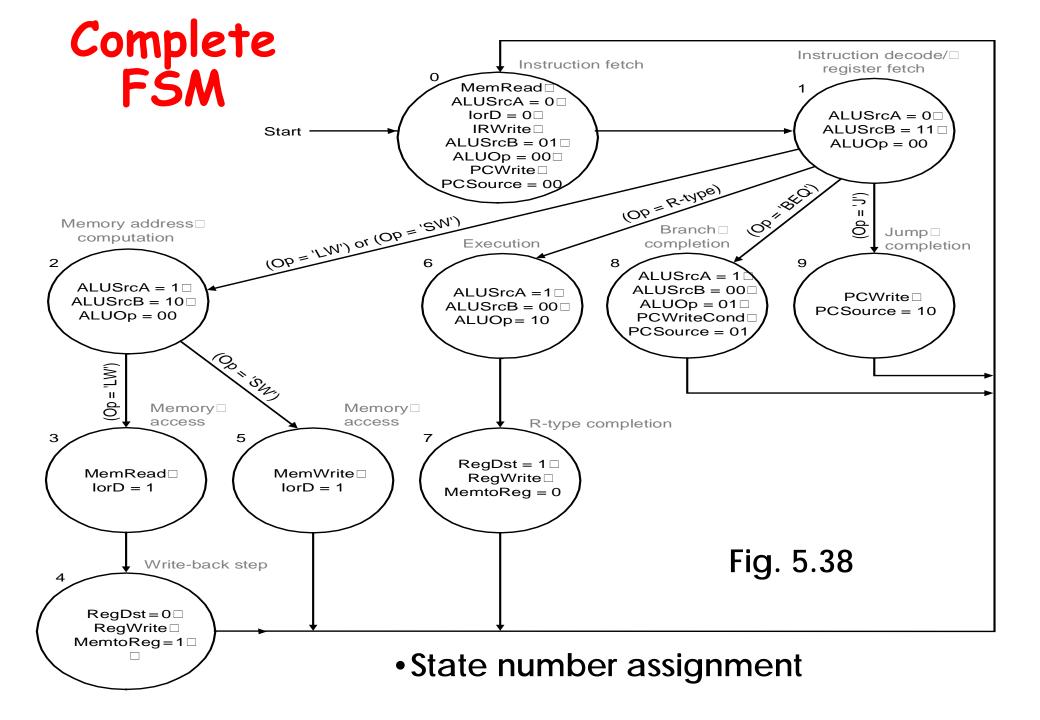
FSM for controlling branch instruction:



Mapping RT to Control Signals

 FSM for controlling jump instruction:





Truth Table

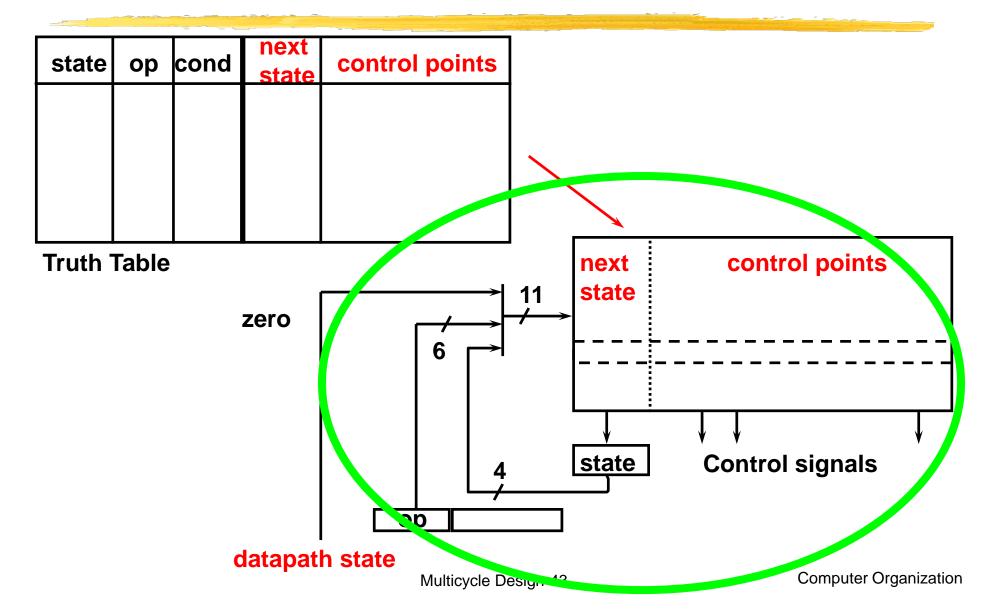
	input	output	
	op S	Datapath control	NS
·	(0000 00000	1001010000001000	0001
	000000 0001	000000000011000	0110
	000000 0010	000000000010100	XXXX
R-type≺	:		
	000000 1010		
	:		
	000000 1111		
	000001 0000		
	:		
	000010 0000	1001010000001000	0001
lumn	000010 0001	0000000000011000	1001
Jump ≺	000010 0010	0000000000010100	XXXX

From FSM to Truth Table

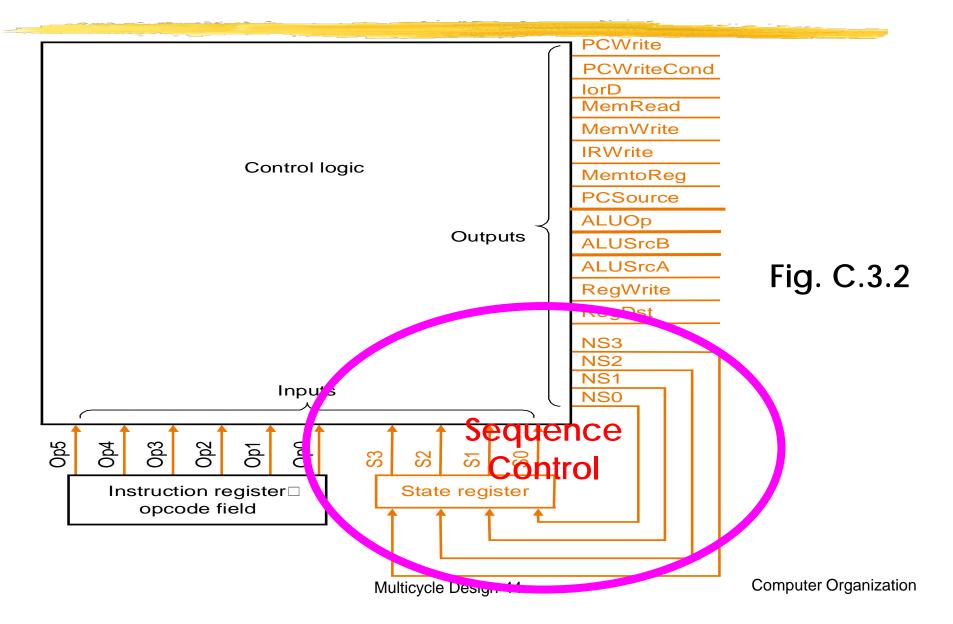
 Please reference the logic equations in Fig. C.3.3 and the truth table in Fig. C.3.6

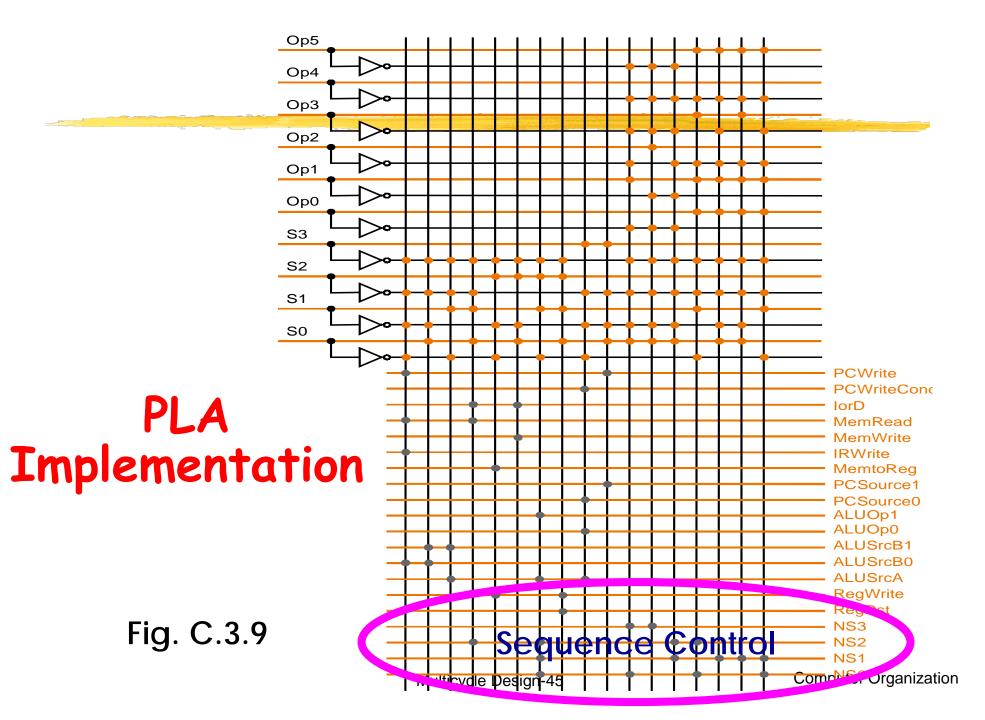
Output	Equation	<u>-</u>								
PCWrite	state0 + state9									
PCWriteCond state8										
IorD	state3 + state5									
•••										
NextState0	Output			Cı	urre	ent	tst	ate	<u>es</u>	
	C) 1	1 2	2 3	4	5	6	7	8	9
NextState1	PCWrite 1	() (0 (0	0	0	0	0	1
NextState2	PCWriteCond 0) () (0 (0	0	0	0	1	0
NextState3	lorD 0) () (1	0	1	0	0	0	0
•••		• •								

Designing FSM Controller



The Control Unit





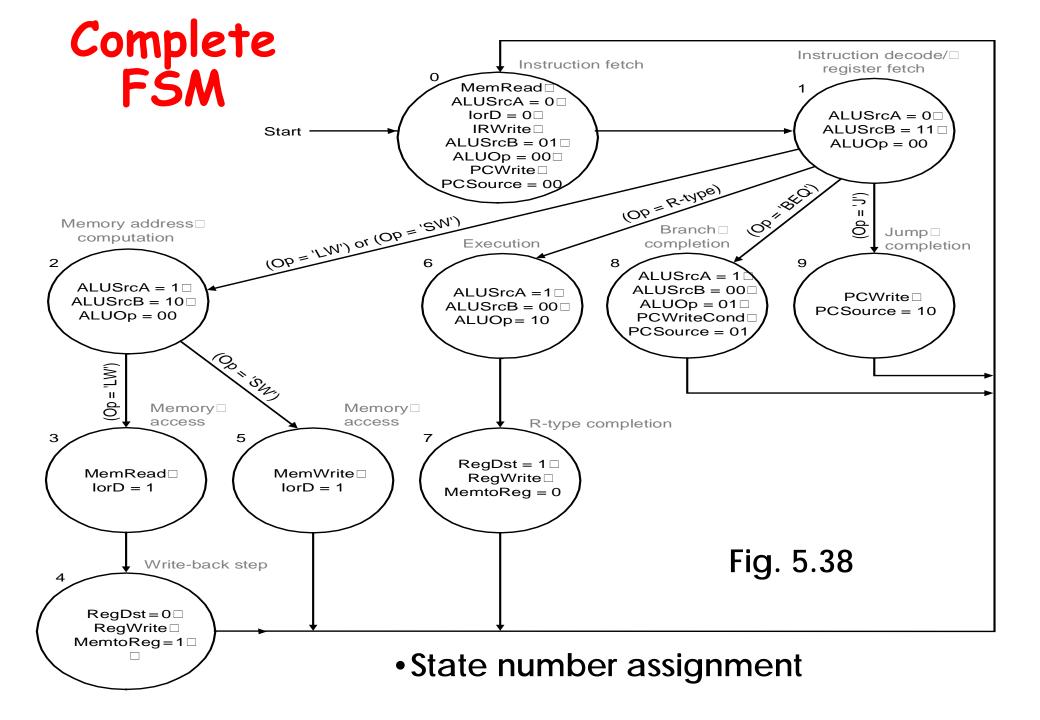
ROM Implementation(Truth Table)

	are a second sec	The second second second
Address	ROM content	NIO
<u>ob</u> 2	Datapath control	NS
00000 0000	1001010000001000	0001
000000 0001	000000000011000	0110
000000 0010	0000000000010100	XXXX
:		
000000 1010		
:		
000000 1111		
000001 0000		
:		
000010 0000	1001010000001000	0001
000010 0001	000000000011000	1001
000000 0010	0000000000010100	XXXX
:		

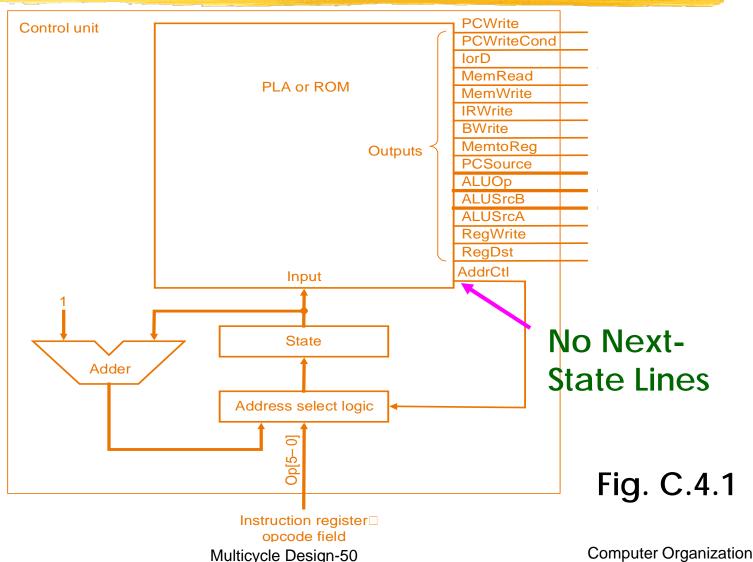
- Rather wasteful, since for lots of entries, outputs are same or are don'tcare
- Could break up into two smaller ROMs (Fig. C.3.7, C.3.8)

ROM vs. PLA

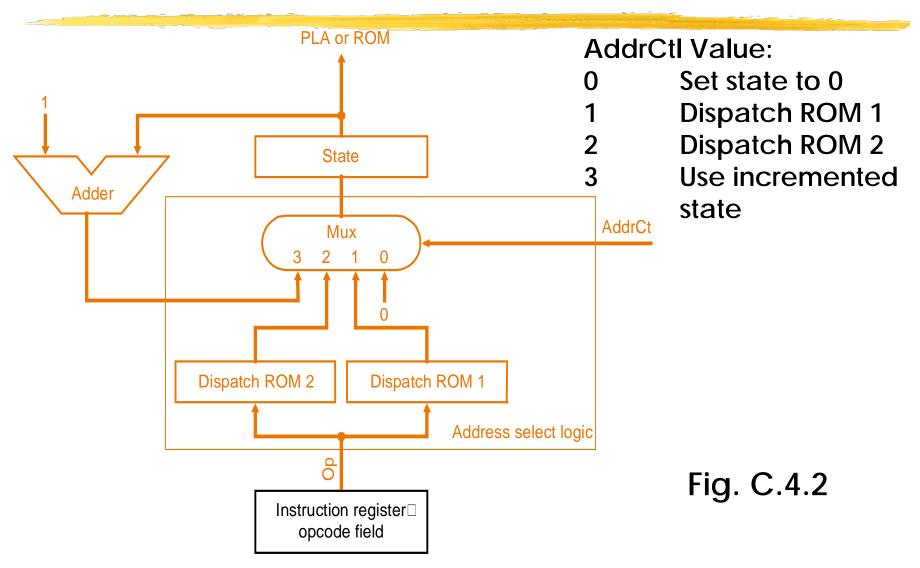
- ROM: use two smaller ROMs (Fig. C.3.7, C.3.8)
 - 4 state bits give the 16 outputs, 24x16 bits of ROM
 - 10 bits (op + state) give 4 next state bits, 2¹⁰x 4 bits of ROM
 - Total = 4.3K bits of ROM (compared to 2¹⁰x 20 bits of single ROM implementation)
- PLA is much smaller
 - can share product terms
 - only need entries that produce an active output
 - can take into account don't-cares
 - Size is (#inputs × 2-input-phase × #product-terms) + (#outputs × #product-terms)
 For this example = (10 × 2 × 17) + (20 × 17) = 680 PLA cells
- PLA cell usually about the size of a ROM cell (slightly bigger)



Use Counter for Sequence Control



Address Select Unit



Control Contents

Dispatch ROM 1					
Op	Opcode name	Value			
000000	R-format	0110			
000010	jmp	1001			
000100	beq	1000			
100011	lw	0010			
101011	SW	0010			

Dispatch ROM 2					
Ор	Opcode name	Value			
100011	lw	0011			
101011	SW	0101			

Fig. C.4.3, C.4.4

State number	Address-control action	Value of AddrCtl
0 (0000)	Use incremented state	3
1 (0001)	Use dispatch ROM 1	1
2 (0010)	Use dispatch ROM 2	2
3 (0011)	Use incremented state	3
4 (0100)	Replace state number by 0	0
5 (0101)	Replace state number by 0	0
6 (0110)	Use incremented state	3
7 (0111)	Replace state number by 0	0
8 (1000)	Replace state number by 0	0
9 (1001)	Replace state number by 0	0

Outline

- Designing a processor
- Building the datapath
- A single-cycle implementation
- A multicycle implementation
 - Multicycle datapath
 - Multicycle execution steps
 - Multicycle control
- Microprogramming: simplifying control (Appendix C.4)
- Exceptions

Controller Design: An Overview

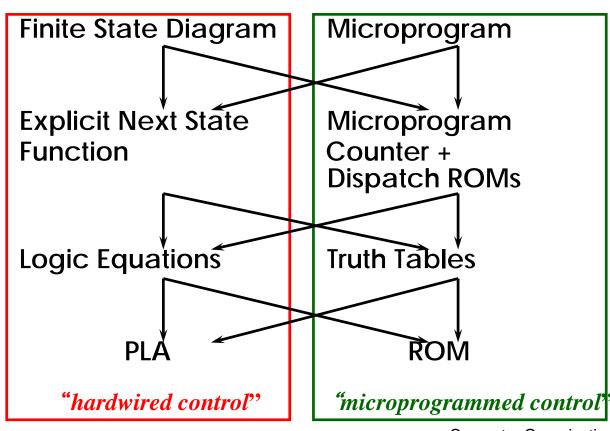
 Several possible initial representations, sequence control and logic representation, and control implementation => all may be determined indep.

Initial Rep.

Sequencing Control

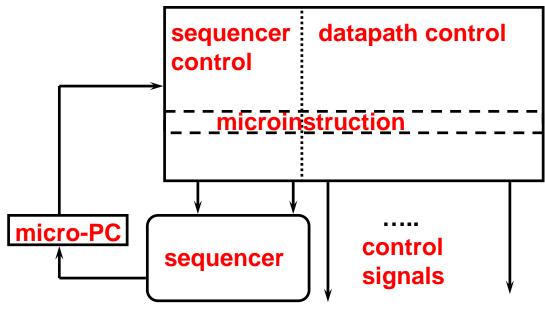
Logic Rep.

Implementation



Microprogram

- Control is the hard part of processor design
 - Datapath is fairly regular and well-organized
 - Memory is highly regular
 - Control is irregular and global
- But, the state diagrams that define the controller for an instruction set processor are highly structured
 - Use this structure to construct a simple "microsequencer"
 - Control reduces to programming this simple device
 - => microprogramming



Microinstruction

- Control signals :
 - Think of the set of control signals that must be asserted in a state as an instruction
 - Executing a microinstruction has the effect of asserting the control signal specified by the microinstruction
- Sequencing
 - What microinstruction should be executed next?
 - Execute sequentially (next state unconditionally)
 - Branch (next state also depends on inputs)
- A microprogram is a sequence of microinstructions executing a program flow chart (finite state machine)

Designing a Microinstruction Set

- 1) Start with a list of control signals
- 2) Group signals together that make sense (vs. random): called *fields*
- Places fields in some logical order (ex: ALU operation & ALU operands first and microinstruction sequencing last)
- 4) Create a symbolic legend for the microinstruction format, showing name of field values and how they set control signals
 - Use computers to design computers
- 5) To minimize the width, encode operations that will never be used at the same time

1-3) Control Signals and Fields

	Signal name	Effect when deasserted	Effect when asserted
10	ALUSrcA	1st ALU operand = PC	1st ALU operand = Reg[rs]
ontr	RegWrite	None	Reg file is written
N	MemtoReg	Reg. write data input = ALU	Reg. write data input = MDR RegDst
C	_	Reg. write dest. no. = rt	Reg. write dest. no. = rd
ij	MemRead	None	Memory at address is read
Bit	MemWrite	None	Memory at address is written
le	IorD	Memory address = PC	Memory address = ALUout
Single	IRWrite	None	IR = Memory
ii	PCWrite	None	PC = PCSource
	PCWriteCor	nd None	If zero then PC = PCSource

o	Signal name	Value	Effect
#	ALUOp	00	ALU adds
N	-	01	ALU subtracts
C		10	ALU operates according to func code
;;	ALUSrcB	00	2nd ALU input = B
B		01	2nd ALU input = 4
j		10	2nd ALU input = sign extended IR[15-0]
d		11	2nd ALU input = sign extended, shift left 2 IR[15-0]
11	PCSource	00	PC = ALU (PC + 4)
		01	PC = ALUout (branch target address)
Z		10	PC = PC+4[31-28] : IR[25-0] << 2

4) Fields and Legend

Field Name	Values for Field	Function of Field with Specific Value
ALU control	Add	ALU adds
	Subt.	ALU subtracts
	Func code	ALU does function code
SRC1	PC	1st ALU input = PC
	Α	1st ALU input = A (Reg[rs])
SRC2	В	2nd ALU input = B (Reg[rt])
	4	2nd ALU input = 4
	Extend	2nd ALU input = sign ext. IR[15-0]
	Extshft	2nd ALU input = sign ext., sl-2 IR[15-0]
Register control	Read	A = Reg[rs]; B = Reg[rt];
	Write ALU	Reg[rd] = ALUout
	Write MDR	Reg[rt] = MDR
Memory	Read PC	IR (MDR) = mem[PC]
J	Read ALU	MDR = mem[ALUout]
	Write ALU	mem[ALUout] = B
PC write	ALU	PC = ALU output
	ALUout-cond.	IF ALU zero then PC = ALUout
	jump addr.	PC = PCSource
Sequencing	Seq	Go to sequential microinstruction
	Fetch	Go to the first microinstruction
	Dispatch 1	Dispatch using ROM1
	Dispatch 2	Dispatch using ROM2
	•	. 5

Control Signals

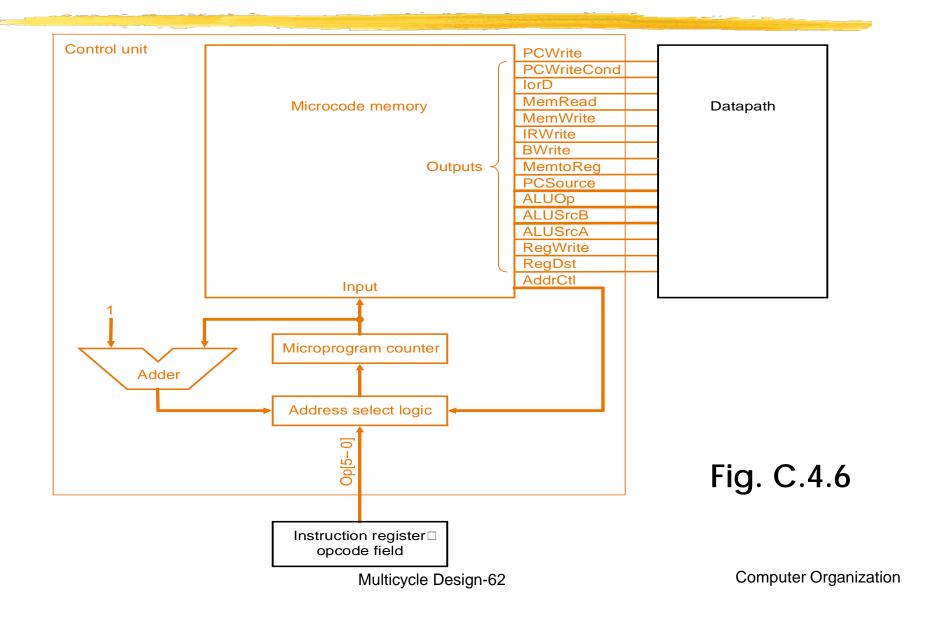
Field name	Value	Signals active	Comment			
	Add	ALUOp = 00	Cause the ALU to add.			
ALU control	Subt	ALUOp = 01	Cause the ALU to subtract; this implements the compare for			
			branches.			
	Func code	ALUOp = 10	Use the instruction's function code to determine ALU control.			
SRC1	PC	ALUSrcA = 0	Use the PC as the first ALU input.			
	A	ALUSrcA = 1	Register A is the first ALU input.			
	В	ALUSrcB = 00	Register B is the second ALU input.			
SRC2	4	ALUSrcB = 01	Use 4 as the second ALU input.			
	Extend	ALUSrcB = 10	Use output of the sign extension unit as the second ALU input.			
	Extshft	ALUSrcB = 11	Use the output of the shift-by-two unit as the second ALU input.			
	Read		Read two registers using the rs and rt fields of the IR as the register			
			numbers and putting the data into registers A and B.			
	Write ALU	RegWrite,	Write a register using the rd field of the IR as the register number and			
Register		RegDst = 1,	the contents of the ALUOut as the data.			
control		MemtoReg = 0				
	Write MDR	RegWrite,	Write a register using the rt field of the IR as the register number and			
		RegDst = 0,	the contents of the MDR as the data.			
		MemtoReg = 1				
	Read PC	MemRead,	Read memory using the PC as address; write result into IR (and			
		lorD = 0	the MDR).			
Memory	Read ALU	MemRead,	Read memory using the ALUOut as address; write result into MDR.			
		lorD = 1				
	Write ALU	MemWrite,	Write memory using the ALUOut as address, contents of B as the			
		lorD = 1	data.			
	ALU	PCSource = 00	Write the output of the ALU into the PC.			
		PCWrite				
PC write control	ALUOut-cond	PCSource = 01,	If the Zero output of the ALU is active, write the PC with the contents			
		PCWriteCond	of the register ALUOut.			
	jump address	PCSource = 10,	Write the PC with the jump address from the instruction.			
		PCWrite				
	Seq	AddrCtl = 11	Choose the next microinstruction sequentially.			
Sequencing	Fetch	AddrCtl = 00	Go to the first microinstruction to begin a new instruction.			
	Dispatch 1	AddrCtl = 01	Dispatch using the ROM 1.			
	Dispatch 2	AddrCtl = 10	Dispatch using the ROM 2.			

The Microprogram

	ALU			Register		PCWrite	
Label	control	SRC1	SRC2	control	Memory	control	Sequencing
Fetch	Add	PC	4		Read PC	ALU	Seq
	Add	PC	Extshft	Read			Dispatch 1
Mem1	Add	Α	Extend				Dispatch 2
LW2					Read ALU		Seq
				Write MDR			Fetch
SW2					Write ALU		Fetch
Rformat1	Func code	Α	В				Seq
				Write ALU			Fetch
BEQ1	Subt	Α	В			ALUOut-cond	Fetch
JUMP1						Jump address	Fetch

Fig. 5.7.3

The Controller



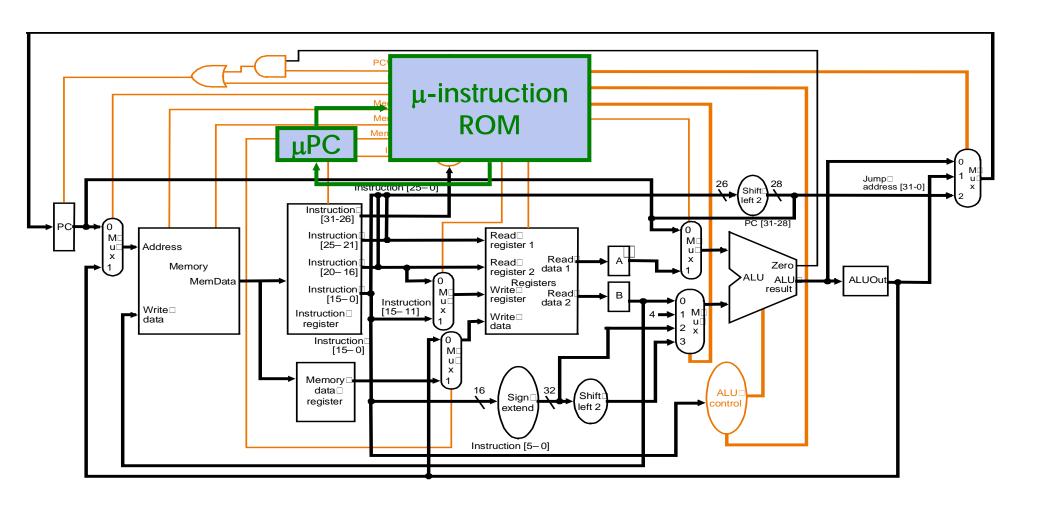
The Dispatch ROMs

Dispatch ROM 1						
Op	Opcode name	Value				
000000	R-format	Rformat1				
000010	jmp	JUMP1				
000100	ped	BEQ1				
100011	lw	Mem1				
101011	SW	Mem1				

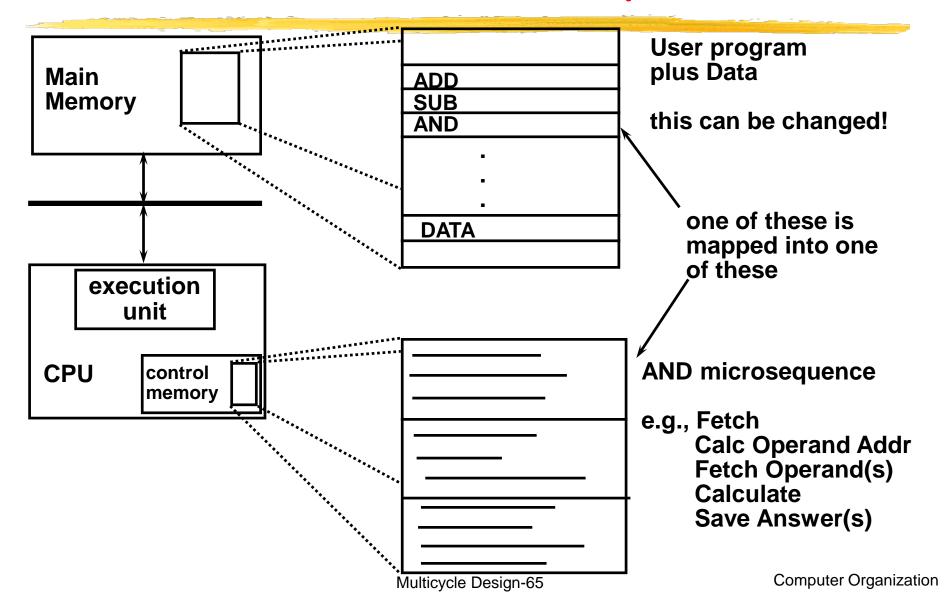
Dispatch ROM 2			
Op Opcode name		Value	
100011	lw	LW2	
101011	SW	SW2	

Fig. C.5.2

Our Plan: Using ROM



Microinstruction Interpretation



Microprogramming Using ROM: Pros and Cons

- Ease of design
- Flexibility
 - Easy to adapt to changes in organization, timing, technology
 - Can make changes late in design cycle, or even in the field
- Generality
 - Implement multiple inst. sets on same machine
 - Can tailor instruction set to application
 - Can implement very powerful instruction sets (just more control memory)
- Compatibility
 - Many organizations, same instruction set
- Costly to implement and slow

5) Microinstruction Encoding

State number	Control bit	ts 17- 2	Control bits 1- 0
0	1001 <mark>010000</mark>	00 <mark>0</mark> 1000	11
1	0000 <mark>0000000000000000000000000000000000</mark>	00 <mark>11000</mark>	01
2	0000 <mark>0000000000000000000000000000000000</mark>	00 <mark>10100</mark>	10
3	0011000000	00000	11
4	0000 <mark>001000</mark>	000010	00
5	0010100000	00000	00
6	0000 00000	10 <mark>00100</mark>	11
7	0000 <mark>0000000000000000000000000000000000</mark>	00 <mark>00011</mark>	00
8	0100 <mark>000010</mark>	0100	00
9	1000 <mark>000100</mark>	00000	00

Fig. C.4.5

 Bits 7-13 can be encoded to 3 bits because there are only 7 patterns of the control word

Minimal vs. Maximal Encoding

- Minimal (Horizontal):
- + more control over the potential parallelism of operations in the datapath
- uses up lots of control store
- Maximal (Vertical):
- + uses less number of control store
- extra level of decoding may slow the machine down

Recap: Designing a Microinstruction Set

- 1) Start with a list of control signals
- 2) Group signals together that make sense (vs. random): called *fields*
- Places fields in some logical order (ex: ALU operation & ALU operands first and microinstruction sequencing last)
- 4) Create a symbolic legend for the microinstruction format, showing name of field values and how they set control signals
 - Use computers to design computers
- 5) To minimize the width, encode operations that will never be used at the same time

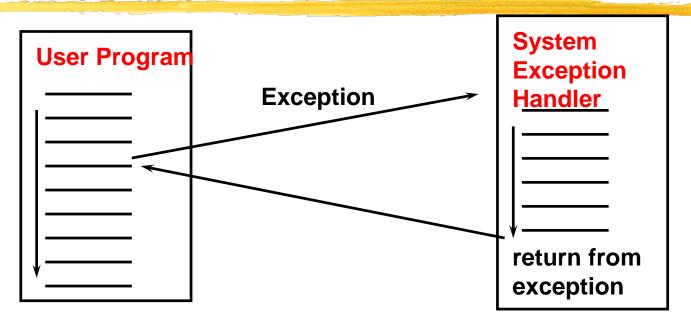
Summary of Control

- Control is specified by a finite state diagram
- Specialized state-diagrams easily captured by microsequencer
 - simple increment and "branch" fields
 - datapath control fields
- Control can also be specified by microprogramming
- Control is more complicated with:
 - complex instruction sets
 - restricted datapaths
- Simple instruction set and powerful datapath => simple control
 - could reduce hardware
 - Or go for speed => many instructions at once!

Outline

- Designing a processor
- Building the datapath
- A single-cycle implementation
- A multicycle implementation
 - Multicycle datapath
 - Multicycle execution steps
 - Multicycle control
- Microprogramming: simplifying control
- Exceptions

Exceptions



- Normal control flow: sequential, jumps, branches, calls, returns
- Exception = unprogrammed control transfer
 - system takes action to handle the exception
 - must record address of the offending instruction
 - should know cause and transfer to proper handler
 - if returns to user, must save & restore user state

User/System Modes

- By providing two modes of execution (user/system), computer may manage itself
 - OS is a special program that runs in the privileged system mode and has access to all of the resources of the computer
 - Presents "virtual resources" to each user that are more convenient than the physical resources
 - files vs. disk sectors
 - virtual memory vs. physical memory
 - protects each user program from others
- Exceptions allow the system to take action in response to events that occur while user program is executing
 - OS begins at the handler

Two Types of Exceptions

- Interrupts:
 - caused by external events and asynchronous to execution
 - → may be handled between instructions
 - simply suspend and resume user program
- Exceptions:
 - caused by internal events and synchronous to execution, ex: exceptional conditions (overflow), errors (parity), faults
 - instruction may be retried or simulated and program continued or program may be aborted

MIPS Convention of Exceptions

MIPS convention:

- exception means any unexpected change in control flow, without distinguishing internal or external
- use interrupt only when the event is externally caused

Type of event	From where?	MIPS terminology
I/O device request	External	Interrupt
Invoke OS from user program	Internal	Exception
Hardware malfunctions	Either	Exception or
		Interrupt
Arithmetic overflow	Internal	Exception
Using an undefined inst.	Internal	Exception

Precise Interrupts

- Precise: machine state is preserved as if program executed upto the offending inst.
 - Same system code will work on different implementations of the architecture
 - Position clearly established by IBM, and taken by MIPS
 - Difficult in the presence of pipelining, out-ot-order execution, ...
- Imprecise: system software has to figure out what is, where is, and put it all back together
- Performance goals often lead designers to forsake precise interrupts
 - system software developers, user, markets etc., usually wish they had not done this

Handling Exceptions in Our Design

- Consider two types of exceptions: undefined instruction & arithmetic overflow
- Basic actions on exception:
 - Save state: save the address of the offending instruction in the exception program counter (EPC)
 - Transfer control to OS at some specified address
 - need to know the cause for the exception
 - → then know the address of exception handler
 - After service, OS can terminate the program or continue its execution, using EPC to return

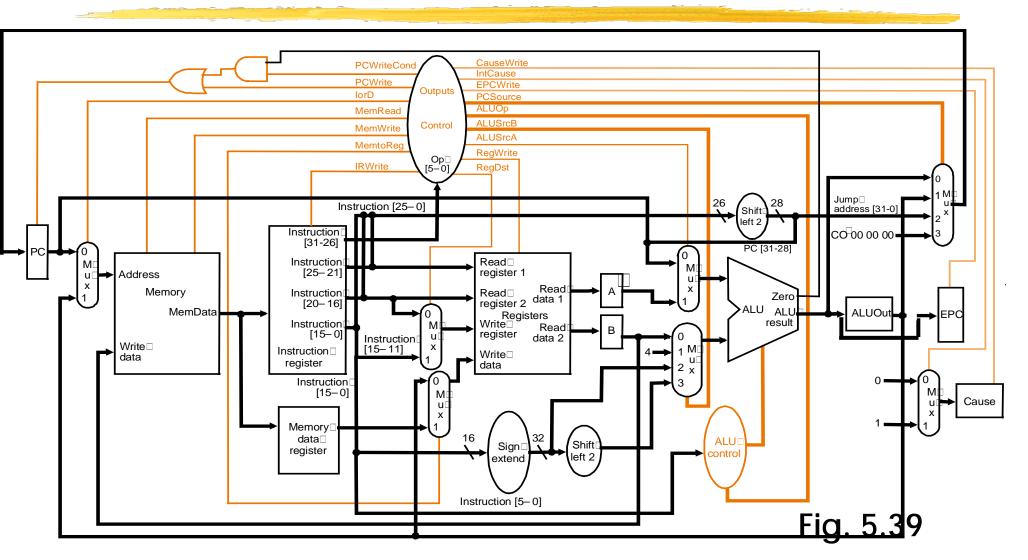
Saving State: General Approaches

- Push it onto the stack
 - Vax, 68k, 80x86
- Save it in special registers
 - MIPS EPC, BadVaddr, Status, Cause
- Shadow Registers
 - M88k
 - Save state in a shadow of the internal pipeline registers

Addressing the Exception Handler

Traditional approach: interrupt vector The cause of exception is a vector giving the address of the handler PC ← MEM[IV_base + cause | | 00] handlei 68000, Vax, 80x86, . . . iv_base code cause RISC Handler Table PC ← IV_base + cause | | 0000 Saves state and jumps Sparc, PA, M88K, . . . MIPS approach: fixed entry iv base use a status register (cause register) cause to hold a field to indicate the cause PC ← EXC_addr

Datapath with Exception Handling



Additions for Our Design

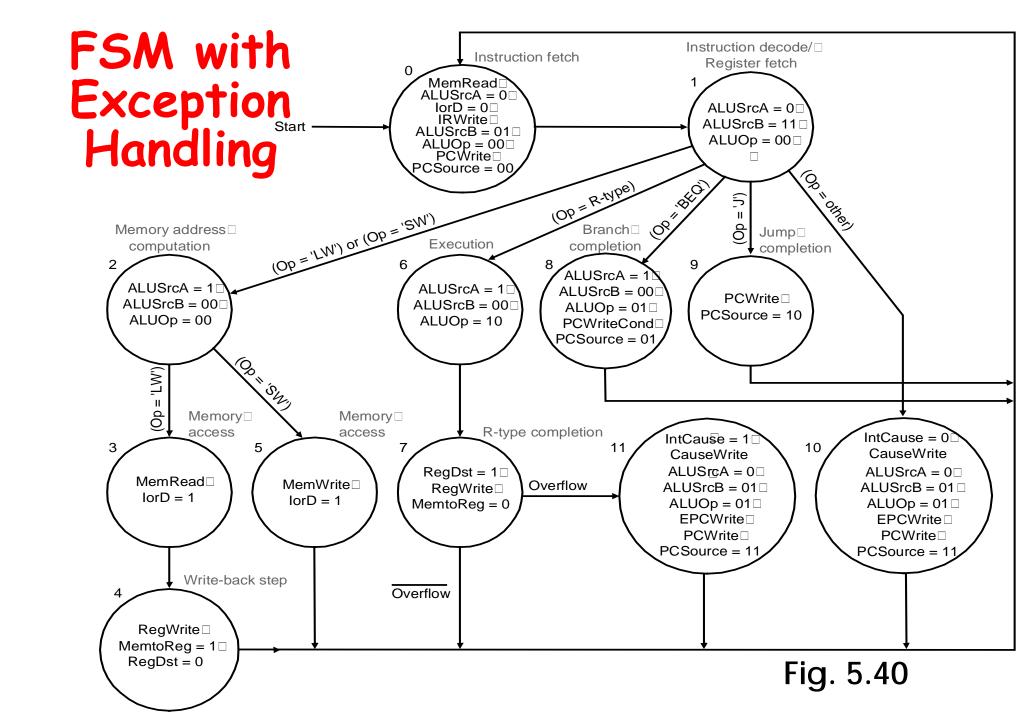
- EPC: reg. to hold address of affected inst.
- Cause: reg. to record cause of exception
 - Assume LSB encodes the two possible exception sources: undefined instruction=0 and arithmetic overflow=1
- Two control signals to write EPC (EPCWrite) and Cause (CauseWrite), and one control signal (IntCause) to set LSB of Cause register
- Be able to write exception address into PC, assuming at C000 0000_{hex}
 needs a 4-way MUX to PC
- May undo PC = PC + 4 (PC = PC 4), since want EPC to point to offending inst. (not its successor)

Exception Detection

- Undefined instruction: detected when no next state is defined from state 1 for the op value
 - Handle this by defining the next state value for all op values other than lw, sw, 0 (R-type), jmp, and beq as a new state, "other"
- Arithmetic overflow: detected with the Overflow signal out of the ALU
 - This signal is used in the modified FSM to specify an additional possible next state

Note: challenge in designing control of a real machine is to handle different interactions between instructions and other exception-causing events such that control logic remains small and fast

 Complex interactions makes the control unit the most challenging aspect of hardware design



Summary

- Specialize state diagrams easily captured by microsequencer
 - simple increment and branch fields
 - datapath control fields
- Control design reduces to microprogramming
- Exceptions are the hard part of control
 - Need to find convenient place to detect exceptions and to branch to state or microinstruction that saves PC and invokes OS
 - Harder with pipelined CPUs that support page faults on memory accesses, i.e., the instruction cannot complete AND you must restart program at exactly the instruction with the exception

Datapath with Exception Handling

