CS2006: 計算機組織

Pipelining

Outline

- An overview of pipelining
- A pipelined datapath
- Pipelined control
- Data hazards and forwarding
- Data hazards and stalls
- Branch hazards
- Exceptions
- Superscalar and dynamic pipelining

Pipelining Is Natural!

Laundry example:

Ann, Brian, Cathy, Dave each have one load of clothes to wash, dry, and fold



Dryer takes 40 minutes

"Folder" takes 20 minutes

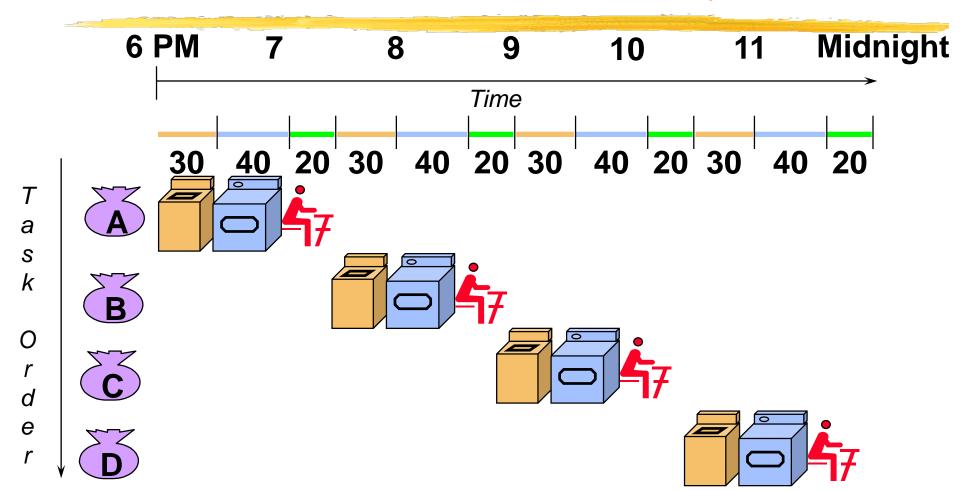








Sequential Laundry

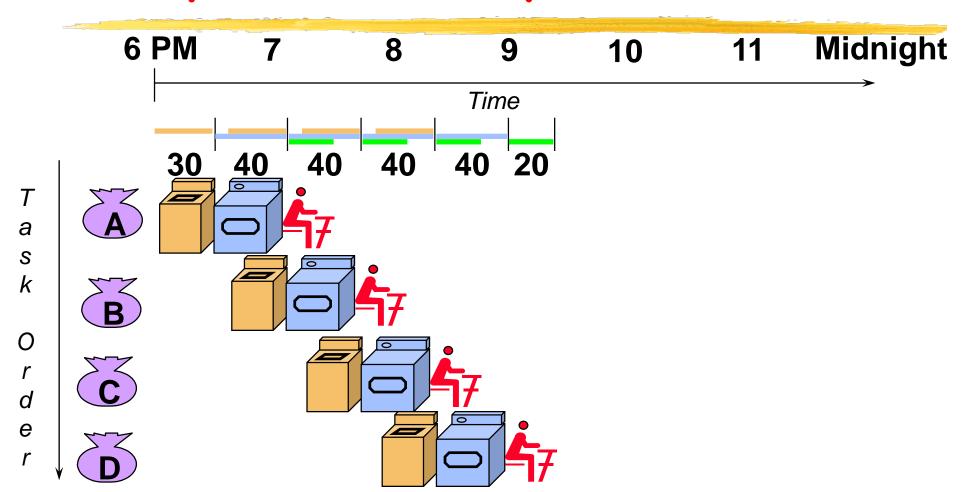


- Sequential laundry takes 6 hours for 4 loads
- If they learned pipelining, how long would it take?

Pipelining-4

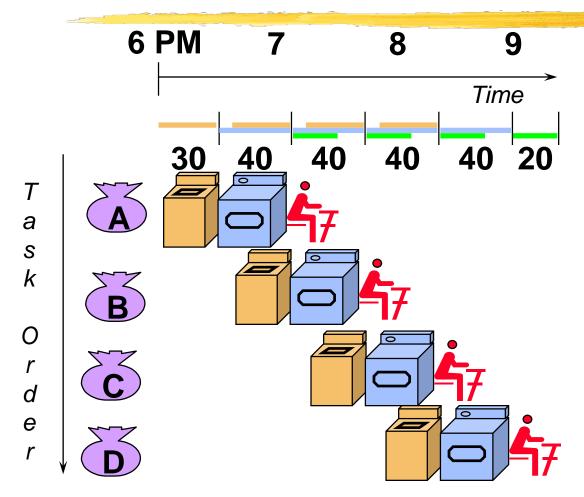
Computer Organization

Pipelined Laundry: Start ASAP



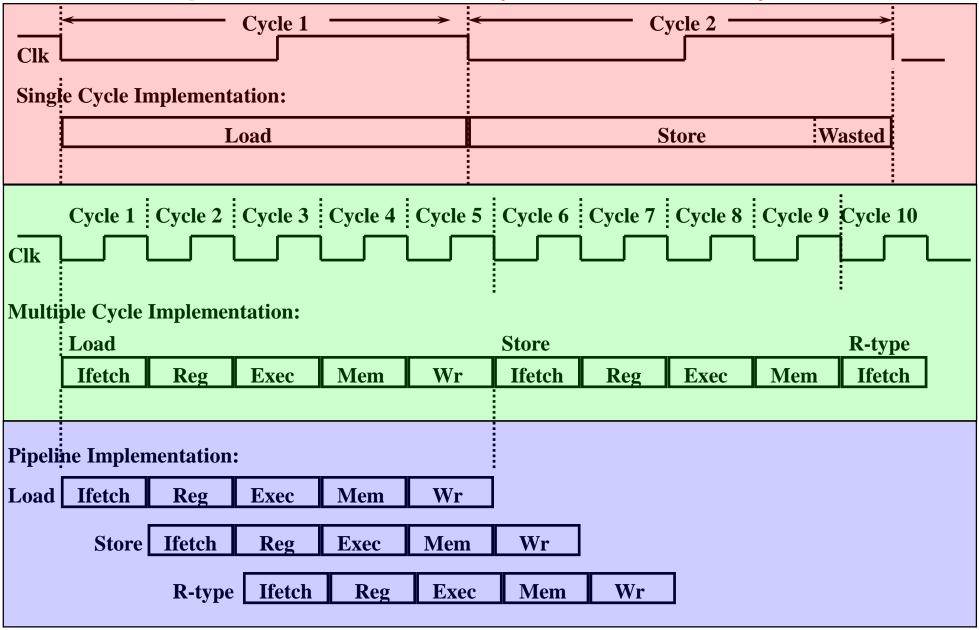
◆ Pipelined laundry takes 3.5 hours for 4 loads

Pipelining Lessons

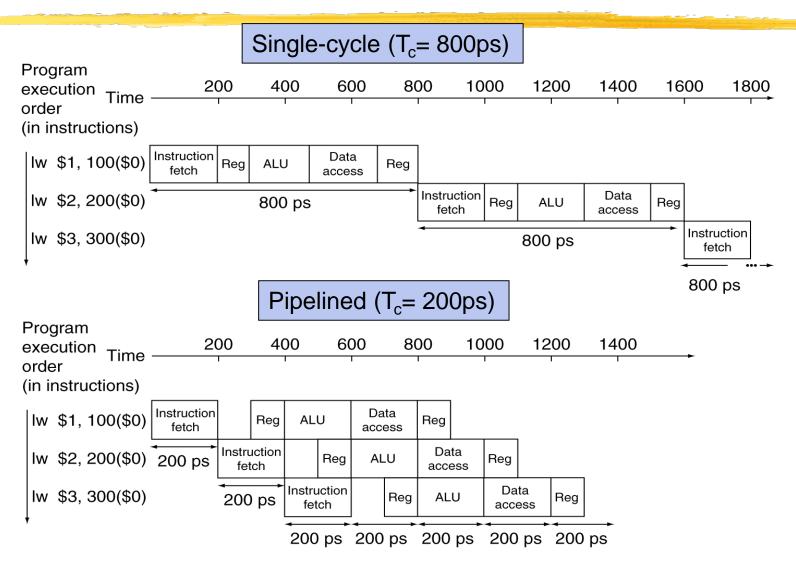


- Doesn't help latency of single task, but throughput of entire
- Pipeline rate limited by slowest stage
- Multiple tasks working at same time using different resources
- Potential speedup = Number pipe stages
- Unbalanced stage length; time to "fill" & "drain" the pipeline reduce speedup
- Stall for dependences

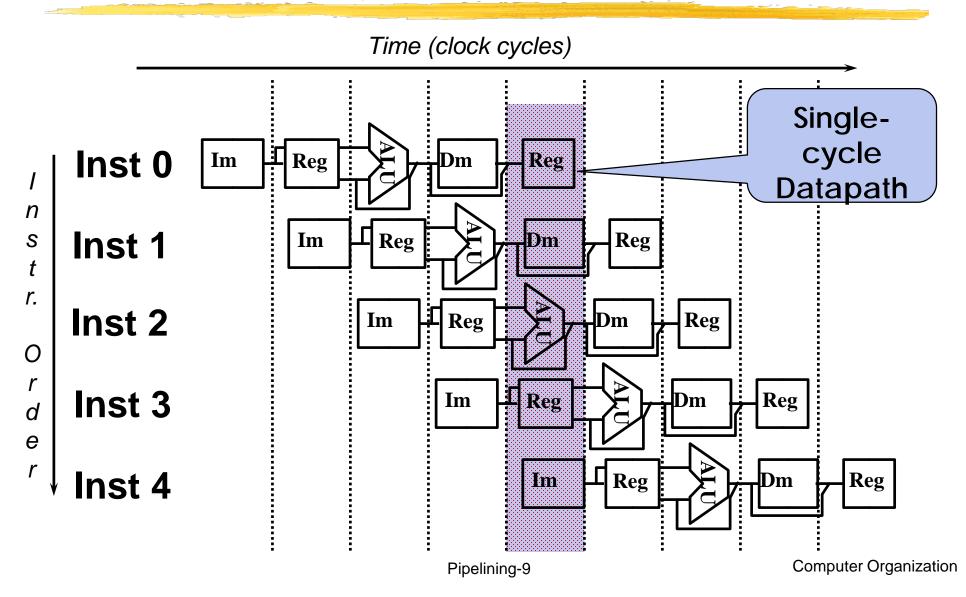
Single-, Multi-Cycle, vs. Pipeline



Pipeline Performance



Why Pipeline? Because the Resources Are There!



Outline

- An overview of pipelining
- A pipelined datapath
- Pipelined control
- Data hazards and forwarding
- Data hazards and stalls
- Branch hazards
- Exceptions
- Superscalar and dynamic pipelining

Designing a Pipelined Processor

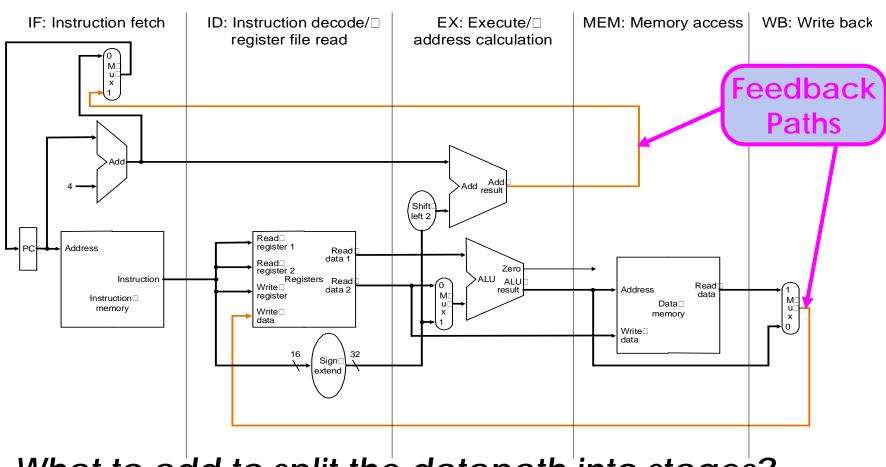
- Examine the datapath and control diagram
 - Starting with single- or multi-cycle datapath?
 - Single- or multi-cycle control?
- Partition datapath into stages:
 - IF (instruction fetch), ID (instruction decode and register file read), EX (execution or address calculation), MEM (data memory access), WB (write back)
- Associate resources with stages
- Ensure that flows do not conflict, or figure out how to resolve
- Assert control in appropriate stage

Use Multicycle Execution Steps

Step name	Action for R-type instructions	Action for memory-reference instructions	Action for branches	Action for jumps
Instruction fetch		IR = Memory[PC] PC = PC + 4		
Instruction decode/register fetch	A = Reg [IR[25-21]] B = Reg [IR[20-16]] ALUOut = PC + (sign-extend (IR[15-0]) << 2)			
Execution, address computation, branch/ jump completion	ALUOut = A op B	ALUOut = A + sign-extend (IR[15-0])	if (A ==B) then PC = ALUOut	PC = PC [31-28] II (IR[25-0]<<2)
Memory access or R-type completion	Reg [IR[15-11]] = ALUOut	Load: MDR = Memory[ALUOut] or Store: Memory [ALUOut] = B		
Memory read completion		Load: Reg[IR[20-16]] = MDR		

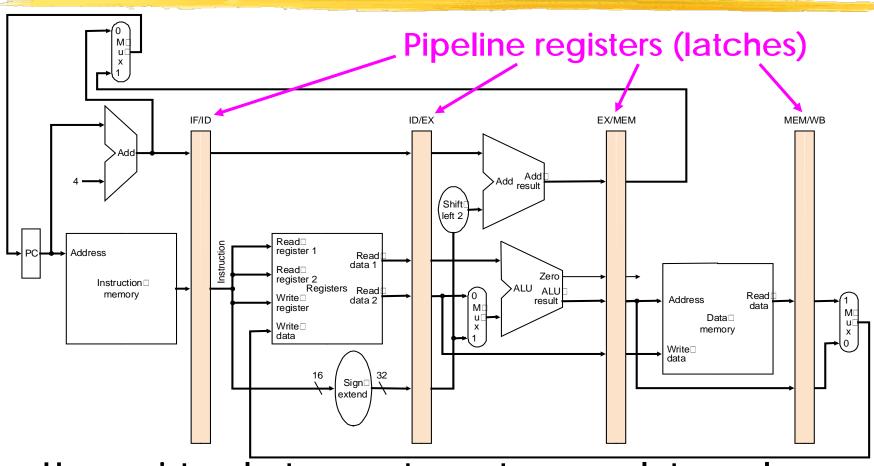
But, use single-cycle datapath ...

Split Single-cycle Datapath



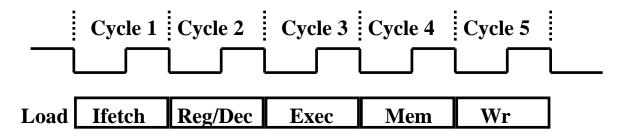
What to add to split the datapath into stages?

Add Pipeline Registers



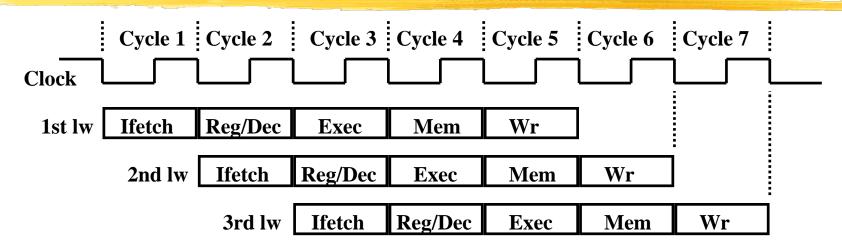
 Use registers between stages to carry data and control

Consider load



- IF: Instruction Fetch
 - Fetch the instruction from the Instruction Memory
- ID: Instruction Decode
 - Registers fetch and instruction decode
- EX: Calculate the memory address
- MEM: Read the data from the Data Memory
- WB: Write the data back to the register file

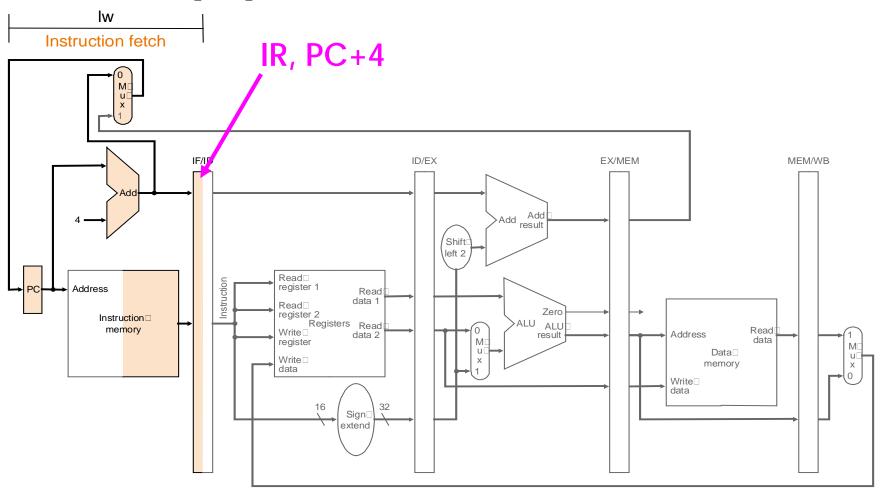
Pipelining load



- 5 functional units in the pipeline datapath are:
 - Instruction Memory for the Ifetch stage
 - Register File's Read ports (busA and busB) for the Reg/Dec stage
 - ALU for the Exec stage
 - Data Memory for the MEM stage
 - Register File's Write port (busW) for the WB stage

IF Stage of load

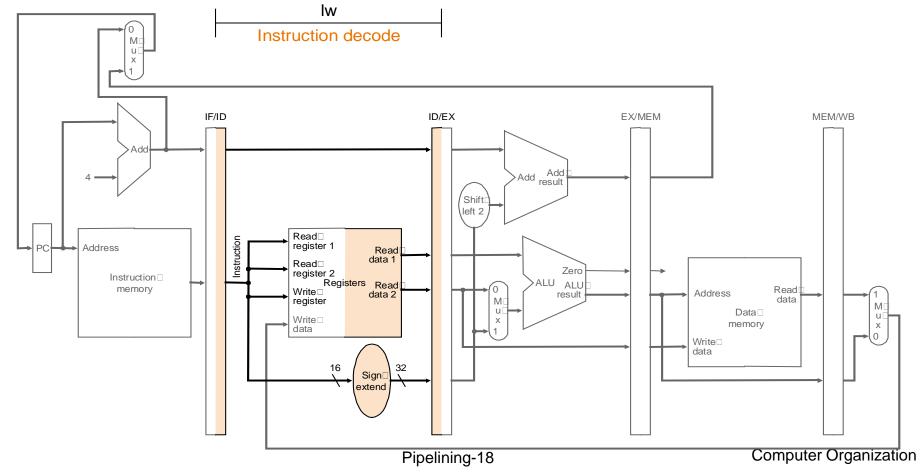
◆ IR = mem[PC]; PC = PC + 4



Pipelining-17

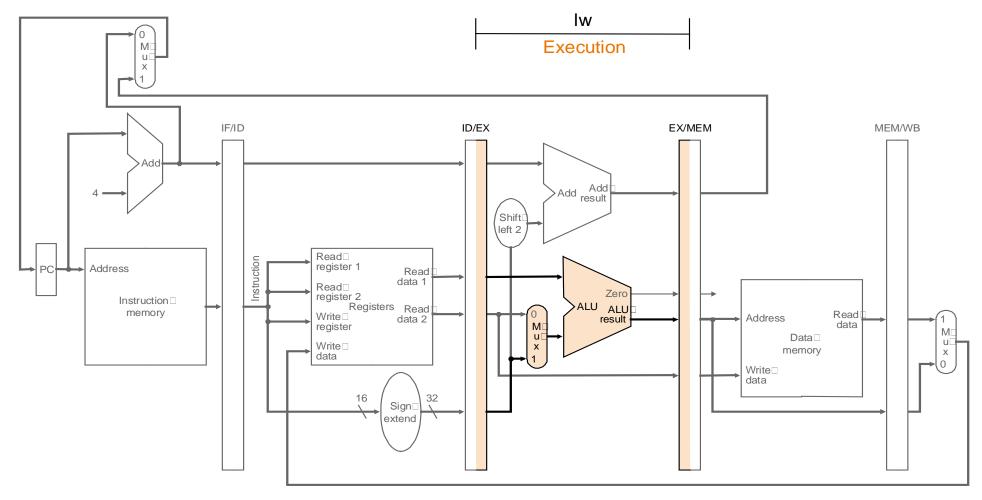
ID Stage of load

A = Reg[IR[25-21]]; B = Reg[IR[20-16]];
 ALUout = PC + (sign-ext(IR[15-0]) << 2) (some ops moved to the next stage)



EX Stage of load

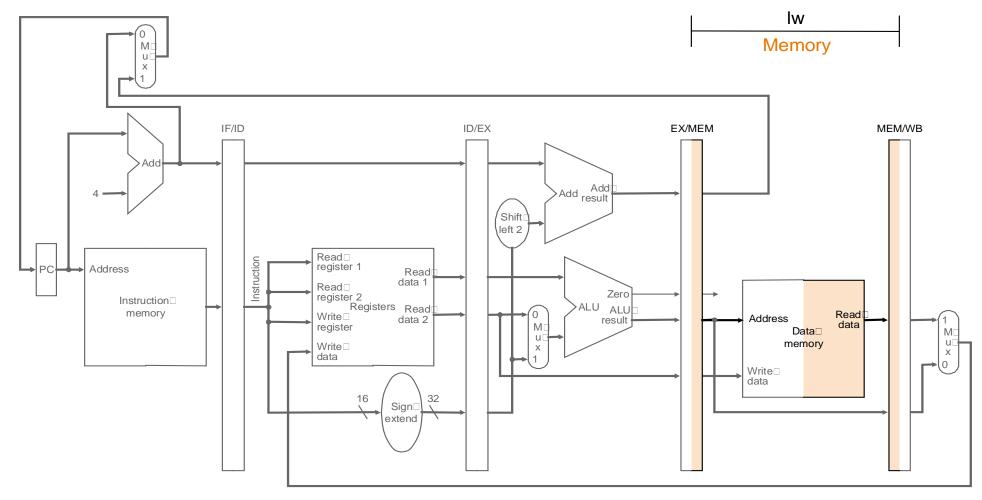
◆ ALUout = A + sign-ext(IR[15-0])



Pipelining-19

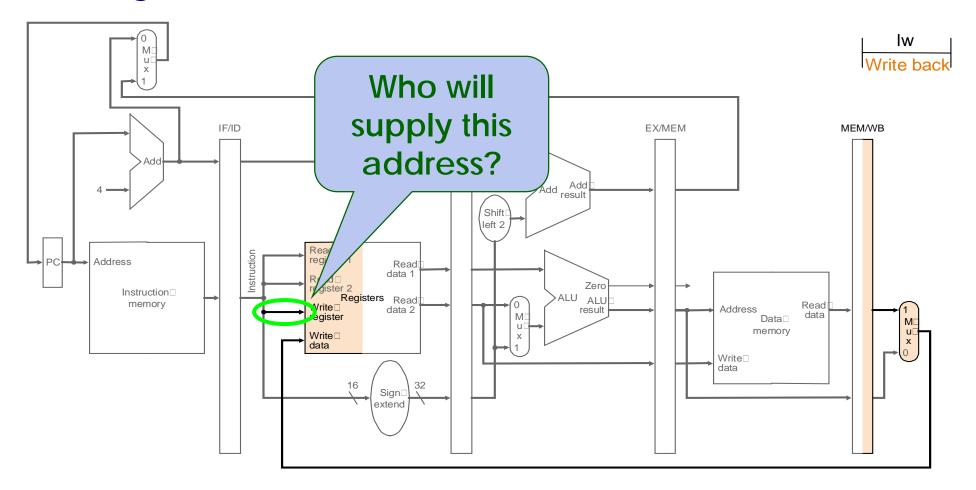
MEM State of load

MDR = mem[ALUout]

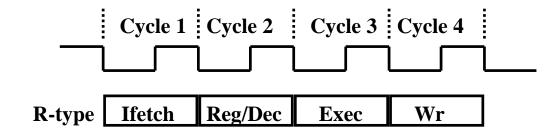


WB Stage of load

♦ Reg[IR[20-16]] = MDR

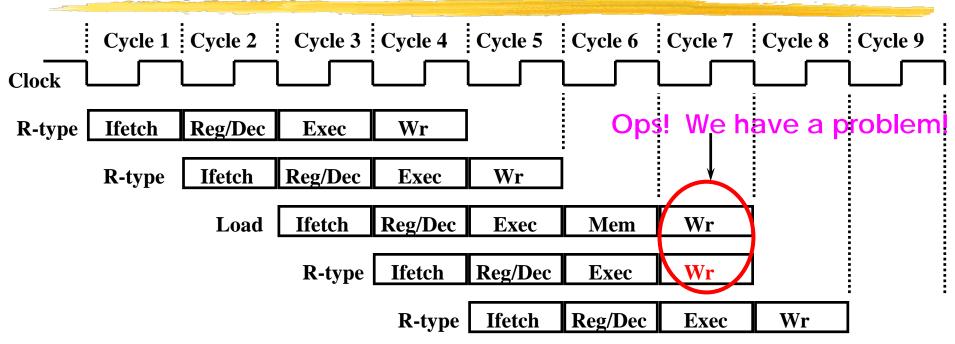


The Four Stages of R-type



- IF: fetch the instruction from the Instruction Memory
- ID: registers fetch and instruction decode
- EX: ALU operates on the two register operands
- WB: write ALU output back to the register file

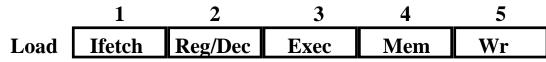
Pipelining R-type and load



- We have a structural hazard:
 - Two instructions try to write to the register file at the same time!
 - Only one write port

Important Observation

- Each functional unit can only be used once per instruction
- Each functional unit must be used at the same stage for all instructions:
 - Load uses Register File's write port during its 5th stage



R-type uses Register File's write port during its 4th stage

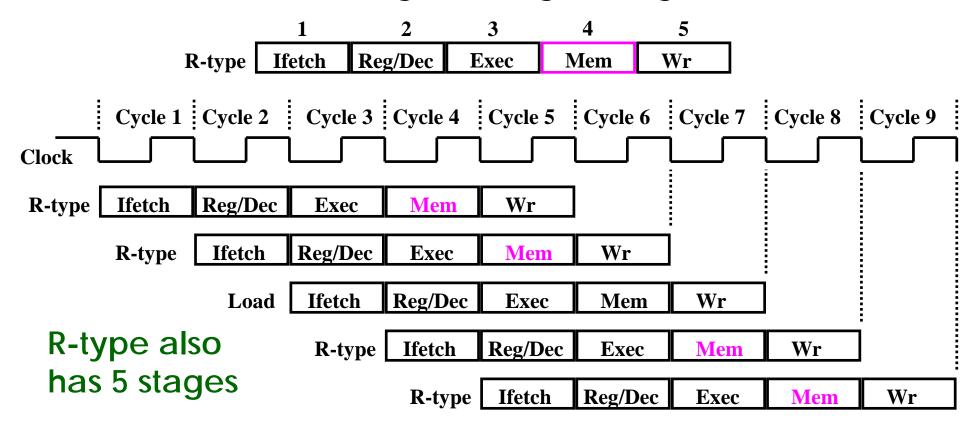
```
        1
        2
        3
        4

        R-type
        Ifetch
        Reg/Dec
        Exec
        Wr
```

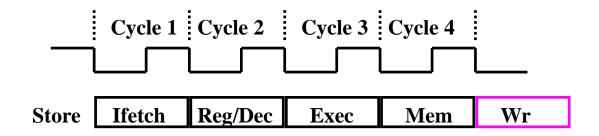
Several ways to solve: forwarding, adding pipeline bubble, making instructions same length

Solution: Delay R-type's Write

- Delay R-type's register write by one cycle:
 - R-type also use Reg File's write port at Stage 5
 - MEM is a NOP stage: nothing is being done



The Four Stages of store

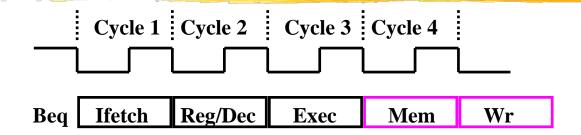


- IF: fetch the instruction from the Instruction Memory
- ID: registers fetch and instruction decode
- EX: calculate the memory address
- MEM: write the data into the Data Memory

Add an extra stage:

WB: NOP

The Three Stages of beq

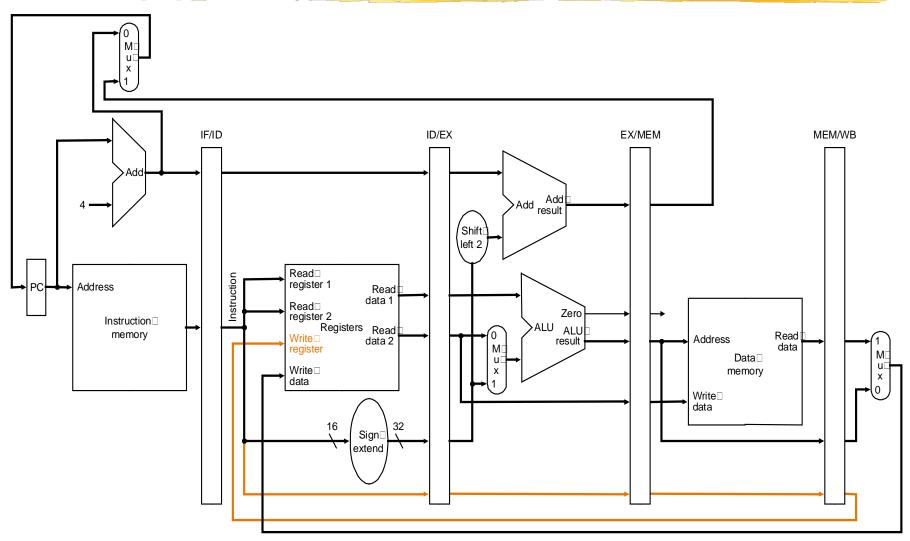


- IF: fetch the instruction from the Instruction Memory
- ID: registers fetch and instruction decode
- **♦ EX**:
 - compares the two register operand
 - select correct branch target address
 - latch into PC

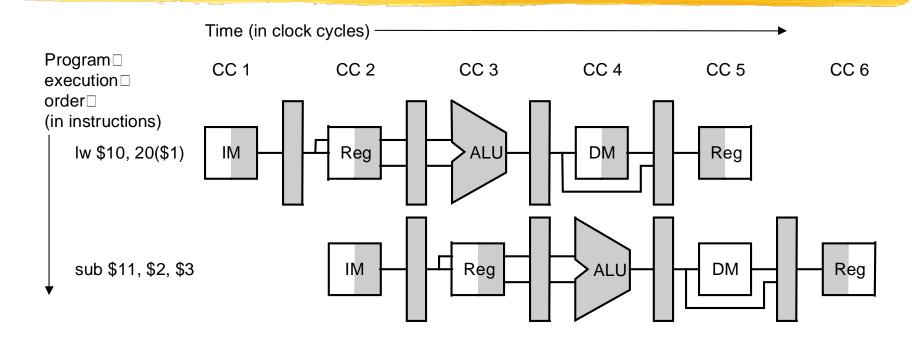
Add two extra stages:

- MEM: NOP
- WB: NOP

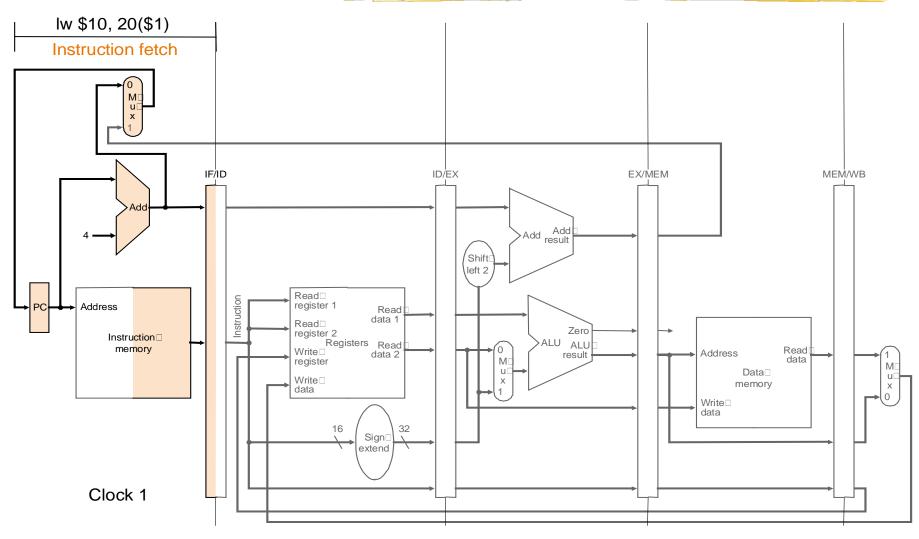
Pipelined Datapath

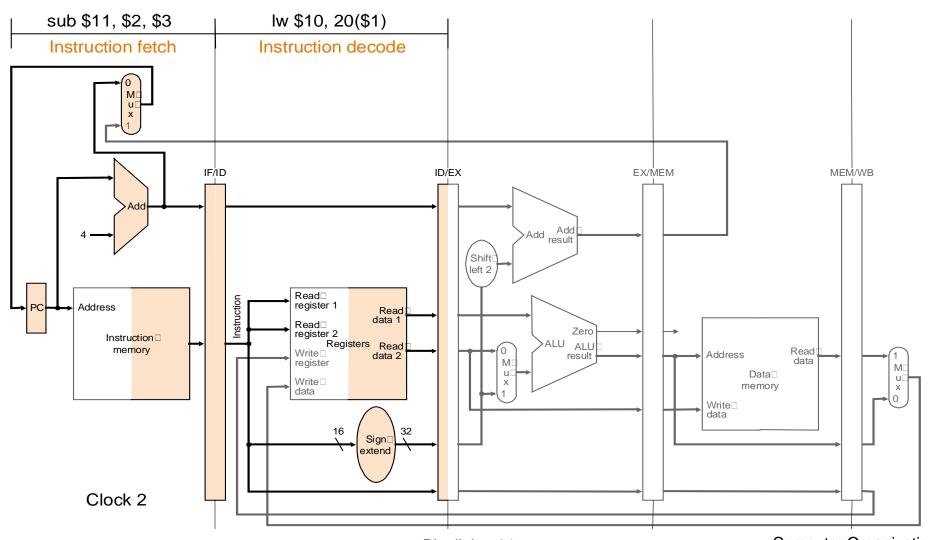


Graphically Representing Pipelines



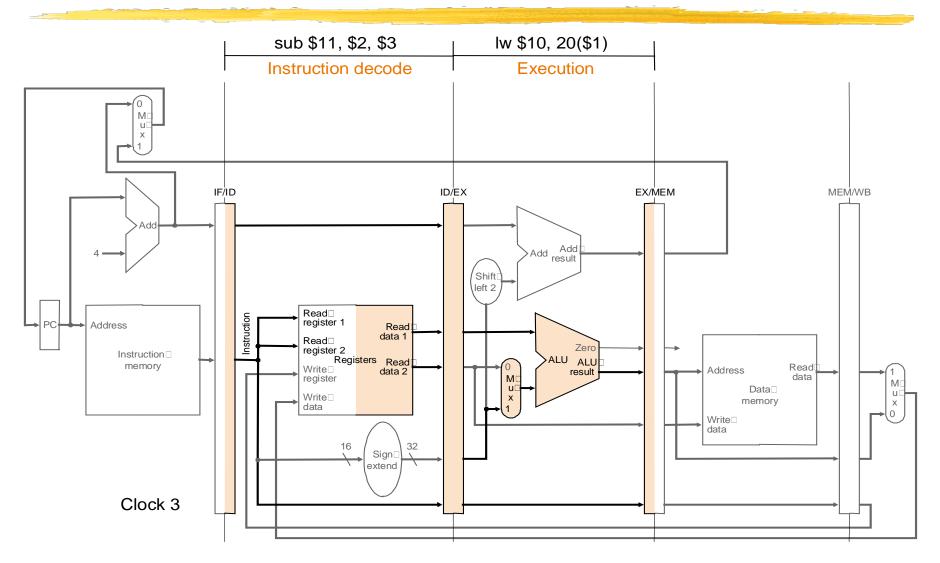
- Can help with answering questions like:
 - How many cycles to execute this code?
 - What is the ALU doing during cycle 4?
 - Help understand datapaths

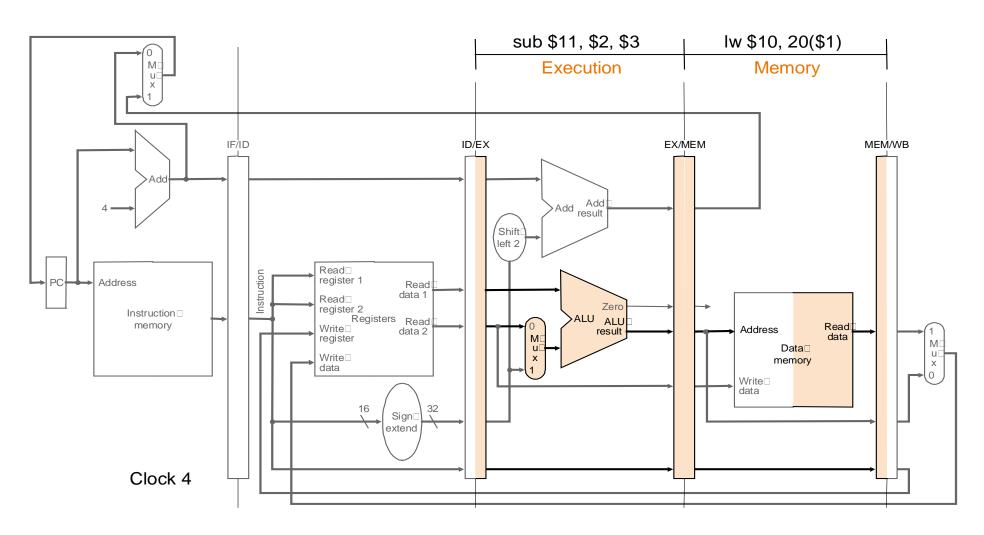


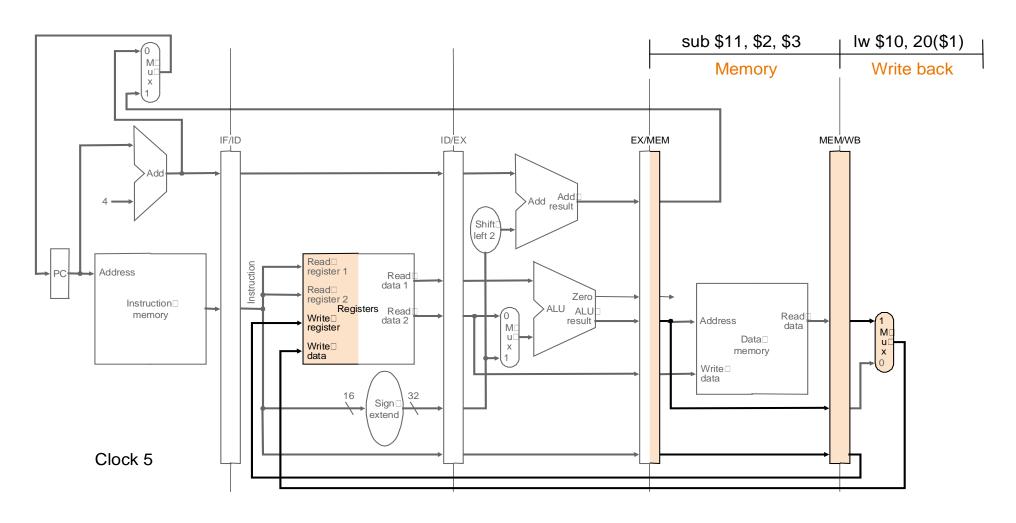


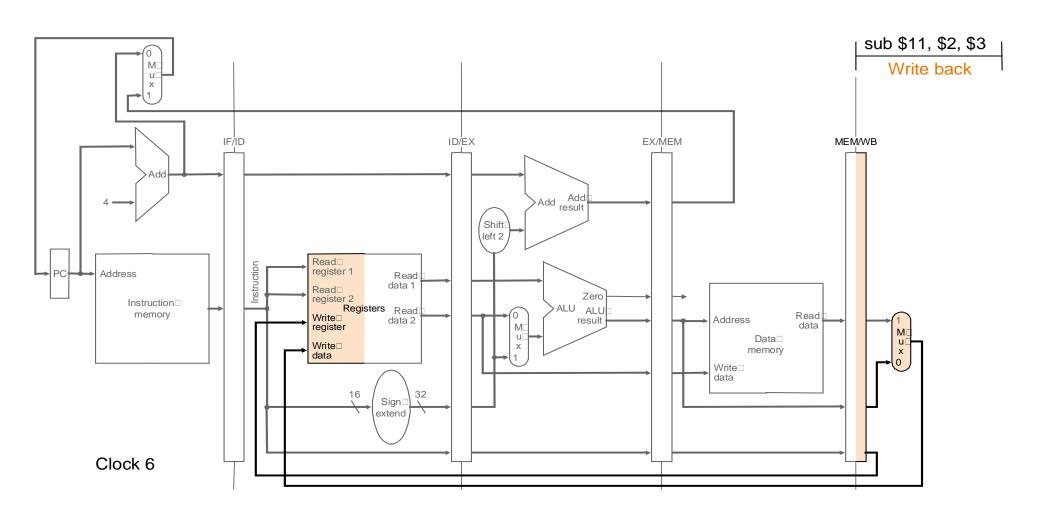
Pipelining-31

Computer Organization





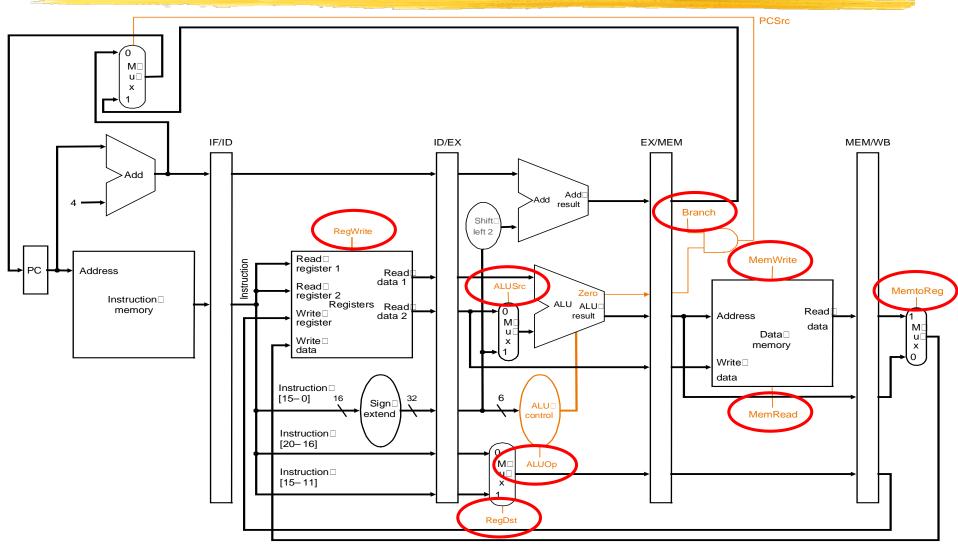




Outline

- An overview of pipelining
- A pipelined datapath
- Pipelined control
- Data hazards and forwarding
- Data hazards and stalls
- Branch hazards
- Exceptions
- Superscalar and dynamic pipelining

Pipeline Control: Control Signals



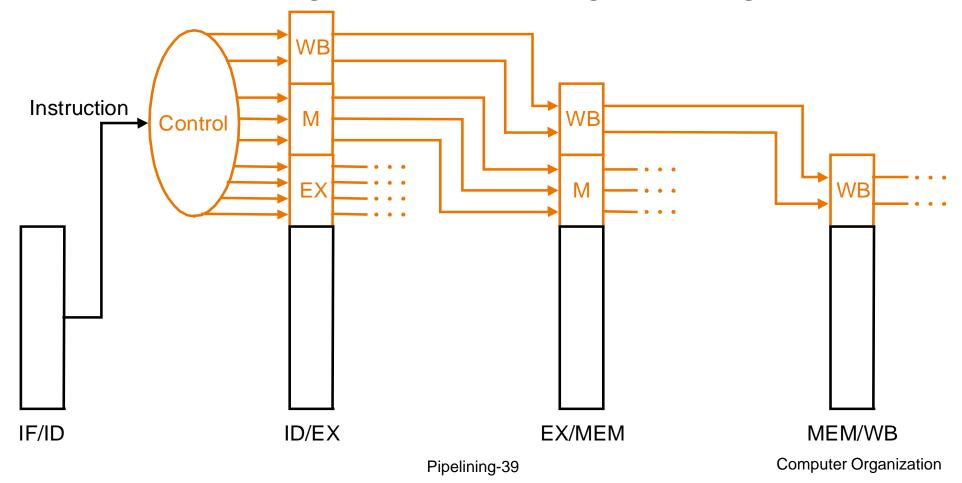
Group Signals According to Stages

We can use control signals of single-cycle CPU

	Execution/Address Calculation				Memory access stage			Write-back stage	
	stage control lines				control lines			control lines	
	Reg	ALU	ALU	ALU		Mem	Mem	Reg	Memto
Instruction	Dst	Op1	Op0	Src	Branch	Read	Write	write	Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

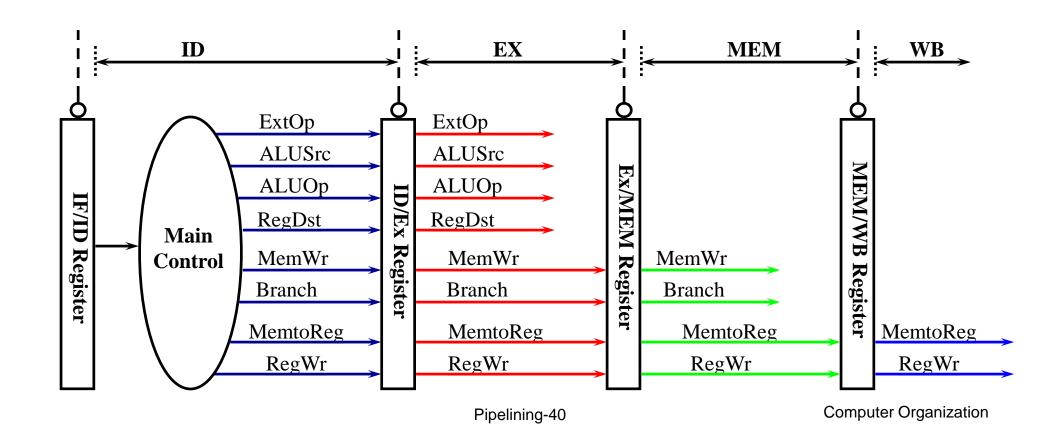
Data Stationary Control

- Pass control signals along just like the data
 - Main control generates control signals during ID

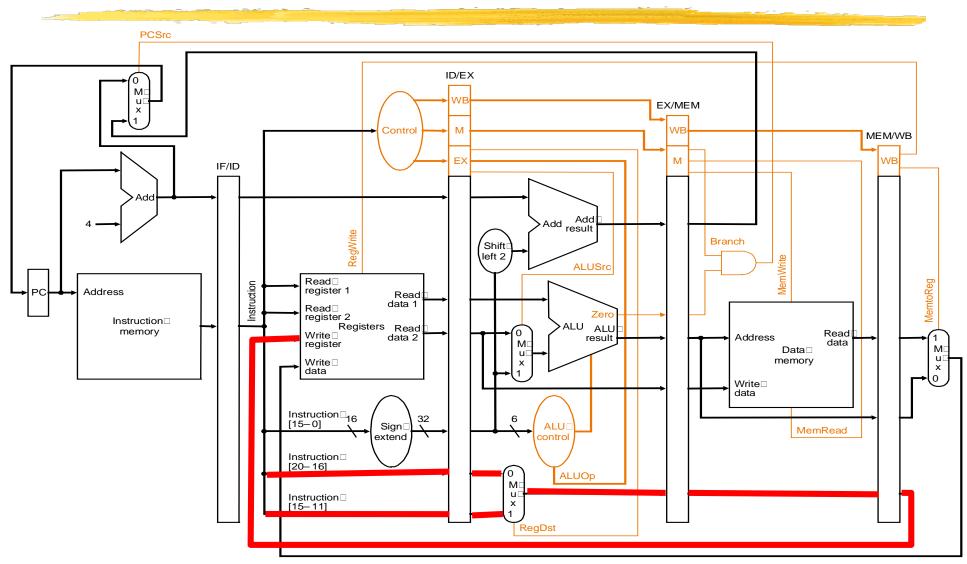


Data Stationary Control (cont.)

- Signals for EX (ExtOp, ALUSrc, ...) are used 1 cycle later
- Signals for MEM (MemWr, Branch) are used 2 cycles later
- Signals for WB (MemtoReg, MemWr) are used 3 cycles later

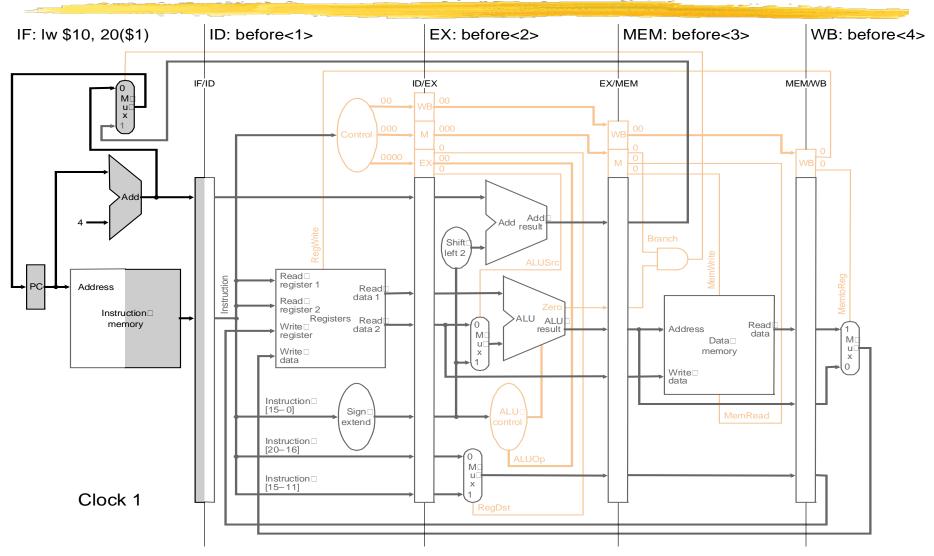


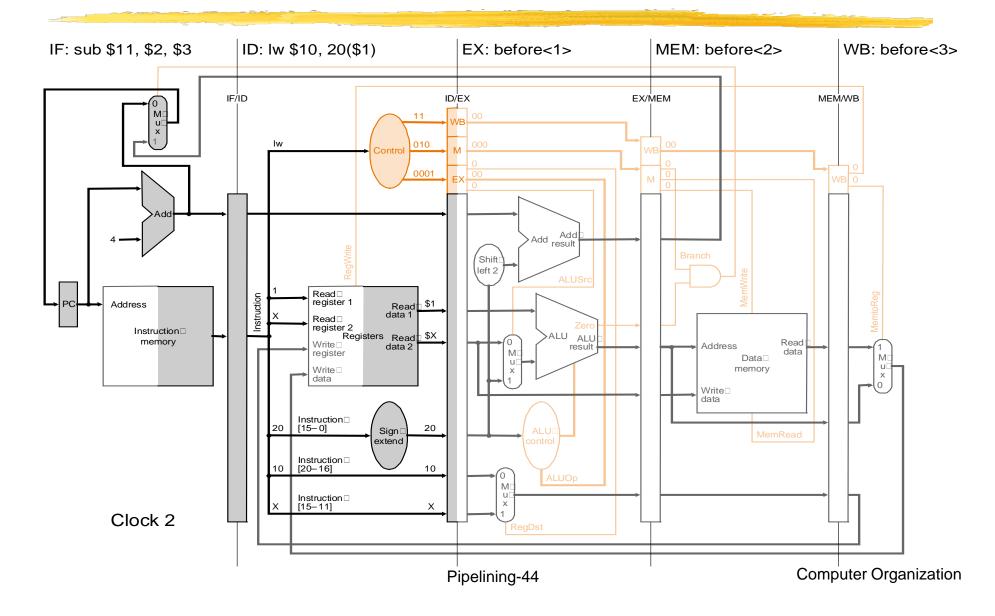
Datapath with Control

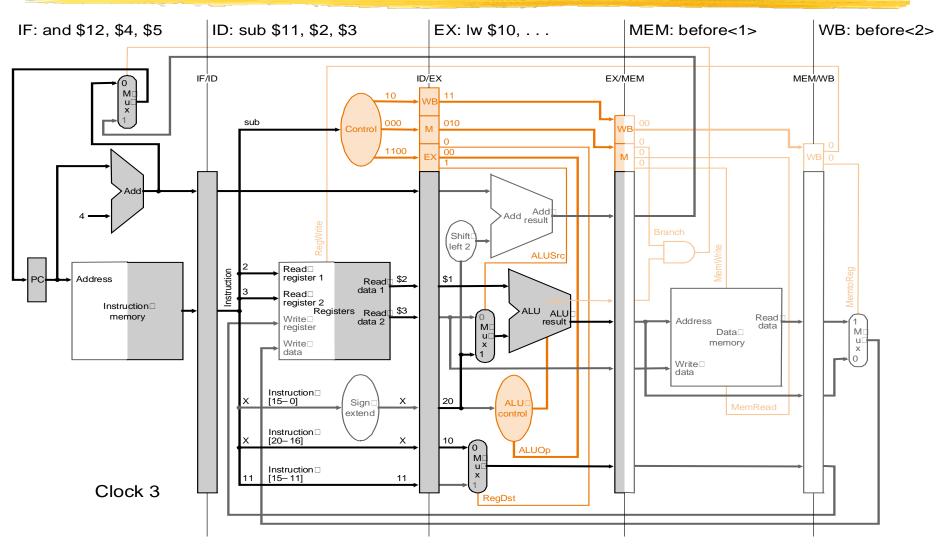


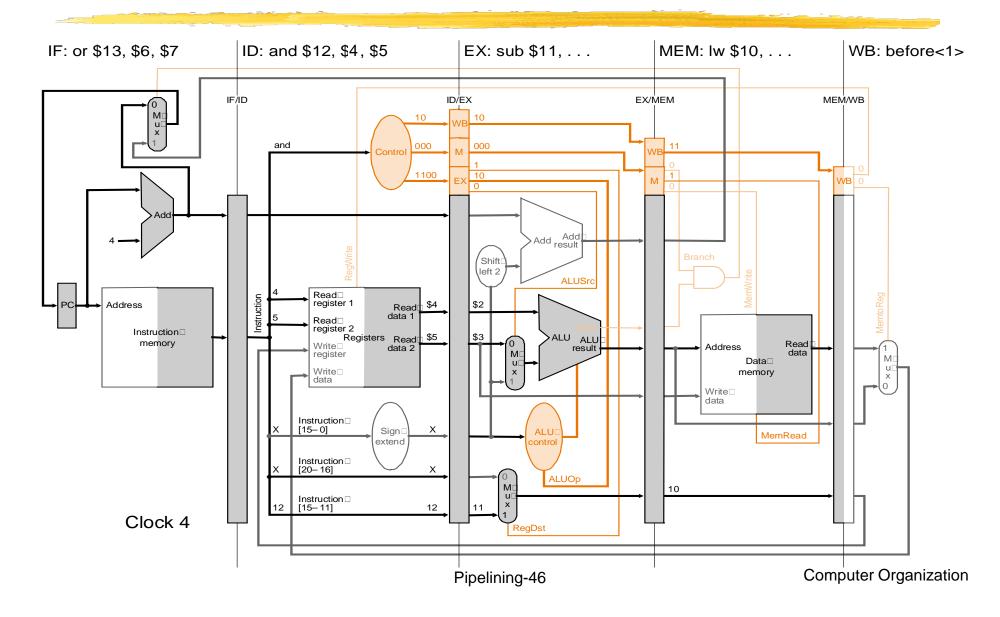
Let's Try it Out

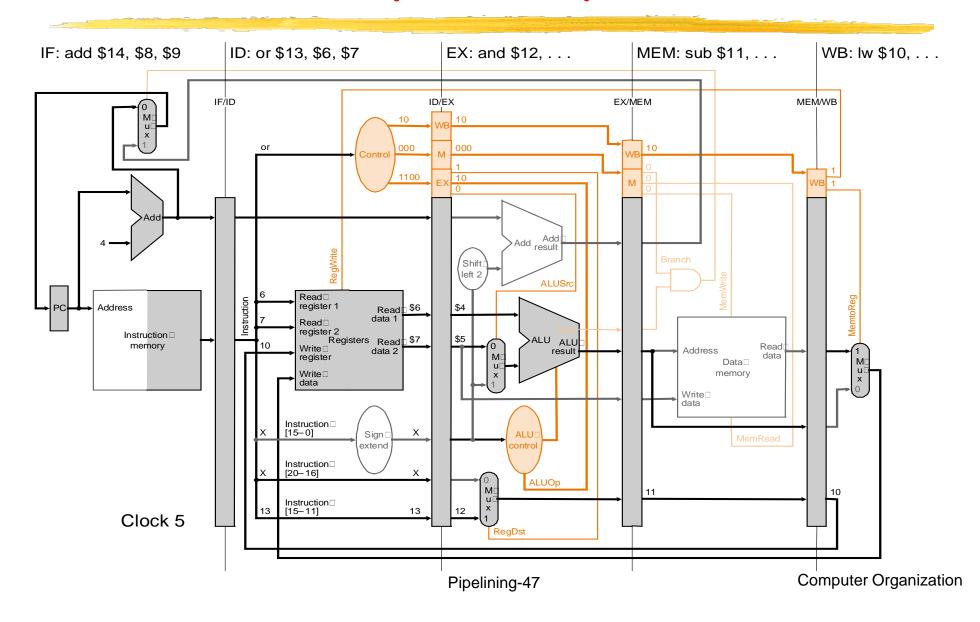
```
lw $10, 20($1)
sub $11, $2, $3
and $12, $4, $5
or $13, $6, $7
add $14, $8, $9
```

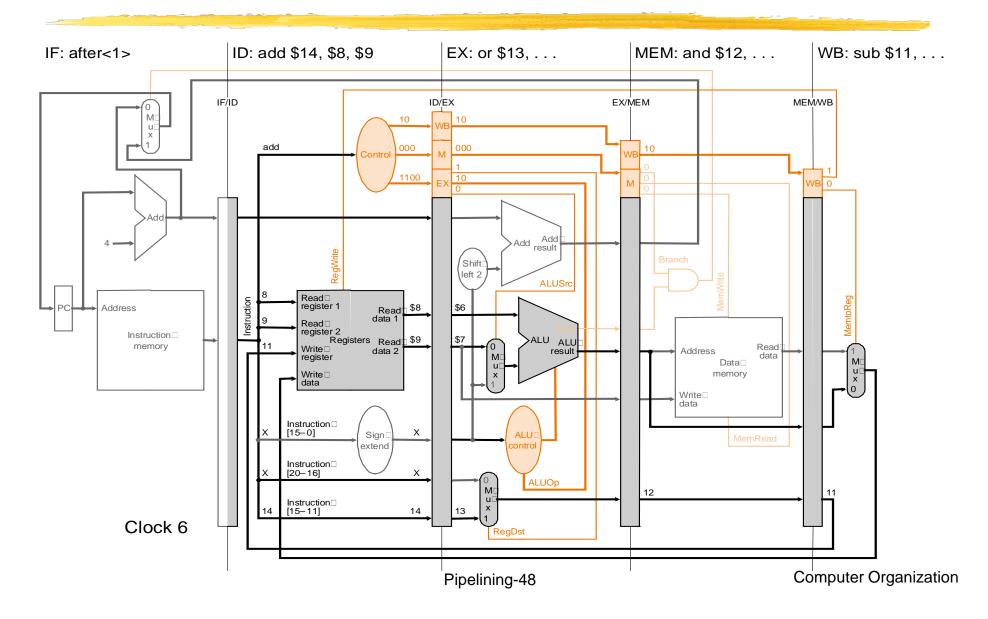


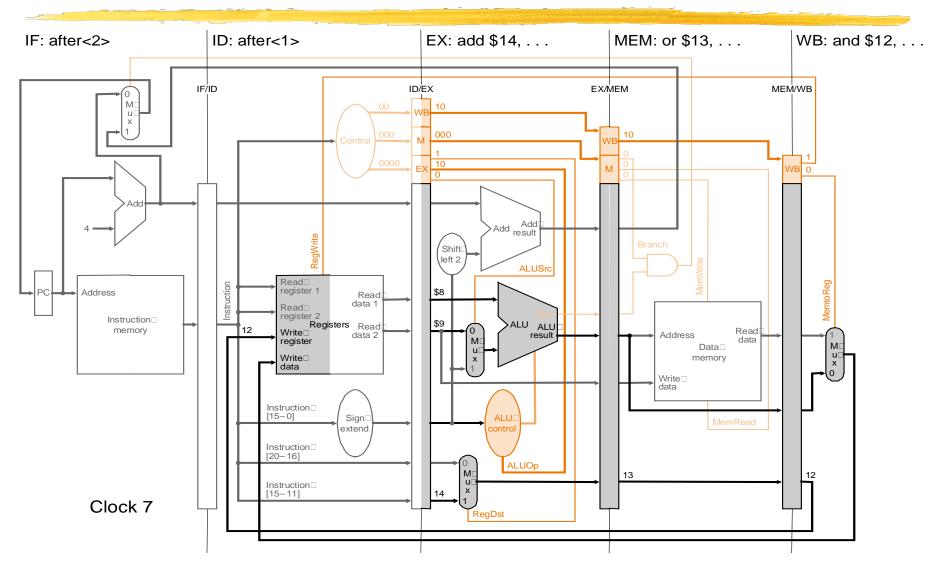


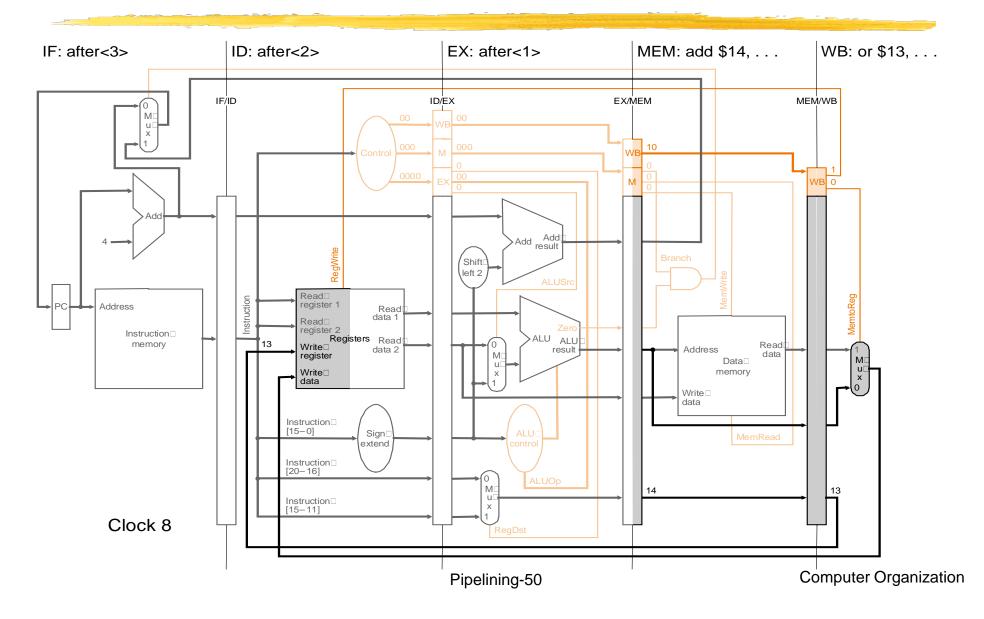


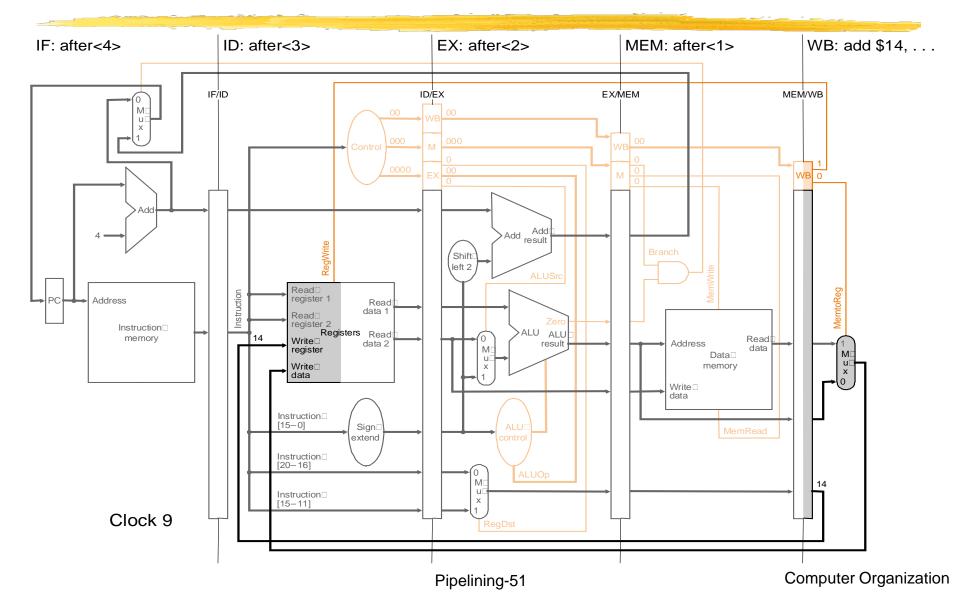












Summary of Pipeline Basics

- Pipelining is a fundamental concept
 - Multiple steps using distinct resources
 - Utilize capabilities of datapath by pipelined instruction processing
 - Start next instruction while working on the current one
 - Limited by length of longest stage (plus fill/flush)
 - Need to detect and resolve hazards
- What makes it easy in MIPS?
 - All instructions are of the same length
 - Just a few instruction formats
 - Memory operands only in loads and stores
- What makes pipelining hard?
 - hazards

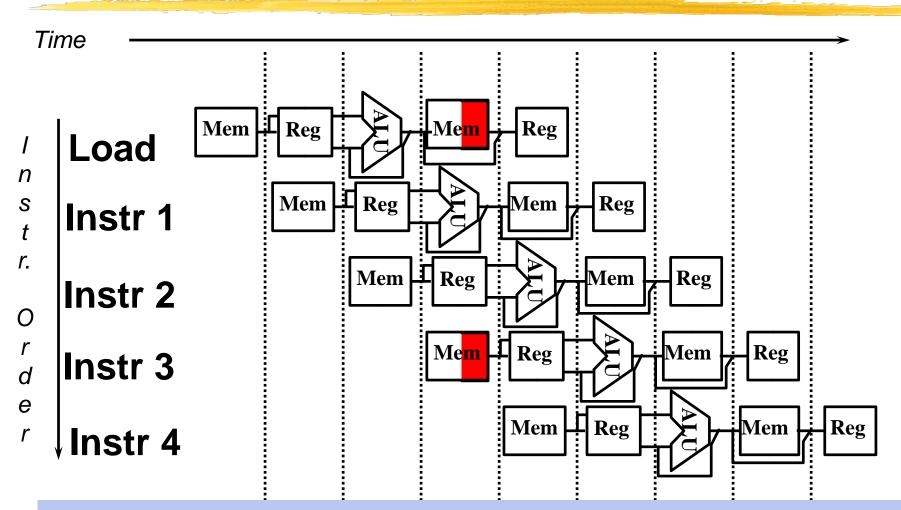
Outline

- An overview of pipelining
- A pipelined datapath
- Pipelined control
- Data hazards and forwarding (R-Type and R-Type)
- Data hazards and stalls (Load and R-type)
- Branch hazards
- Exceptions
- Superscalar and dynamic pipelining

Pipeline Hazards

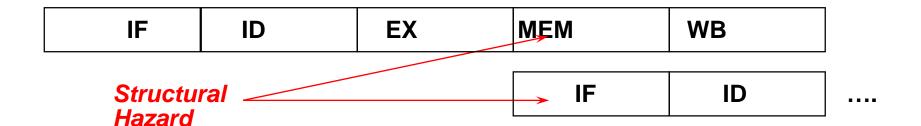
- Pipeline Hazards:
 - Structural hazards: attempt to use the same resource in two different ways at the same time
 - Ex: combined washer/dryer or folder busy doing something else (watching TV)
 - Data hazards: attempt to use item before ready
 - Instruction depends on result of prior instruction still in the pipeline
 - Control hazards: attempt to make decision before condition is evaluated
 - Ex: wash football uniforms and need to see result of previous load to get proper detergent level
 - Branch instructions
- Can always resolve hazards by waiting
 - pipeline control must detect the hazard
 - take action (or delay action) to resolve hazards

Structural Hazard: Single Memory

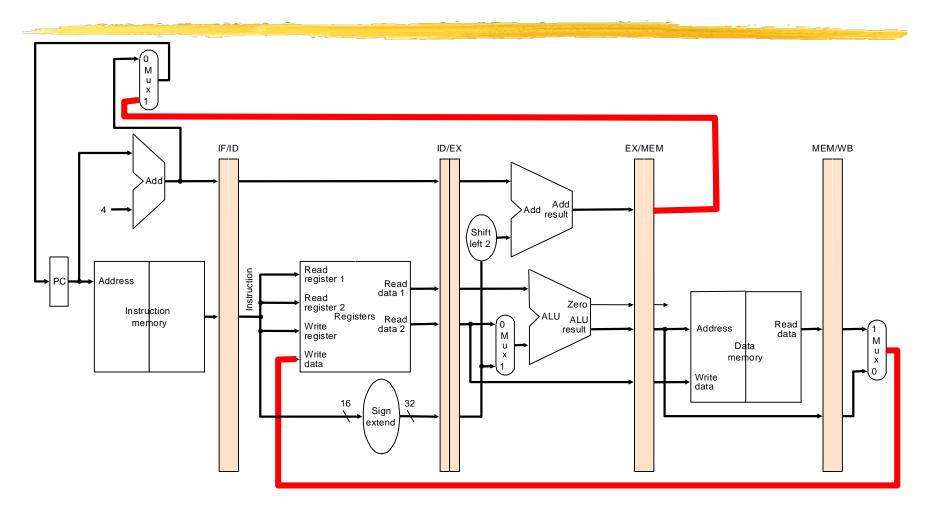


Use 2 memories: data memory and instruction memory

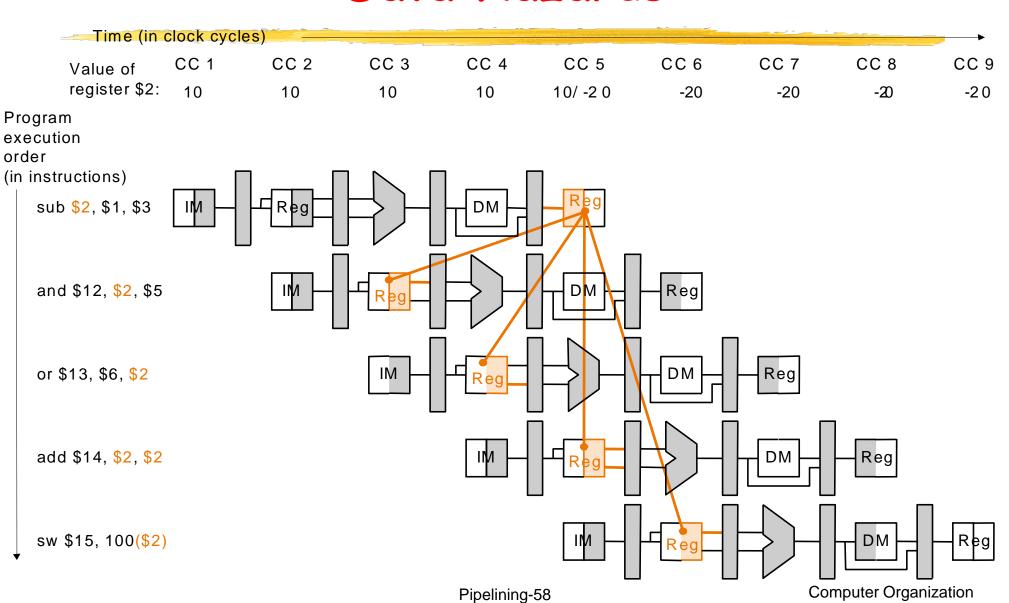
Pipeline Hazards Illustrated



Feedback Path



Data Hazards

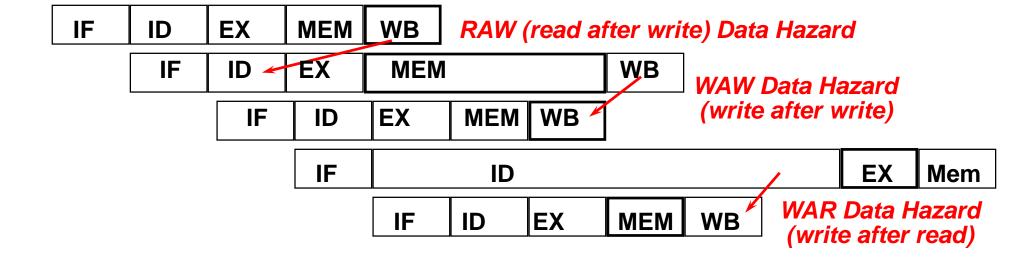


Types of Data Hazards

Three types: (inst. i1 followed by inst. i2)

- RAW (read after write):
 i2 tries to read operand before i1 writes it
- ♦ WAR (write after read):
 - i2 tries to write operand before i1 reads it
 - Gets wrong operand, ex: autoincrement addr.
 - Won't happen in MIPS 5-stage pipeline because:
 - All instructions take 5 stages, and reads are always in stage 2, and writes are always in stage 5
- WAW (write after write):
 - i2 tries to write operand before i1 writes it
 - Leaves wrong result (i1's not i2's); occur only in pipelines that write in more than one stage
 - Won't happen in MIPS 5-stage pipeline because:
 - All instructions take 5 stages, and writes are always in stage 5

Pipeline Hazards Illustrated



Handling Data Hazards

- Use simple, fixed designs
 - Eliminate WAR by always fetching operands early (ID) in pipeline
 - Eliminate WAW by doing all write backs in order (last stage, static)
 - These features have a lot to do with ISA design
- Internal forwarding in register file:
 - Write in first half of clock and read in second half
 - Read delivers what is written, resolve hazard between sub and add
- Detect and resolve remaining ones
 - Compiler inserts NOP (software solution)
 - Stall (hardware solution)
 - Forward (hardware solution)

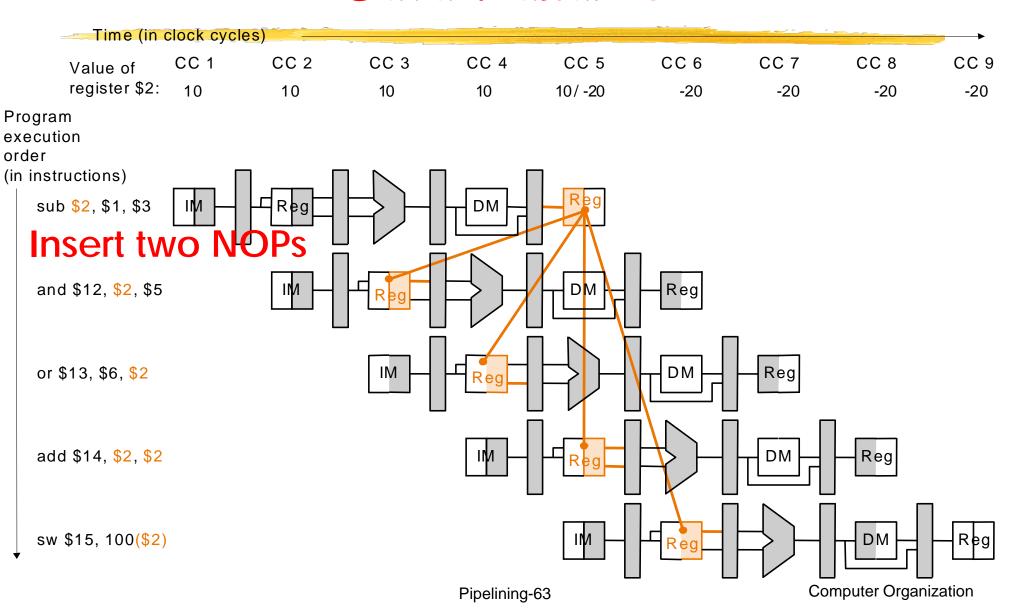
Software Solution

- Have compiler guarantee no hazards
- Where do we insert the NOPs?

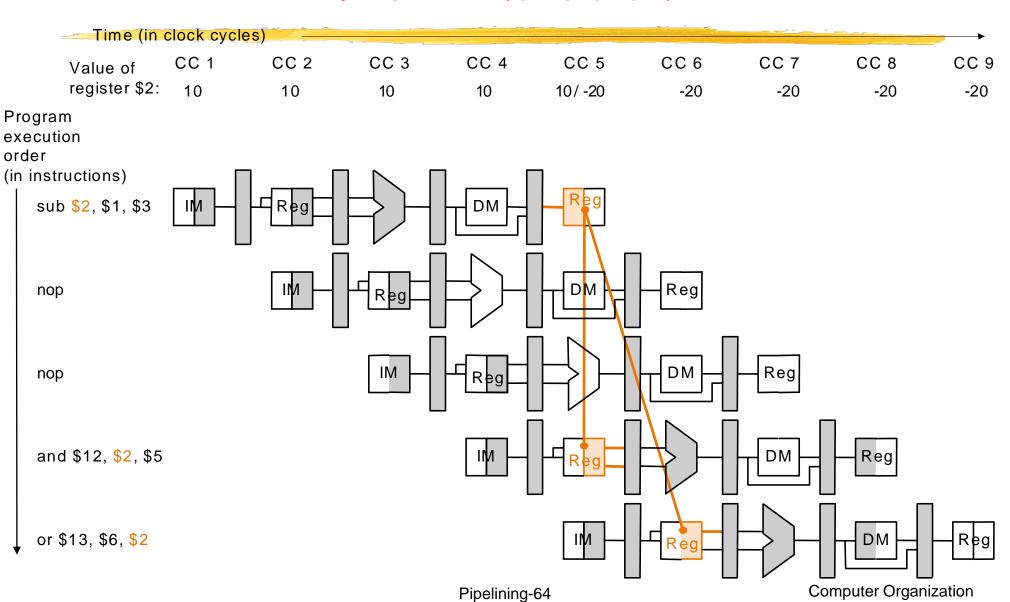
```
sub $2, $1, $3
and $12, $2, $5
or $13, $6, $2
add $14, $2, $2
sw $15, 100($2)
```

Problem: this really slows us down!

Data Hazards



NOP Insertion



Data Hazards: Forwarding

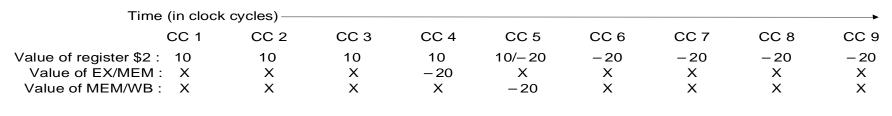
Time (in clock cycles) CC 2 CC₃ CC 4 CC 5 CC₆ CC₁ Value of CC 7 CC8 CC9 register \$2: 10 10/ -20 10 10 10 -20 -20 -20 -20 Program execution order (in instructions) sub \$2, \$1, \$3 DM Reg and \$12, \$2, \$5 Reg DMor \$13, \$6, \$2 Reg add \$14, \$2, \$2 DM sw \$15, 100(\$2) DM

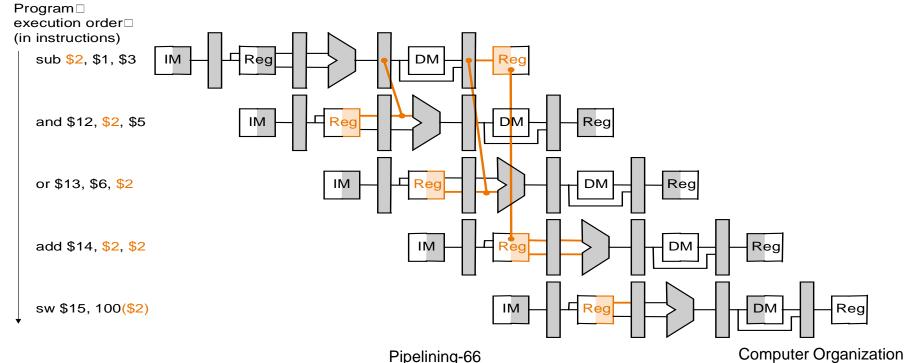
Pipelining-65

Computer Organization

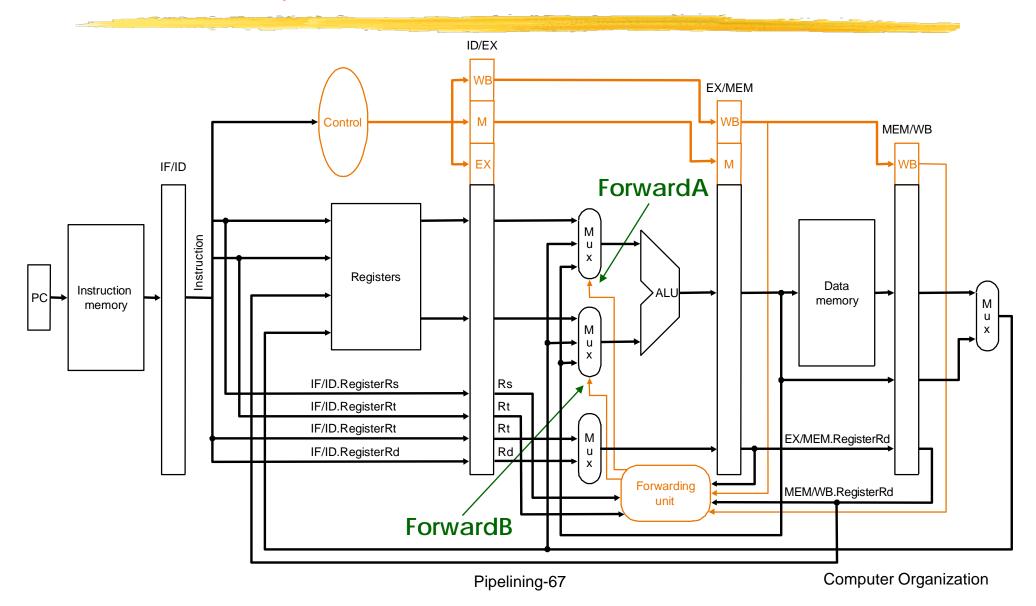
Resolving Hazards: Forwarding

 Use temporary results, ex: those in pipeline registers, don't wait for them to be written into register file





Pipeline with Forwarding



Detecting Data Hazards

- Hazard conditions:
 - 1a. EX/MEM.RegisterRd = ID/EX.RegisterRs
 - 1b. EX/MEM.RegisterRd = ID/EX.RegisterRt
 - 2a. MEM/WB.RegisterRd = ID/EX.RegisterRs
 - 2b. MEM/WB.RegisterRd = ID/EX.RegisterRt
- Two more conditions:

 - Don't forward if destination register is \$0
 → check if RegisterRd = 0

Detecting Data Hazards (cont.)

- Hazard conditions using control signals:
 - At EX stage: EX/MEM.RegWrite and (EX/MEM.RegRd≠0) and (EX/MEM.RegRd=ID/EX.RegRs)
 - At MEM stage: MEM/WB.RegWrite and (MEM/WB.RegRd≠0) and (MEM/WB.RegRd=ID/EX.RegRs)
 - (replace ID/EX.RegRt for ID/EX.RegRs for the other two conditions)

Forwarding Logic

- Forwarding: input to ALU from any pipe reg.
 - Add multiplexors to ALU input
 - Control forwarding in EX → carry Rs in ID/EX
- Control signals for forwarding:
 - If both WB and MEM forward, ex: add \$1,\$1,\$2;
 add \$1,\$1,\$3; add \$1,\$1,\$4; → let MEM forward
 - EX hazard:

MEM hazard:

```
• if (MEM/WB.RegWrite and (MEM/WB.RegRd≠0)

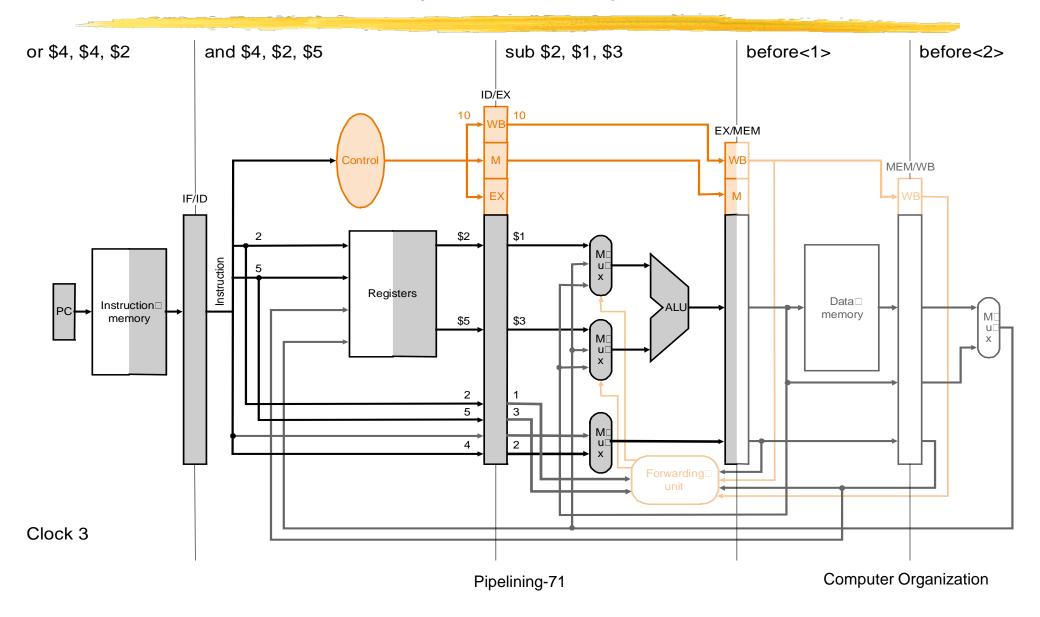
and not(EX/MEM.RegWrite and (EX/MEM.RegRd≠0))

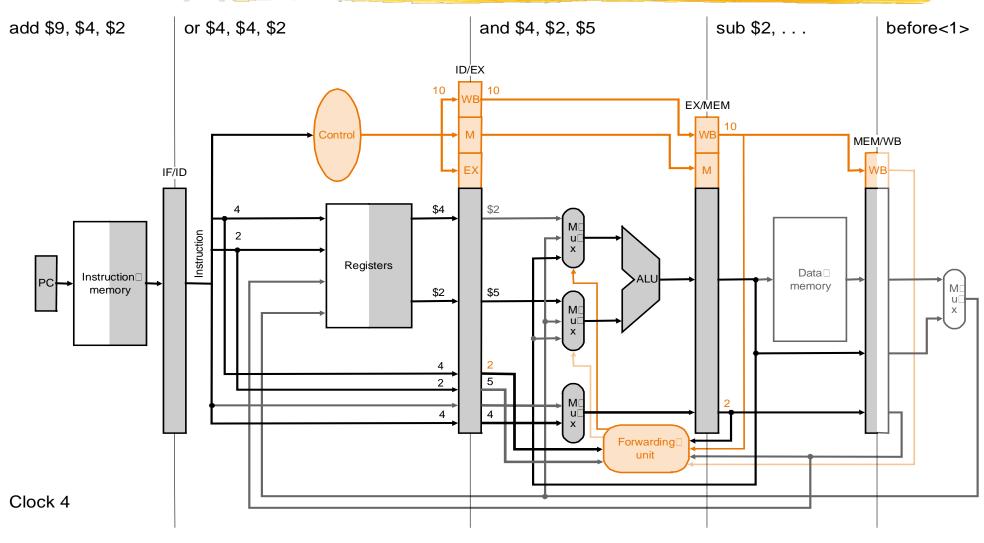
and (EX/MEM.RegRd ≠ ID/EX.RegRs)

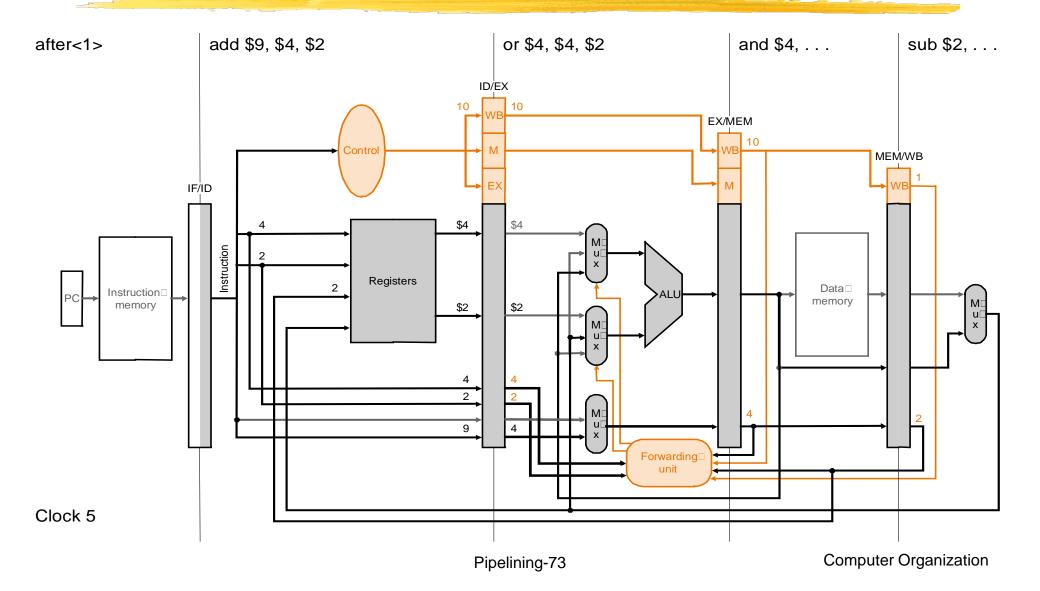
and (MEM/WB.RegRd=ID/EX.RegRs))

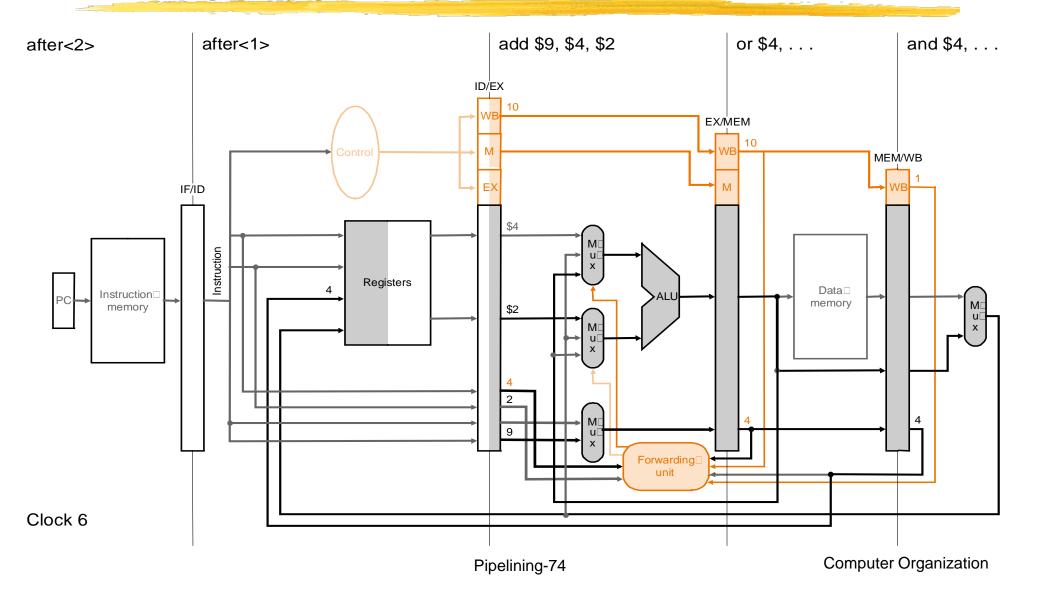
ForwardA=01
```

(ID/EX.RegRt ↔ ID/EX.RegRs, ForwardB ↔ ForwardA)







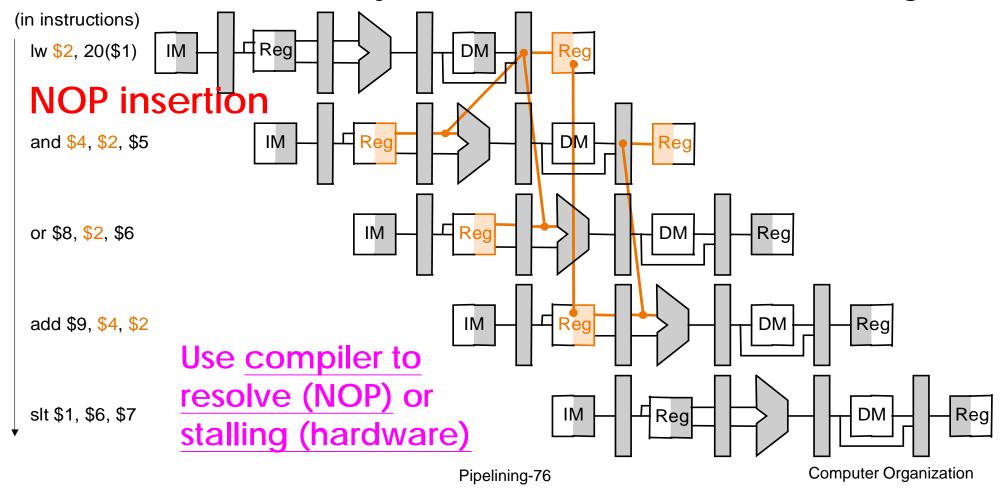


Outline

- An overview of pipelining
- A pipelined datapath
- Pipelined control
- Data hazards and forwarding (R-Type and R-Type)
- Data hazards and stalls (Load and R-type)
- Branch hazards
- Exceptions
- Superscalar and dynamic pipelining

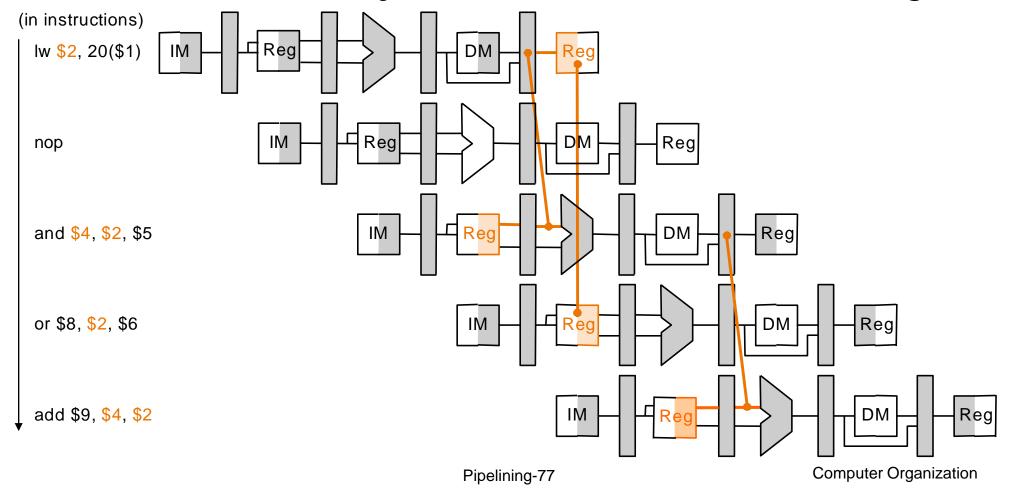
Can't Always Forward

- 1w can still cause a hazard:
 - if it is followed by an instruction to read the loaded reg.



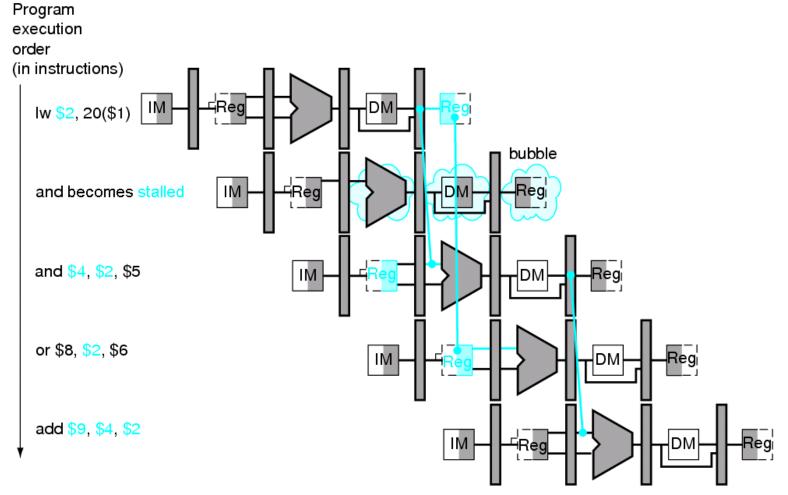
NOP Insertion

- 1w can still cause a hazard:
 - if it is followed by an instruction to read the loaded reg.



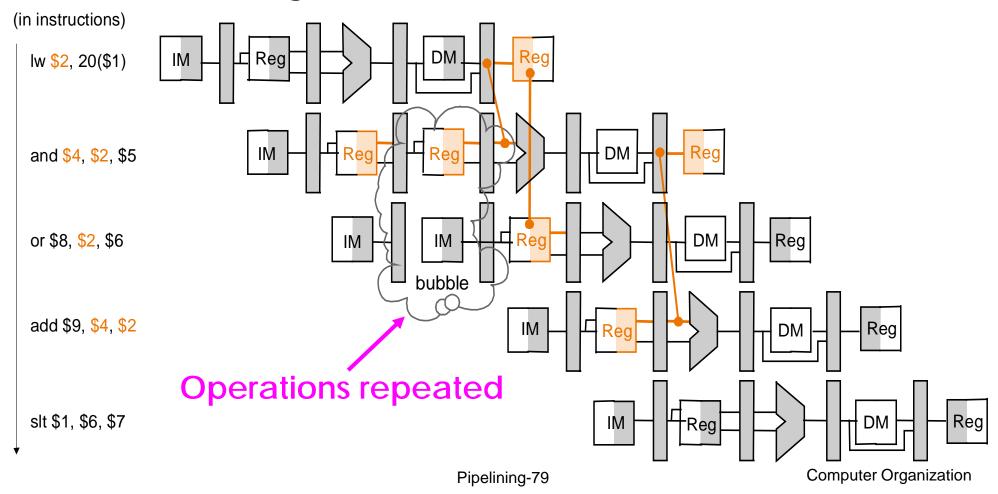
Stalling

The way stalls (bubbles) are inserted into the pipeline



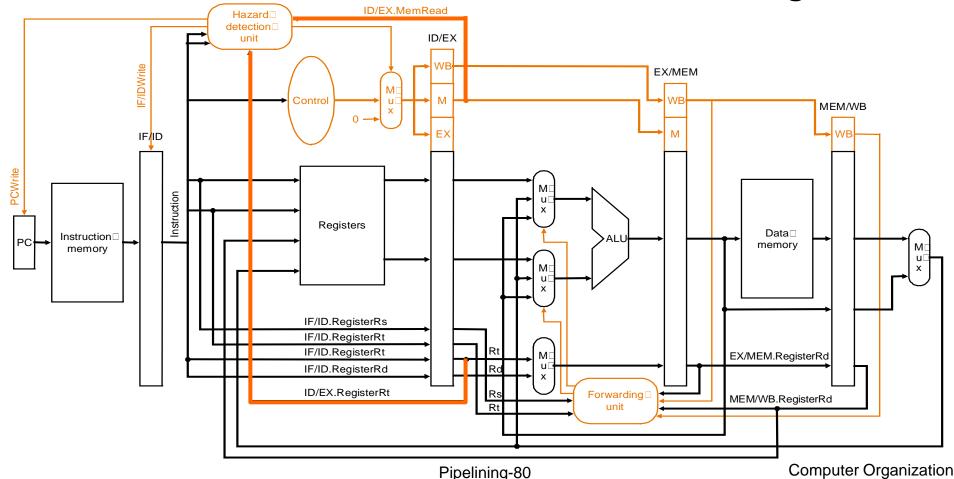
Stalling

 Stall pipeline by keeping instructions in same stage and inserting an NOP instead



Pipeline with Stalling Unit

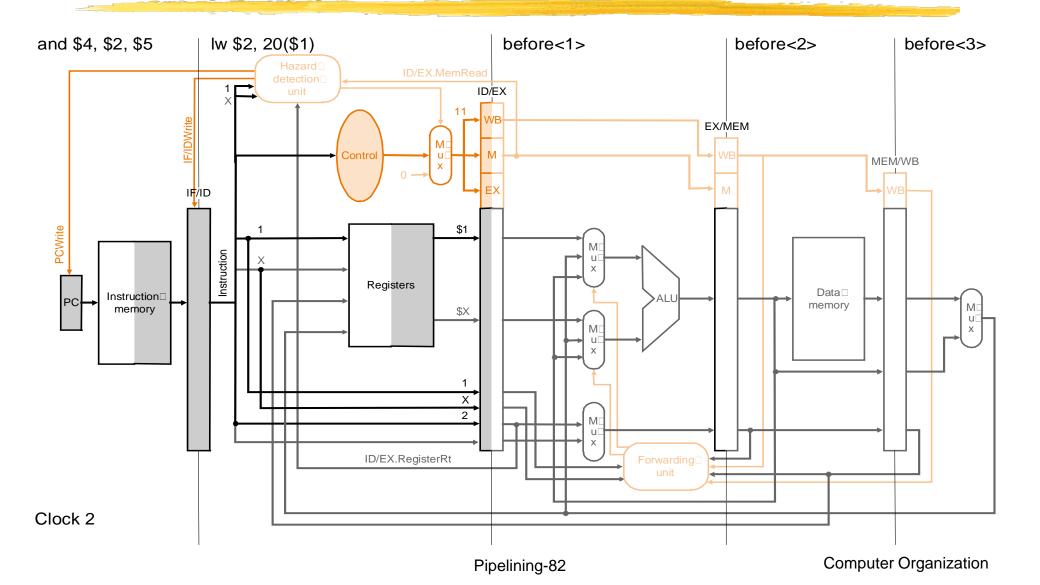
- Forwarding controls ALU inputs
- Hazard detection controls PC, IF/ID, control signals

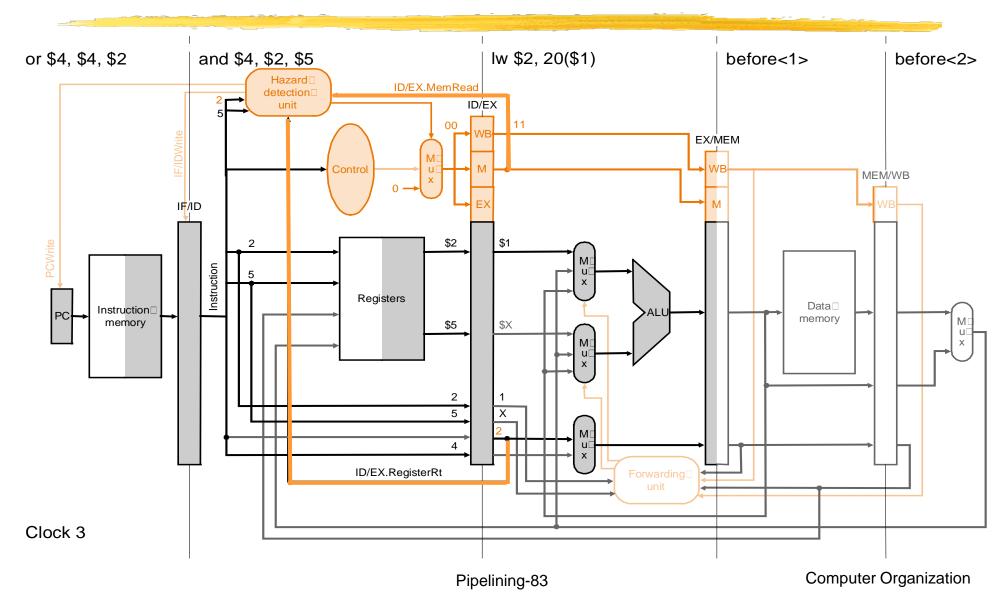


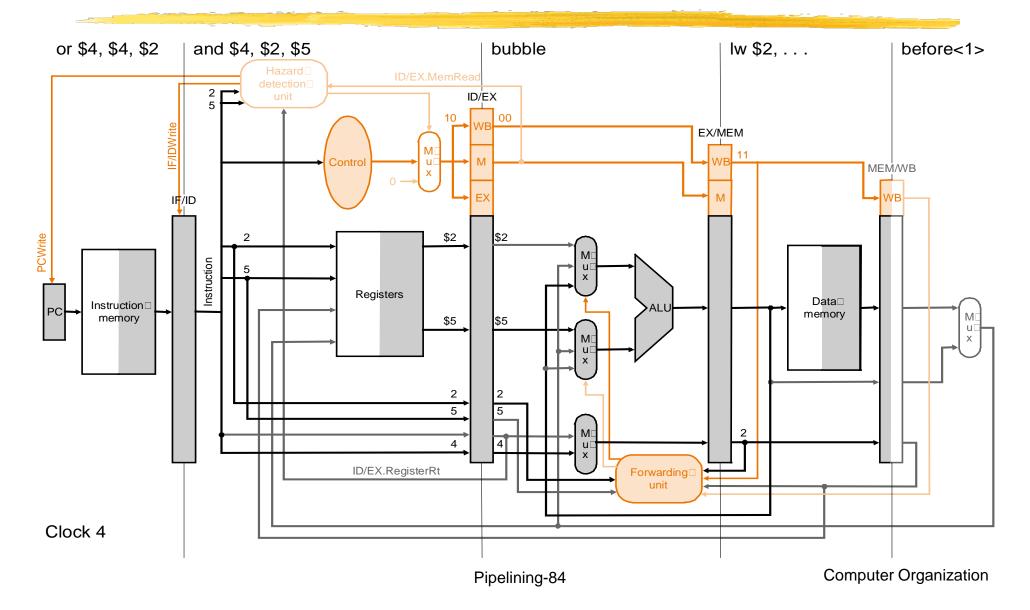
Handling Stalls

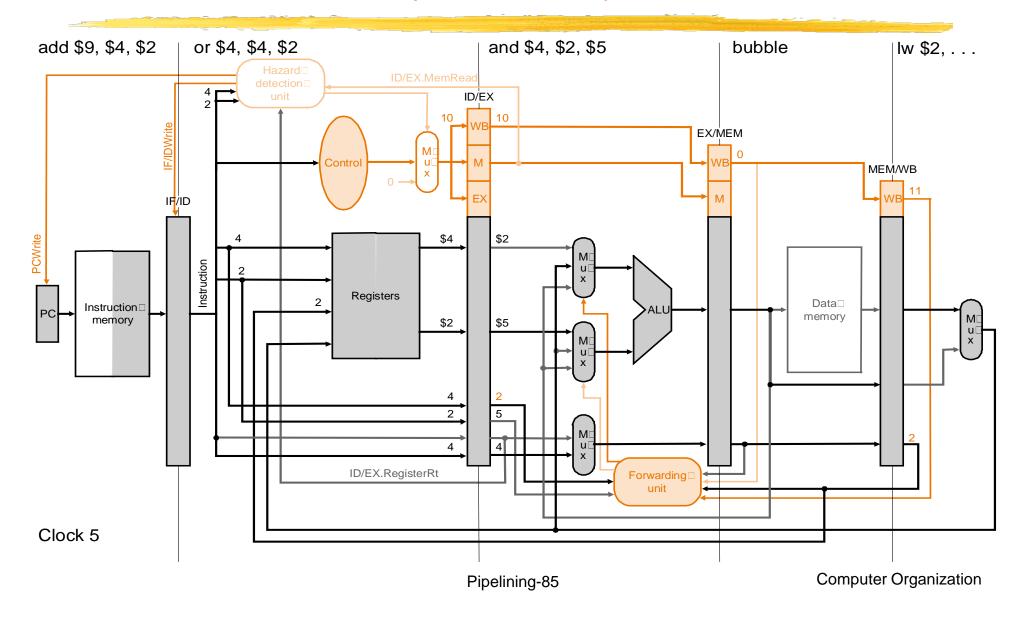
Hazard detection unit in ID to insert stall between a load instruction and its use:

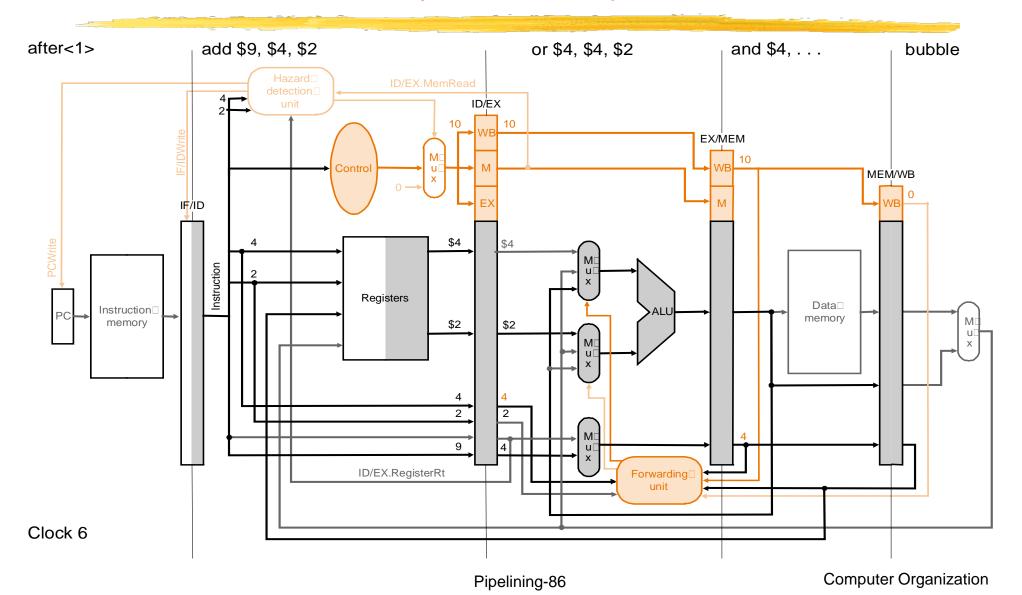
- How to stall?
 - Stall instruction in IF and ID: not change PC and IF/ID
 the stages re-execute the instructions
 - What to move into EX: insert an NOP by changing EX, MEM, WB control fields of ID/EX pipeline register to 0
 - as control signals propagate, all control signals to EX, MEM, WB are deasserted and no registers or memories are written

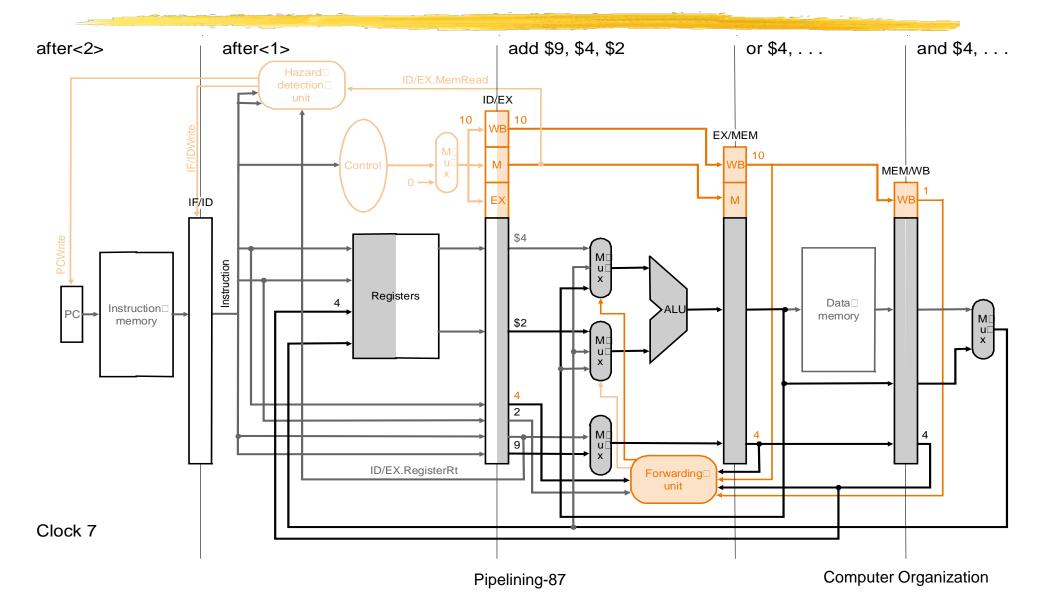








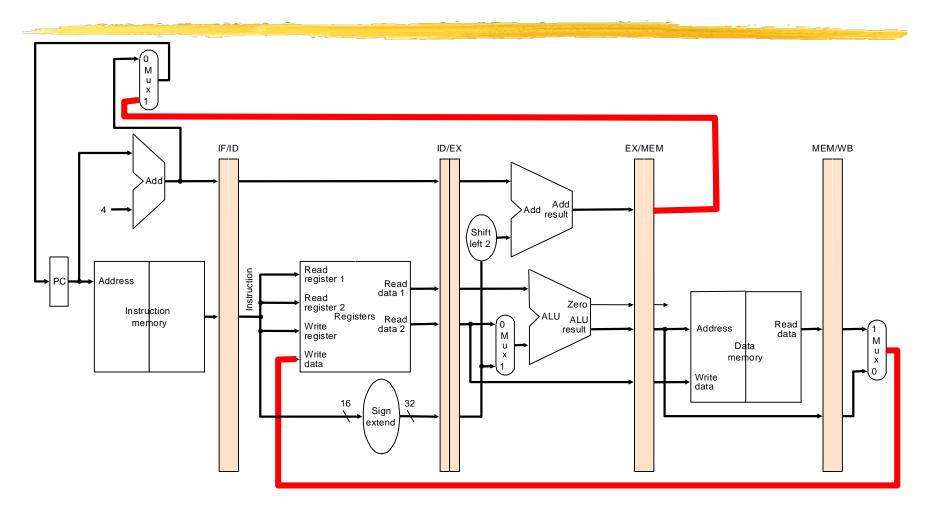




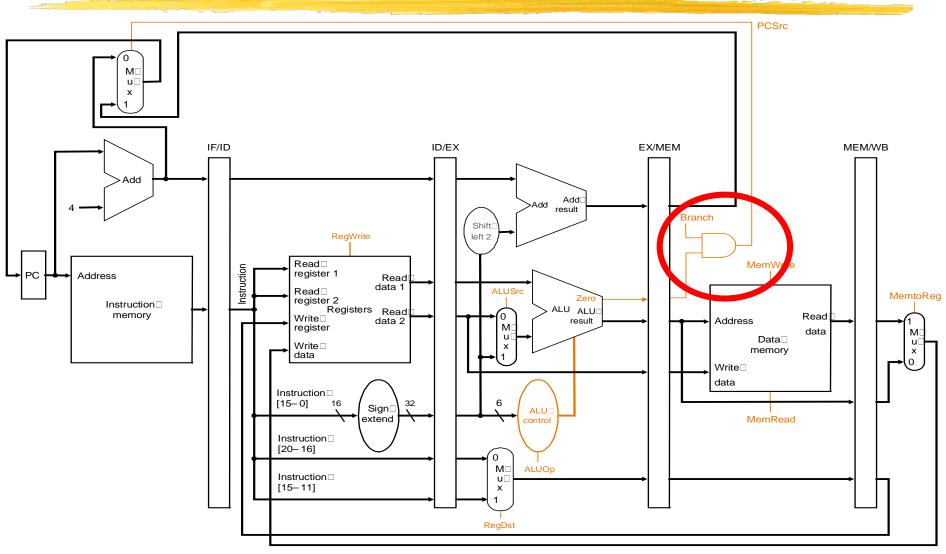
Outline

- An overview of pipelining
- A pipelined datapath
- Pipelined control
- Data hazards and forwarding
- Data hazards and stalls
- Branch hazards
- Exceptions
- Superscalar and dynamic pipelining

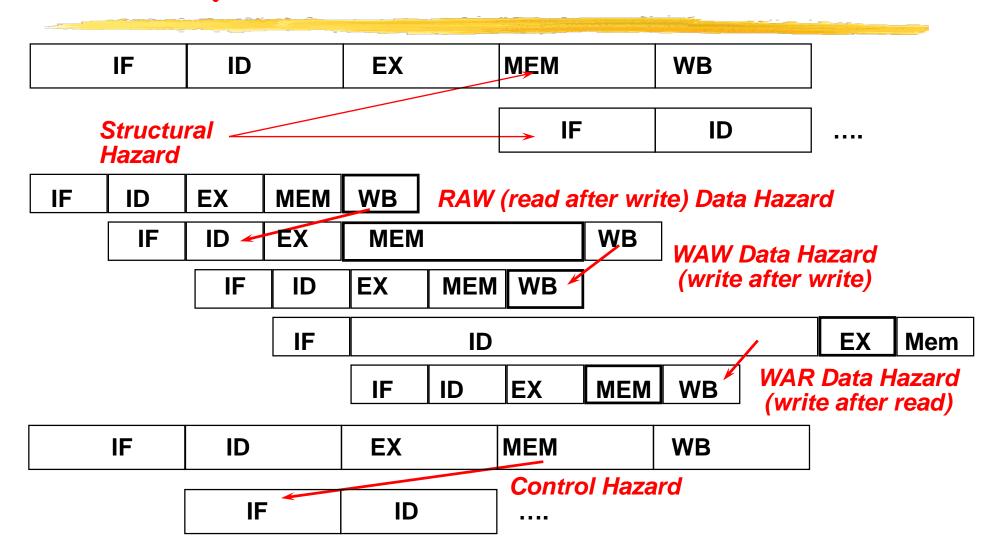
Feedback Path



Pipeline Datapath with Control Signals



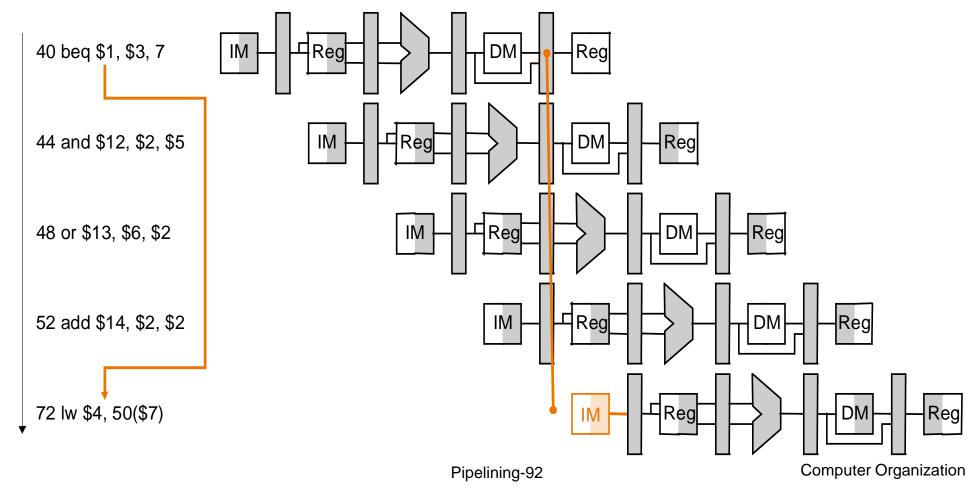
Pipeline Hazards Illustrated



Branch Hazards

When decide to branch, other inst. are in pipeline!

(in instructions)

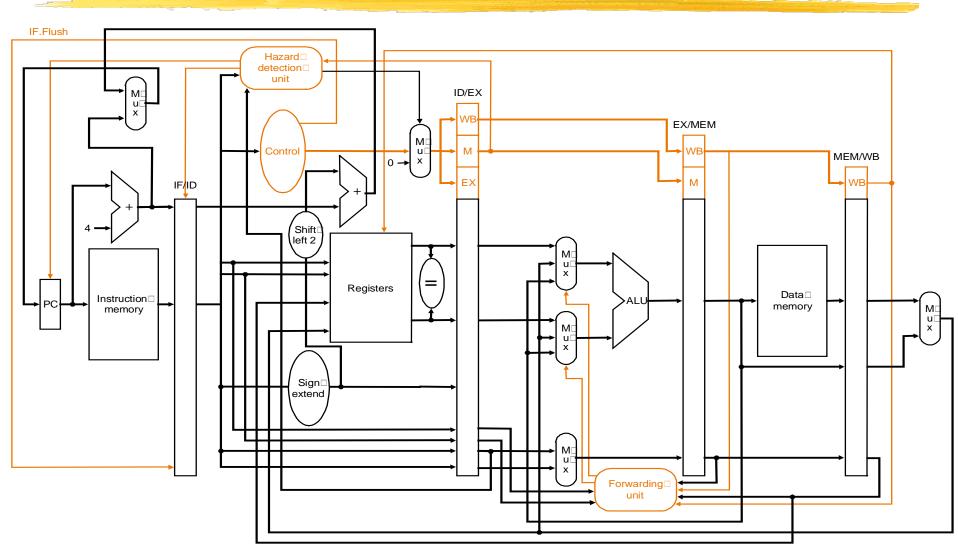


Handling Branch Hazard

- Predict branch always not taken
 - Need to add hardware for flushing inst. if wrong
 - Branch decision made at MEM

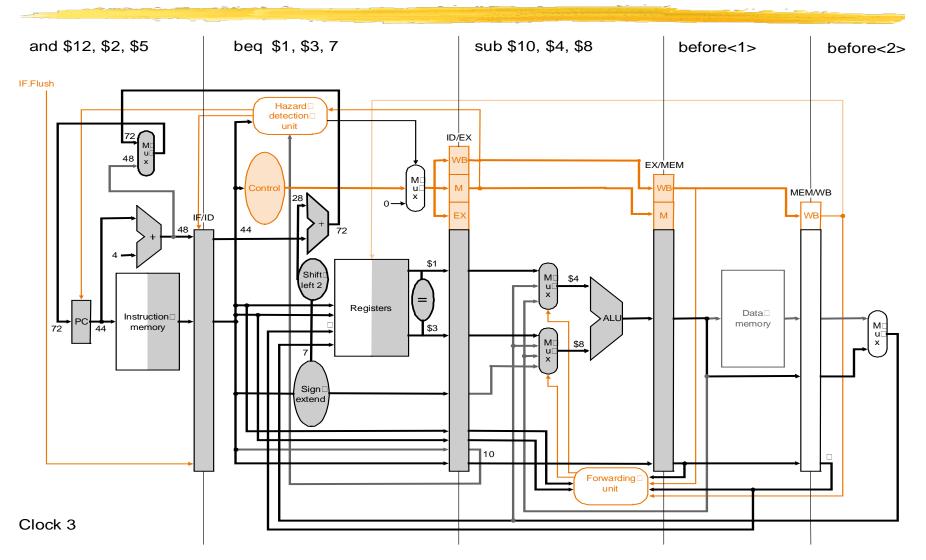
 need to flush inst. in IF/ID, ID/EX, EX/MEM by changing control values to 0
- Reduce delay of taken branch by moving branch execution earlier in the pipeline
 - Move up branch address calculation to ID
 - Check branch equality at ID (using XOR) by comparing the two registers read during ID
 - Branch decision made at ID -> one instruction to flush
 - Add a control signal, IF.Flush, to zero instruction field of IF/ID → making the instruction an NOP
- Dynamic branch prediction
- Compiler rescheduling, delay branch

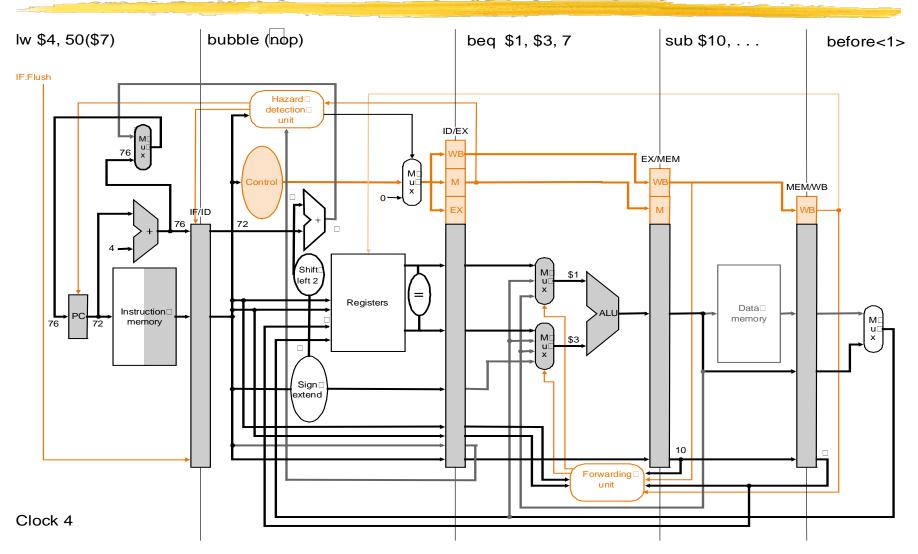
Pipeline with Flushing



Is that all for branch-on-equal? Pipelining-94

Computer Organization



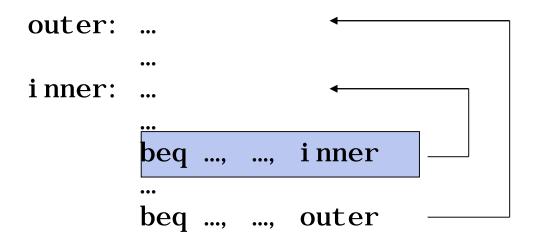


Dynamic Branch Prediction

- In deeper and superscalar pipelines, branch penalty is more significant
- Use dynamic prediction
 - Branch prediction buffer (aka branch history table)
 - Indexed by recent branch instruction addresses
 - Stores outcome (taken/not taken)
 - To execute a branch
 - Check table, expect the same outcome
 - Start fetching from fall-through or target
 - If wrong, flush pipeline and flip prediction

1-Bit Predictor: Shortcoming

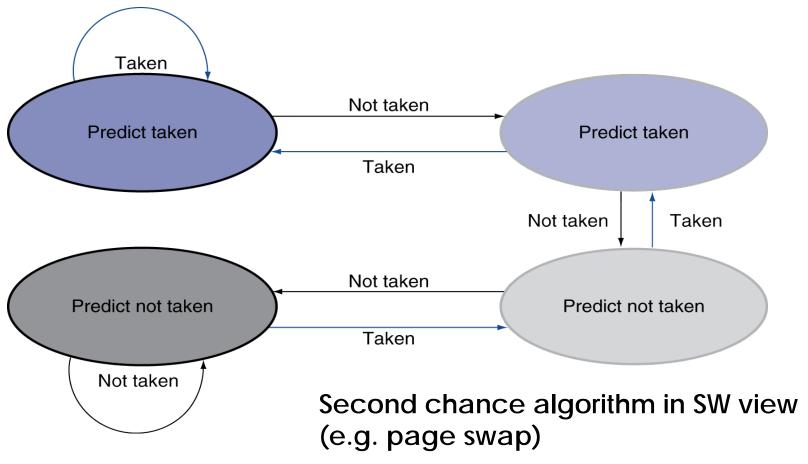
Inner loop branches mispredicted twice!



- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of inner loop next time around

2-Bit Predictor

 Only change prediction on two successive mispredictions

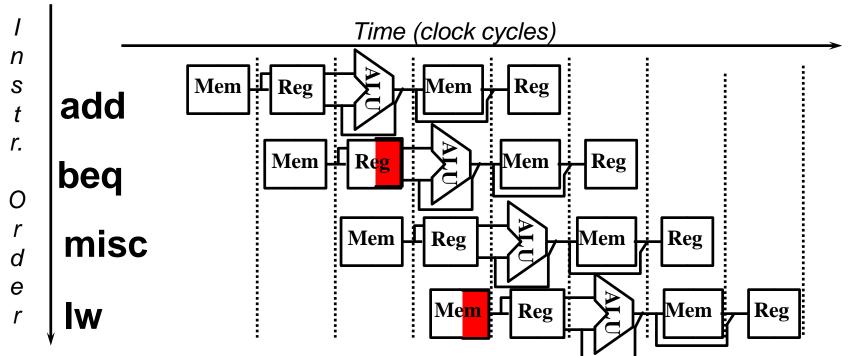


Calculating the Branch Target

- Even with predictor, still need to calculate the target address
 - 1-cycle penalty for a taken branch
- Branch target buffer
 - Cache of target addresses
 - Indexed by PC when instruction fetched
 - If hit and instruction is branch predicted taken, can fetch target immediately

Delayed Branch

- Predict-not-taken + branch decision at ID
 - → the following instruction is always executed
 - → branches take effect 1 cycle later



 0 clock cycle penalty per branch instruction if we can find an instruction to put in slot (≅50% of time)

Outline

- An overview of pipelining
- A pipelined datapath
- Pipelined control
- Data hazards and forwarding
- Data hazards and stalls
- Branch hazards
- Exceptions
- Superscalar and dynamic pipelining

Exceptions and Interrupts

- "Unexpected" events requiring change in flow of control
 - Different ISAs use the terms differently
- Exception
 - Arises within the CPU
 - Ex: undefined opcode, overflow, syscall, ...
- Interrupt
 - From an external I/O controller
- Dealing with them without sacrificing performance is hard

Handling Exceptions

- In MIPS, exceptions managed by a System Control Coprocessor (CP0)
- Save PC of offending (or interrupted) instruction
 - In MIPS: Exception Program Counter (EPC)
- Save indication of the problem
 - In MIPS: Cause register
 - We'll assume 1-bit
 - 0 for undefined opcode, 1 for overflow
- Jump to handler at 8000 00180

An Alternate Mechanism

- Vectored Interrupts
 - Handler address determined by the cause
- Example:

Undefined opcode: C000 0000

Overflow: C000 0020

• ...: C000 0040

- Instructions either
 - Deal with the interrupt, or
 - Jump to real handler

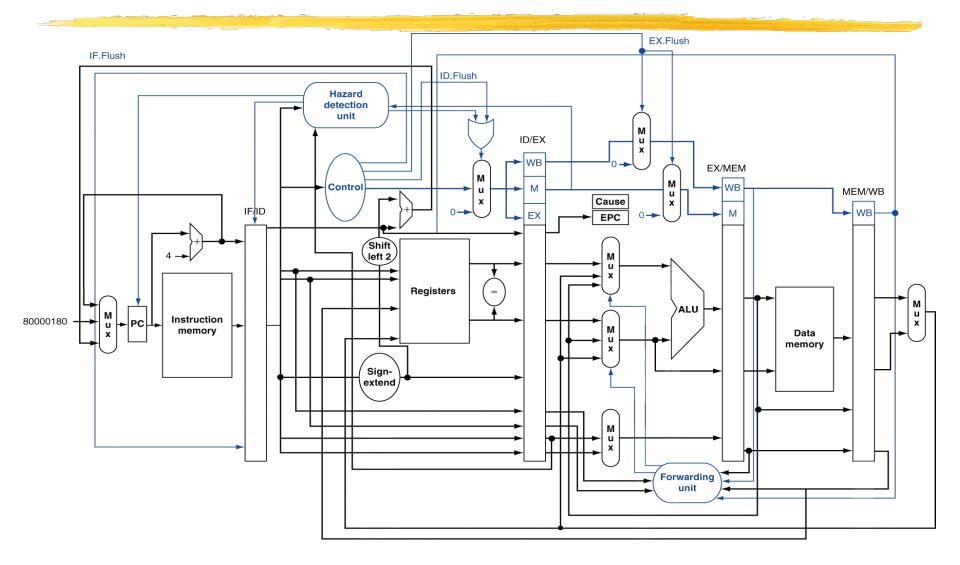
Handler Actions

- Read cause, and transfer to relevant handler
- Determine action required
- If restartable
 - Take corrective action
 - use EPC to return to program
- Otherwise
 - Terminate program
 - Report error using EPC, cause, ...

Exceptions in a Pipeline

- Another form of control hazard
- Consider overflow on add in EX stage
 add \$1, \$2, \$1
 - Prevent \$1 from being clobbered
 - Complete previous instructions
 - Flush add and subsequent instructions
 - Set Cause and EPC register values
 - Transfer control to handler
- Similar to mispredicted branch
 - Use much of the same hardware

Pipeline with Exceptions



Exception Properties

- Restartable exceptions
 - Pipeline can flush the instruction
 - Handler executes, then returns to the instruction
 - Refetched and executed from scratch
- PC saved in EPC register
 - Identifies causing instruction
 - Actually PC + 4 is saved
 - Handler must adjust

Exception Example

Exception on add in

```
40 sub $11, $2, $4
44 and $12, $2, $5
48 or $13, $2, $6
4C add $1, $2, $1
50 slt $15, $6, $7
54 lw $16, 50($7)
```

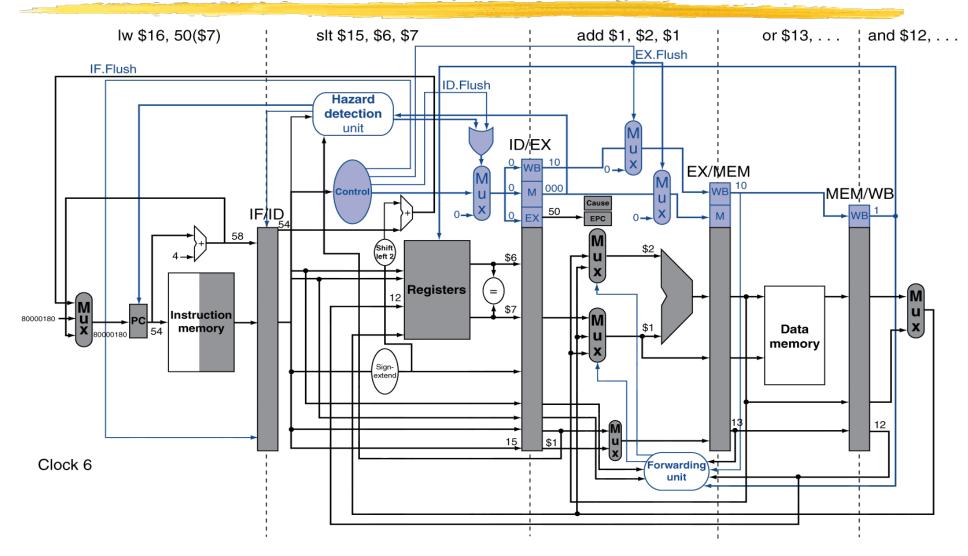
•••

Handler

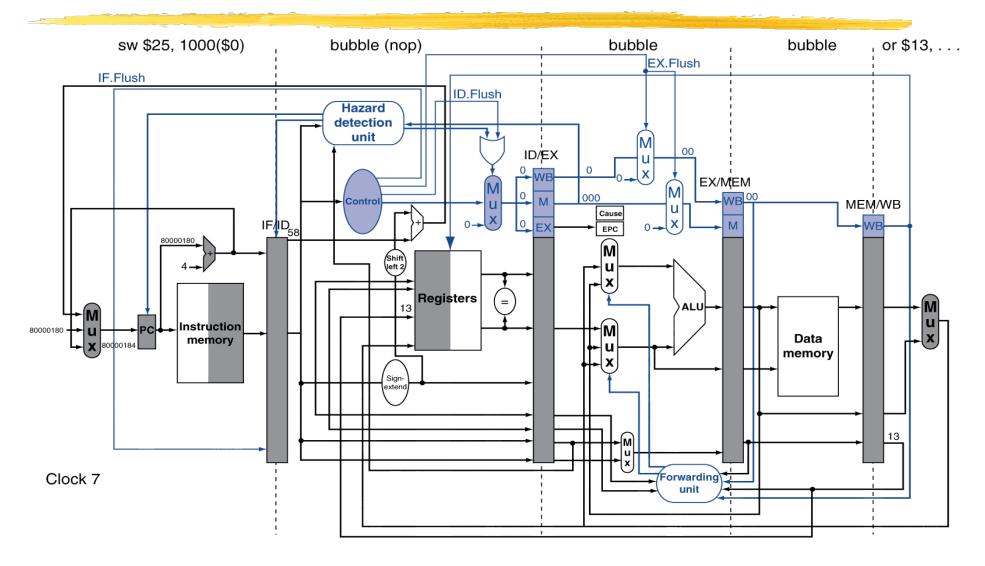
```
80000180 sw $25, 1000($0)
80000184 sw $26, 1004($0)
```

•••

Exception Example



Exception Example



Multiple Exceptions

- Pipelining overlaps multiple instructions
 - Could have multiple exceptions at once
- Simple approach: deal with exception from earliest instruction
 - Flush subsequent instructions
 - "Precise" exceptions
- In complex pipelines
 - Multiple instructions issued per cycle
 - Out-of-order completion
 - Maintaining precise exceptions is difficult!

Imprecise Exceptions

- Just stop pipeline and save state
 - Including exception cause(s)
- Let the handler work out
 - Which instruction(s) had exceptions
 - Which to complete or flush
 - May require "manual" completion
- Simplifies hardware, but more complex handler software
- Not feasible for complex multiple-issue out-of-order pipelines

Outline

- An overview of pipelining
- A pipelined datapath
- Pipelined control
- Data hazards and forwarding
- Data hazards and stalls
- Branch hazards
- Exceptions
- Superscalar and dynamic pipelining

Instruction-Level Parallelism (ILP)

- Pipelining: executing multiple instructions in parallel
- To increase ILP
 - Deeper pipeline
 - Less work per stage ⇒ shorter clock cycle
 - Multiple issue
 - Replicate pipeline stages ⇒ multiple pipelines
 - Start multiple instructions per clock cycle
 - CPI < 1, so use Instructions Per Cycle (IPC)</p>
 - Ex: 4GHz 4-way multiple-issue
 - 16 BIPS, peak CPI = 0.25, peak IPC = 4
 - But dependencies reduce this in practice

Multiple Issue

- Static multiple issue
 - Compiler groups instructions to be issued together
 - Packages them into "issue slots"
 - Compiler detects and avoids hazards
- Dynamic multiple issue
 - CPU examines instruction stream and chooses instructions to issue each cycle
 - Compiler can help by reordering instructions
 - CPU resolves hazards using advanced techniques at runtime

Speculation

- "Guess" what to do with an instruction
 - Start operation as soon as possible
 - Check whether guess was right
 - If so, complete the operation
 - If not, roll-back and do the right thing
- Common to static and dynamic multiple issue
- Examples
 - Speculate on branch outcome
 - Roll back if path taken is different
 - Speculate on load
 - Roll back if location is updated

Compiler/Hardware Speculation

- Compiler can reorder instructions
 - Ex: move load before branch
 - Can include "fix-up" instructions to recover from incorrect guess
- Hardware can look ahead for instructions to execute
 - Buffer results until it determines they are actually needed
 - Flush buffers on incorrect speculation

Speculation and Exceptions

- What if exception occurs on a speculatively executed instruction?
 - Ex: speculative load before null-pointer check
- Static speculation
 - Can add ISA support for deferring exceptions
- Dynamic speculation
 - Can buffer exceptions until instruction completion (which may not occur)

Static Multiple Issue

- Compiler groups instructions into "issue packets"
 - Group of instructions that can be issued on a single cycle
 - Determined by pipeline resources required
- Think of an issue packet as a very long instruction
 - Specifies multiple concurrent operations
 - ⇒ Very Long Instruction Word (VLIW)

Scheduling Static Multiple Issue

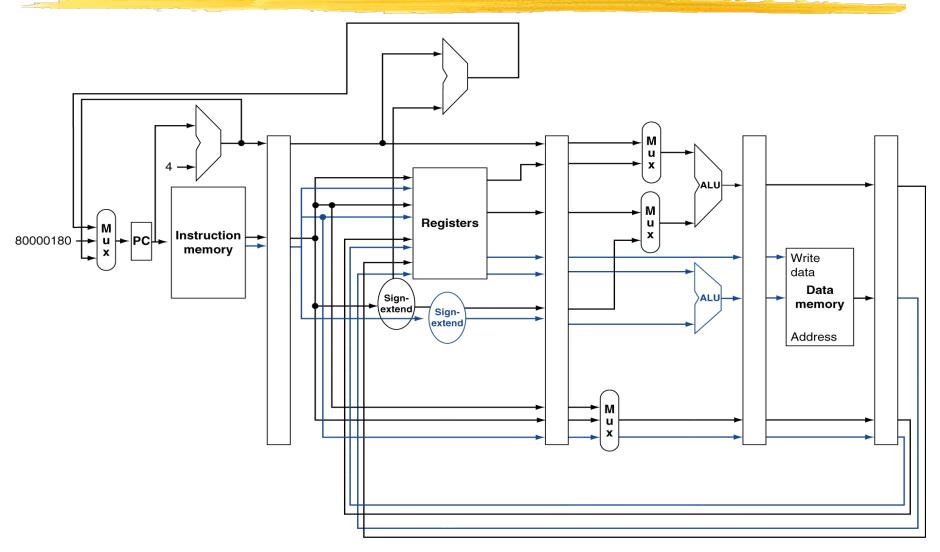
- Compiler must remove some/all hazards
 - Reorder instructions into issue packets
 - No dependencies within a packet
 - Possibly some dependencies between packets
 - Varies between ISAs; compiler must know!
 - Pad with nop if necessary

MIPS with Static Dual Issue

- Two-issue packets
 - One ALU/branch instruction
 - One load/store instruction
 - 64-bit aligned
 - ALU/branch, then load/store
 - Pad an unused instruction with nop

Address	Instruction type	Pipeline Stages						
n	ALU/branch	IF	ID	EX	MEM	WB		
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	MEM	WB	
n + 12	Load/store		IF	ID	EX	MEM	WB	
n + 16	ALU/branch			IF	ID	EX	MEM	WB
n + 20	Load/store			IF	ID	EX	MEM	WB

MIPS with Static Dual Issue



Hazards in the Dual-Issue MIPS

- More instructions executing in parallel
- EX data hazard
 - Forwarding avoided stalls with single-issue
 - Now can't use ALU result in load/store in same packet

```
add $t0, $s0, $s1
load $s2, 0($t0)
```

- Split into two packets, effectively a stall
- Load-use hazard
 - Still one cycle use latency, but now two instructions
- More aggressive scheduling required

Scheduling Example

Schedule this for dual-issue MIPS

```
Loop: lw $t0, 0($s1) # $t0=array element addu $t0, $t0, $s2 # add scalar in $s2 sw $t0, 0($s1) # store result addi $s1, $s1,-4 # decrement pointer bne $s1, $zero, Loop # branch $s1!=0
```

	ALU/branch	Load/store	cycle
Loop:	nop	lw \$t0, 0(\$s1)	1
	addi \$s1, \$s1,-4	nop	2
	addu \$t0, \$t0, \$s2	nop	3
	bne \$s1, \$zero, Loop	sw \$t0, 4(\$s1)	4

IPC = 5/4 = 1.25 (c.f. peak IPC = 2)

Loop Unrolling

- Replicate loop body to expose more parallelism
 - Reduces loop-control overhead
- Use different registers per replication
 - Called "register renaming"
 - Avoid loop-carried "anti-dependencies"
 - Store followed by a load of the same register
 - Aka "name dependence"
 - Reuse of a register name

Loop Unrolling Example

	ALU/branch			Load	cycle			
Loop:	addi	\$s1,	\$s1,-	-16	lw	\$t0,	0(\$s1)	1
	nop				lw	\$t1,	12(\$s1)	2
	addu	\$t0,	\$t0,	\$s2	lw	\$t2,	8(\$s1)	3
	addu	\$t1,	\$t1,	\$s2	lw	\$t3,	4(\$s1)	4
	addu	\$t2,	\$t2,	\$s2	sw	\$t0,	16(\$s1)	5
	addu	\$t3,	\$t4,	\$s2	sw	\$t1,	12(\$s1)	6
	nop				sw	\$t2,	8(\$s1)	7
	bne	\$s1,	\$zero	, Loop	sw	\$t3,	4(\$s1)	8

- ◆ IPC = 14/8 = 1.75
 - Closer to 2, but at cost of registers and code size

Dynamic Multiple Issue

- "Superscalar" processors
- CPU decides whether to issue 0, 1, 2, ... each cycle
 - Avoiding structural and data hazards
- Avoids the need for compiler scheduling
 - Though it may still help
 - Code semantics ensured by the CPU

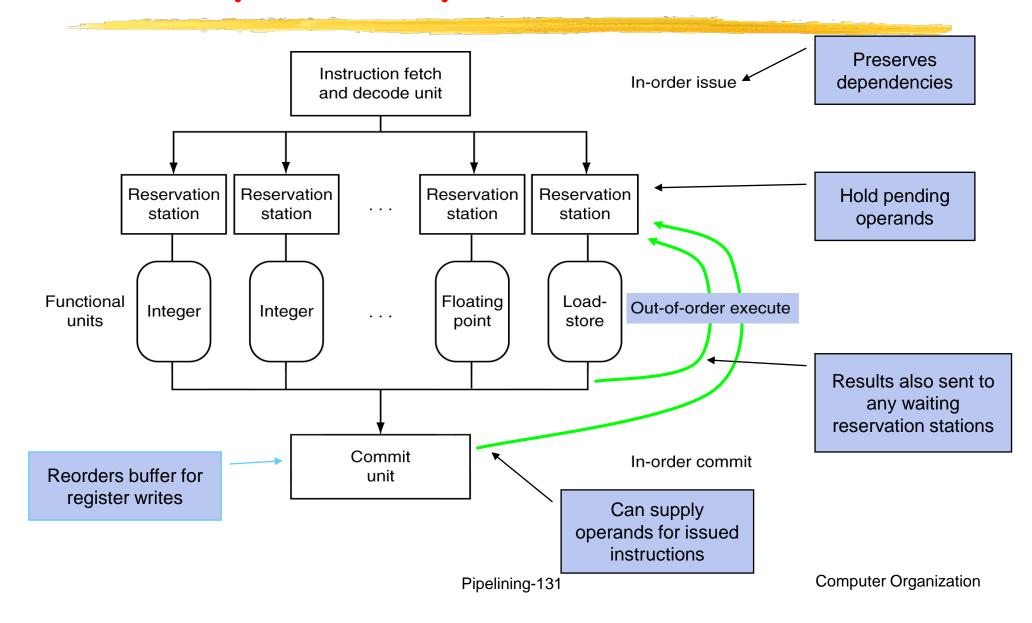
Dynamic Pipeline Scheduling

- Allow the CPU to execute instructions out of order to avoid stalls
 - But commit result to registers in order
- Example

```
lw $t0, 20($s2)
addu $t1, $t0, $t2
sub $s4, $s4, $t3
slti $t5, $s4, 20
```

Can start sub while addu is waiting for 1w

Dynamically Scheduled CPU



Register Renaming

- Reservation stations and reorder buffer effectively provide register renaming
- On instruction issue to reservation station
 - If operand is available in register file or reorder buffer
 - Copied to reservation station
 - No longer required in the register; can be overwritten
 - If operand is not yet available
 - It will be provided to the reservation station by a function unit
 - Register update may not be required

Speculation

- Predict branch and continue issuing
 - Don't commit until branch outcome determined
- Load speculation
 - Avoid load and cache miss delay
 - Predict the effective address
 - Predict loaded value
 - Load before completing outstanding stores
 - Bypass stored values to load unit
 - Don't commit load until speculation cleared

Why Do Dynamic Scheduling?

- Why not just let the compiler schedule code?
 - Not all stalls are predicable
 - Ex: cache misses
 - Can't always schedule around branches
 - Branch outcome is dynamically determined
 - Different implementations of an ISA have different latencies and hazards

Different Pipelined Designs

□Pipeline

□Super-pipeline

- Issue one instruction per (fast) cycle
- ALU takes multiple cycles

 IF
 D
 Ex
 M
 W

 IF
 D
 Ex
 M
 W

 IF
 D
 Ex
 M
 W

 IF
 D
 Ex
 M
 W

Limitation

Issue rate, FU stalls, FU depth

IF1 | IF2 | ID1 | ID2 | EX1 | EX2 | EX3 | M1 | M2 | W1 | W2

IF1 | IF2 | ID1 | ID2 | EX1 | EX2 | EX3 | M1 | M2 | W1 | W2

Clock skew, FU stalls, FU depth

□Super-scalar

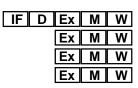
- Issue multiple scalar instructions per cycle



Hazard resolution

□VLIW (EPIC)

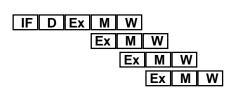
- Each instruction specifies multiple scalar operations
- Compiler determines parallelism



Packing

□Vector operations

- Each instruction specifies series of identical operations



Applicability

Does Multiple Issue Work?

The BIG Picture

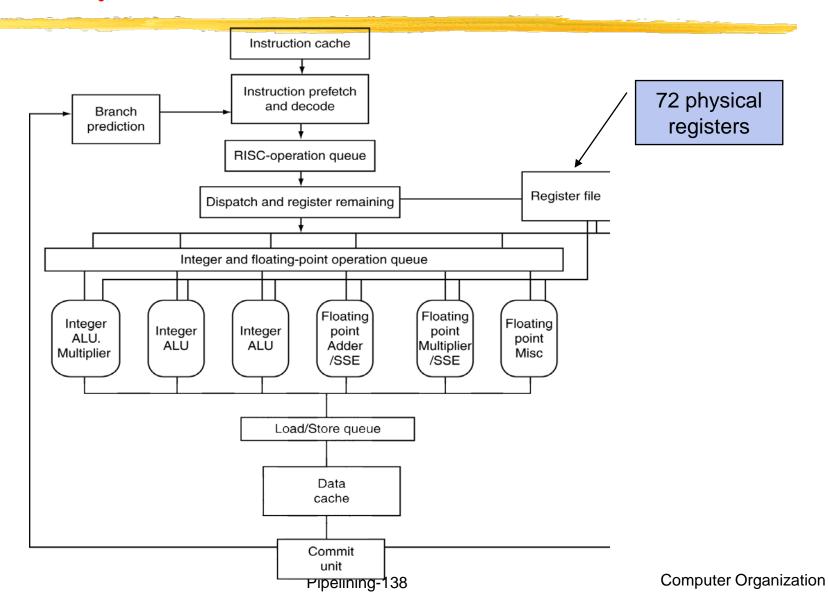
- Yes, but not as much as we'd like
- Programs have real dependencies that limit ILP
- Some dependencies are hard to eliminate
 - Ex: pointer aliasing
- Some parallelism is hard to expose
 - Limited window size during instruction issue
- Memory delays and limited bandwidth
 - Hard to keep pipelines full
- Speculation can help if done well

Power Efficiency

- Complexity of dynamic scheduling and speculations requires power
- Multiple simpler cores may be better

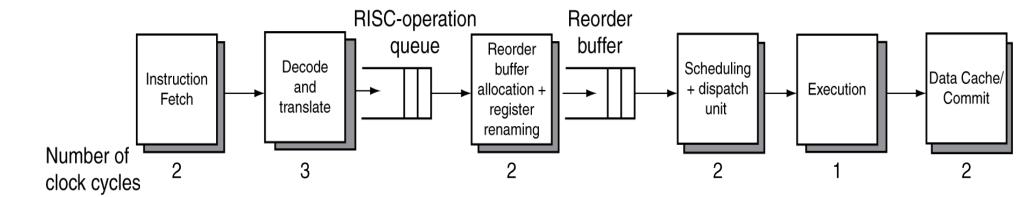
Microprocessor	Year	Clock Rate	Pipeline Stages	Issue width	Out-of- order/ Speculation	Cores	Power
i486	1989	25MHz	5	1	No	1	5W
Pentium	1993	66MHz	5	2	No	1	10W
Pentium Pro	1997	200MHz	10	3	Yes	1	29W
P4 Willamette	2001	2000MHz	22	3	Yes	1	75W
P4 Prescott	2004	3600MHz	31	3	Yes	1	103W
Core	2006	2930MHz	14	4	Yes	2	75W
UltraSparc III	2003	1950MHz	14	4	No	1	90W
UltraSparc T1	2005	1200MHz	6	1	No	8	70W

The Opteron X4 Microarchitecture



The Opteron X4 Pipeline Flow

For integer operations



- FP is 5 stages longer
- Up to 106 RISC-ops in progress
- Bottlenecks
 - Complex instructions with long dependencies
 - Branch mispredictions
 - Memory access delays

Fallacies

- Pipelining is easy (!)
 - The basic idea is easy
 - The devil is in the details
 - Ex: detecting data hazards
- Pipelining is independent of technology
 - So why haven't we always done pipelining?
 - More transistors make more advanced techniques feasible
 - Pipeline-related ISA design needs to take account of technology trends
 - Ex: predicated instructions

Pitfalls

- Poor ISA design can make pipelining harder
 - Ex: complex instruction sets (VAX, IA-32)
 - Significant overhead to make pipelining work
 - IA-32 micro-op approach
 - Ex: complex addressing modes
 - Register update side effects, memory indirection
 - Ex: delayed branches
 - Advanced pipelines have long delay slots

Concluding Remarks

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
 - More instructions completed per second
 - Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
 - Dependencies limit achievable parallelism
 - Complexity leads to the power wall