

Topic II: Resources and Resource Access Control

謝仁偉 教授
jenwei@mail.ntust.edu.tw
國立台灣科技大學 資訊工程系
2017 Fall

Outline

- **Assumptions on Resources and Their Usage**
- Effects of Resource Contention and Resource Access Control
- Nonpreemptive Critical Sections Protocol
- Basic Priority-Inheritance Protocol
- Basic Priority-Ceiling Protocol
- Stack-Based, Priority-Ceiling (Ceiling-Priority) Protocol
- Preemption-Ceiling Protocols
- Controlling Accesses to Multiple-Unit Resources

Assumptions


- The system contains only one processor.
- The system contains ρ types of serially reusable resources named R_1, R_2, \dots, R_ρ .
 - There are v_i indistinguishable units of resource (of type) R_i , for $1 \leq i \leq \rho$.
 - Serially reusable resources are typically allocated to jobs on a nonpreemptive basis and used in a **mutual exclusive** manner.

Enforcement of Mutual Exclusion and Critical Sections (1/3)

- When a job wants to use η_i units of resource R_i , it executes a **lock** to request them, denoted by $L(R_i, \eta_i)$.
- When the job no longer needs the resource, it releases the resource by executing an **unlock**, denoted by $U(R_i, \eta_i)$.
- When a resource R_i has only one unit, we use the simpler notations $L(R_i)$ and $U(R_i)$ for lock and unlock.

Enforcement of Mutual Exclusion and Critical Sections (2/3)

- We call a segment of a job that begins at a lock and ends at a matching unlock a **critical section**.
 - Since resources are released in the LIFO order, overlapping critical sections are **properly nested**.
 - A critical section that is not included in other critical sections is an **outermost critical section**.




Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

5

Enforcement of Mutual Exclusion and Critical Sections (3/3)

- We denote each critical section by $[R, \eta; e]$.
 - R : the name of the resource
 - η : the number of units of the resource
 - e : the (maximum) execution time of the critical section
 - When $\eta = 1$, we use the simpler notation $[R; e]$.
 - We denote nested critical sections by nested square brackets.



6

Outline

- Assumptions on Resources and Their Usage
- Effects of Resource Contention and Resource Access Control**
- Nonpreemptive Critical Sections Protocol
- Basic Priority-Inheritance Protocol
- Basic Priority-Ceiling Protocol
- Stack-Based, Priority-Ceiling (Ceiling-Priority) Protocol
- Preemption-Ceiling Protocols
- Controlling Accesses to Multiple-Unit Resources

Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

7

Resource Access-Control Protocol


- Effects of resource contention
 - Priority inversion
 - Timing anomalies
 - Deadlock
- A **resource access-control protocol** is a set of rules that govern
 - When and under what conditions each request for resource is granted, and
 - How jobs requiring resources are scheduled

Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

8

Priority Inversion

- Priority inversion can occur when the execution of some jobs or portions of jobs is nonpreemptable.
- Resource contentions among jobs can also cause priority inversion.




Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

9

Timing Anomalies

- When priority inversion occurs, timing anomalies invariably follow.
- Timing anomaly:** When we shorten the execution time of some critical section, rather than complete sooner, some job might complete later.




Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

10

Unbounded Priority Inversion

- Without good resource access control, the duration of a priority inversion can be unbounded.



Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

11

Deadlock

- Nonpreemptivity of resource allocation can also lead to deadlocks.
 - Two jobs both require resources X and Y . The jobs are in deadlock when one of them holds X and requests for Y , while the other holds Y and requests for X .

Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

12

Criterion of Resource Access-Control Protocols

- No resource access-control protocol can eliminate the priority inversion and anomalous behavior caused by resource contention.
- A criterion we use to measure the performance of a resource access-control protocol is the **blocking time of each job**.
- ➔ A good resource access-control protocol should control priority inversions and prevent deadlock and keep the blocking time of every job bounded from the above.

Terms and Assumptions


- **Assumption:** No job ever suspends itself and every job is preemptable on the processor.
- A higher-priority job J_h is said to be **directly blocked** by a lower-priority job J_l when J_l holds some resource which J_h requests and is not allocated.
- We sometimes denote the periodic task by the tuple $(\phi_i, p_i, e_i, D_i, [R, x; y])$.
 - In general, when jobs in the task require more than one resource and hence have more than one critical section, we put all the critical sections in the tuple.

Wait-for Graph (1/2)

- In the **wait-for graph** of a system
 - Every **job** that requires some resource: A vertex labeled by the name of the job.
 - Every **resource**: A vertex labeled by the name and the number of units of the resource.
 - **Ownership** edge: An edge with label x from a resource vertex to a job vertex if x units of the resource are allocated to the job at the time.
 - **Wait-for** edge: An edge with label y from a job vertex to a resource vertex if the job requested y units of the resource earlier and the request was denied.

Wait-for Graph (2/2)


- A **path** in a wait-for graph from a higher-priority job to a lower-priority job represents the fact that the former is **directly blocked** by the later.
- A **cyclic path** in a wait-for graph indicates a **deadlock**.



Resource Requirements (1/2)

- We will specify **resource requirements** of a system by a **bipartite graph**
 - There is a vertex for every job and every resource named by the name of the job (or task) or resource it represents.
 - The integer next to each resource vertex R_i gives the number v_i of units of the resource.
 - A job J (or task T) **requires** a resource R_i is represented by an edge from the job (or task) vertex to the resource vertex.
 - We may label each edge by 2-tuple: (the number of units used in the critical section, the duration of the critical section)

Resource Requirements (2/2)




Outline

- Assumptions on Resources and Their Usage
- Effects of Resource Contention and Resource Access Control
- Nonpreemptive Critical Sections Protocol**
- Basic Priority-Inheritance Protocol
- Basic Priority-Ceiling Protocol
- Stack-Based, Priority-Ceiling (Ceiling-Priority) Protocol
- Preemption-Ceiling Protocols
- Controlling Accesses to Multiple-Unit Resources

Nonpreemptive Critical Sections

- Nonpreemptive Critical Section (NPCS) protocol**
 - The simplest way to control access of resources
 - When a job requests a resource, it is always allocated the resource.
 - When a job holds any resource, it executes at a priority higher than the priority of all jobs.
 - No job is ever preempted when it holds any resource, deadlock can never occur.
 - J_h can be blocked only once, and its blocking time due to resource conflict is at most equal to the maximum execution time of the critical sections of all lower-priority jobs.

Example



Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

21

Pros and Cons of NPCS

- The advantage is its **simplicity**, especially when the numbers of resource units are arbitrary.
 - Does not need any prior knowledge about resource requirements of jobs
 - Simple to implement and can be used in both fixed-priority and dynamic-priority systems
 - Good protocol when all the critical sections are short and when most of the jobs conflict with each other
- The shortcoming is that every job **can be blocked by every lower-priority job** that requires some resource even when there is no resource conflict between them.

Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

22

Outline

- Assumptions on Resources and Their Usage
- Effects of Resource Contention and Resource Access Control
- Nonpreemptive Critical Sections Protocol
- **Basic Priority-Inheritance Protocol**
- Basic Priority-Ceiling Protocol
- Stack-Based, Priority-Ceiling (Ceiling-Priority) Protocol
- Preemption-Ceiling Protocols
- Controlling Accesses to Multiple-Unit Resources

Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

23

Basic Priority-Inheritance Protocol (PIP)

- A simple protocol proposed by Sha, *et al.**
- Work with any preemptive, priority-driven scheduling algorithm
- Does not require prior knowledge on resource requirements of jobs
- Does **NOT** prevent deadlock
- Assume every resource has only 1 unit

*Sha, L., R. Rajkumar, and J. P. Lehoczky, "Priority Inheritance Protocols: An Approach to Real-Time Synchronization," *IEEE Transactions on Computers*, Vol. 39, 1990.

Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

24

Definition of Basic PIP

- We call the priority that is assigned to a job according to the scheduling algorithm its **assigned priority**.
- At any time t , each ready job J_l is scheduled and executes at its **current priority** $\pi_l(t)$, which may differ from its assigned priority and may vary with time.
- The current priority $\pi_l(t)$ of a job J_l may be raised to the higher priority $\pi_h(t)$ of another job J_h . We say that the lower-priority job J_l **inherits** the priority of the higher priority job J_h and that J_l executes at its **inherited priority** $\pi_h(t)$.

Rules of the Basic PIP (1/2)


- Scheduling Rule:** Ready jobs are scheduled on the processor preemptively in a priority-driven manner according to their current priorities. At its release time t , the current priority $\pi(t)$ of every job J is equal to its assigned priority. The job remains at this priority except under the condition stated in rule 3.
- Allocation Rule:** When a job J requires a resource R at time t ,
 - If R is free, R is allocated to J until J releases the resource
 - If R is not free, the request is denied and J is blocked

Rules of the Basic PIP (2/2)

- Priority-Inheritance Rule:** When the requesting job J becomes blocked, the job J_l which blocks J inherits the current priority $\pi(t)$ of J . The job J_l executes at its inherited priority $\pi(t)$ until it releases R ; at that time, the priority of J_l returns to its priority $\pi_l(t')$ at the time t' when it acquires the resource R .
- The protocol ensures that the duration of priority inversion is never longer than the duration of an outermost critical section **each time** a job is blocked.

Example

Job	r_j	e_j	π_j	Critical Sections
J_1	7	3	1	[Shaded; 1]
J_2	5	3	2	[Black; 1]
J_3	4	2	3	
J_4	2	6	4	[Shaded; 4] [Black; 1.5]
J_5	0	6	5	[Black; 4]




Properties of the PIP (1/2)

- There are two types of blocking:
 - Direct blocking**: J_2 is directly blocked by J_5 in $(6, 11]$ and by J_4 in $(11, 12.5]$
 - Priority-inheritance blocking** (or simply **inheritance blocking**): J_3 is blocked by J_5 in $(6, 7]$ because the latter inherits a higher priority in this interval.
- ➡ Priority-inheritance blocking suffered by jobs that are **not involved in resource contention** is the cost for controlling the durations of priority inversion suffered by jobs that are involved in resource contention.

Properties of the PIP (2/2)

- Jobs can **transitively block** each other:
 - At time 9, J_5 blocks J_4 , and J_4 blocks J_1 .
 - In the time interval $(9, 11)$, J_5 inherits J_4 's priority, which J_4 inherited from J_1 .
- Disadvantages of the PIP:
 - Does not prevent deadlock**: suppose J_5 request the resource *Shaded* at time 6.5.
 - Does not reduce the blocking times suffered by jobs as small as possible**.

Worst-Case Blocking Scenario for PIP



Outline

- Assumptions on Resources and Their Usage
- Effects of Resource Contention and Resource Access Control
- Nonpreemptive Critical Sections Protocol
- Basic Priority-Inheritance Protocol
- Basic Priority-Ceiling Protocol**
- Stack-Based, Priority-Ceiling (Ceiling-Priority) Protocol
- Preemption-Ceiling Protocols
- Controlling Accesses to Multiple-Unit Resources

Basic Priority-Ceiling Protocol (PCP)

- The **priority-ceiling protocol*** extends the priority-inheritance protocol to **prevent deadlocks** and to **further reduce the blocking time**.
- Two key assumptions:
 - The assigned priorities of all jobs are fixed.
 - The resources required by all jobs are known a priori before the execution of any job begins.

*Sha, L., R. Rajkumar, and J. P. Lehoczky, "Priority Inheritance Protocols: An Approach to Real-Time Synchronization," *IEEE Transactions on Computers*, Vol. 39, 1990.

Definitions

- The **priority ceiling** of any resource R_i is the highest priority of all the jobs that require R_i and is denoted by $\Pi(R_i)$.
- At any time t , the **current priority ceiling** (or the **ceiling**) $\hat{\Pi}(t)$ of the system is equal to the highest priority ceiling of the resources that are in use at the time, if some resources are in use.
- If all the resource are free at the time, the current ceiling $\hat{\Pi}(t)$ is equal to Ω , a nonexisting priority level that is lower than the lowest priority of all jobs.

Rules of Basic Priority-Ceiling Protocol (1/3)

1. Scheduling Rule:

- a) At its release time t , the current priority $\pi(t)$ of every job J is equal to its assigned priority. The job remains at this priority except under the condition stated in rule 3.
- b) Every ready job J is scheduled preemptively and in a priority-driven manner at its current priority $\pi(t)$.

Rules of Basic Priority-Ceiling Protocol (2/3)

2. Allocation Rule: Whenever a job J requests a resource R at time t , one of the following two conditions occurs:

- a) The resource R is held by another job. J 's request fails and J becomes blocked.
- b) The resource R is free.
 - 1) If J 's priority $\pi(t)$ is higher than the current priority ceiling $\hat{\Pi}(t)$, R is allocated to J .
 - 2) If J 's priority $\pi(t)$ is not higher than the ceiling $\hat{\Pi}(t)$ of the system, R is allocated to J only if J is the job holding the resource(s) whose priority ceiling is equal to $\hat{\Pi}(t)$; otherwise, J 's request is denied, and J becomes blocked.

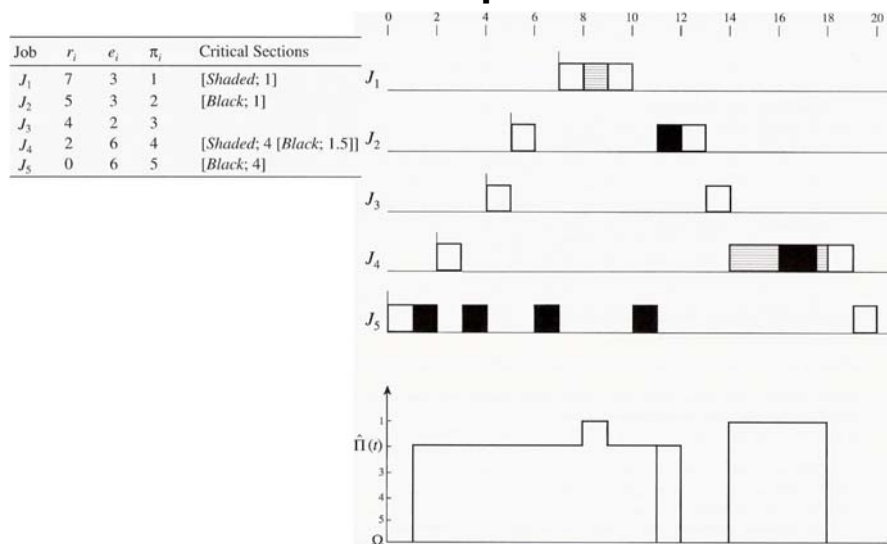
Rules of Basic Priority-Ceiling Protocol (3/3)

- 3. Priority-Inheritance Rule:** When J becomes blocked, the job J_l which blocks J inherits the current priority $\pi(t)$ of J . J_l executes at its inherited priority until the time when it releases every resource whose priority ceiling is equal to or higher than $\pi(t)$; at that time, the priority of J_l returns to its priority $\pi_l(t')$ at the time t' when it was granted the resource(s).

Assumptions in Rules

- We note that [2\) in part b\) of rule 2](#) assumes that only one job holds all the resources with priority ceiling equal to $\hat{\Pi}(t)$.
- [Rule 3](#) assumes that only one job is responsible for J 's request being denied, because it holds either the requested resource or a resource with priority ceiling $\hat{\Pi}(t)$.

Example

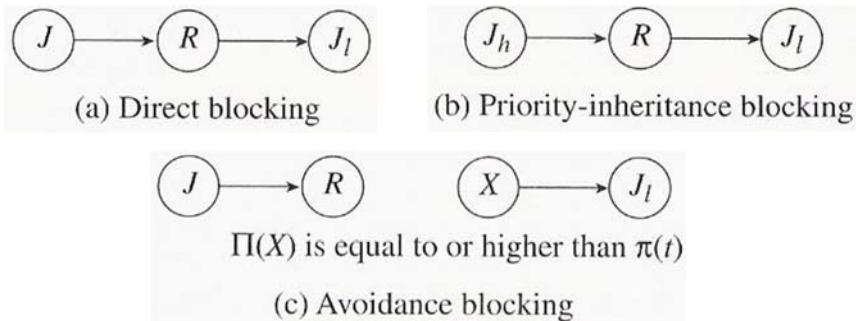


PIP vs. PCP

- A fundamental difference is that the PIP is greedy while the PCP is not.
- The priority-inheritance rules of these two protocols are essentially the same.
 - The difference arises because of the nongreedy nature of the PCP.
 - It is possible for J to be blocked by a lower-priority job which does not hold the requested resource according to the PCP, while this is impossible according to the PIP.

Ways for a Job to Block Another Job in PCP

- **Priority-ceiling blocking (avoidance blocking)**: the requesting job J is blocked by a lower-priority job J_l when J requests a resource R that is free at the time. The reason is that J_l holds another resource X whose priority ceiling is equal to or higher than J 's priority $\pi(t)$.




Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

41

Deadlock Avoidance by PCP (1/3)

- The set of priority ceilings of resources impose a linear order on all the resources.



Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

42

Deadlock Avoidance by PCP (2/3)

- At any time t , the priority $\pi(t)$ of a job J being higher than the current ceiling $\hat{\Pi}(t)$ of the system means that
 - Job J will not require any of the resources in use at t
 - Jobs with priorities equal to or higher than J will not require any of these resource
- ➡ The priority ceiling $\hat{\Pi}(t)$ of the system tells us the subset of all jobs to which we can safely grant free resources at time t ; this subset contains all the jobs that have higher priorities than $\hat{\Pi}(t)$.

Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

43

Deadlock Avoidance by PCP (3/3)

- **2) in part b) of rule 2** states an exception to the rule that J 's request for any resource is denied if its priority is not higher than the ceiling of the system.
 - The exception applies when J is the job holding the resource(s) whose priority ceiling(s) is equal to $\hat{\Pi}(t)$;
- **Theorem 1.** When resource accesses of a system of preemptive, priority-driven jobs on one processor are controlled by the priority-ceiling protocol, deadlock can never occur.

Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

44

Duration of Blocking


- Under the PCP, a job may be directly blocked, avoidance blocked, and inheritance blocked by lower-priority jobs.
- ➡ Whether the PCP would cause a job to be blocked for a longer duration than the PIP (as a cost of its ability to prevent deadlock)?
- Theorem 2.** When resource accesses of preemptive, priority-driven jobs on one processor are controlled by the PCP, a job can be blocked for at most the duration of one critical section.

Informal Proof of Theorem 2

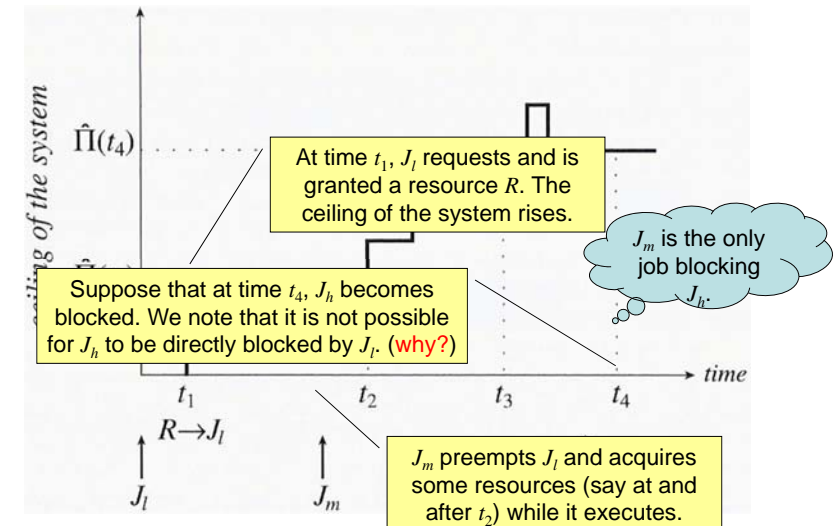
There are two parts to this argument:

- When a job becomes blocked, it is blocked by only one job
 - A job which blocks another job cannot be blocked in turn by some other job.
- ➡ There can be no transitive blocking under the priority-ceiling protocol.

Why 1. Is True?

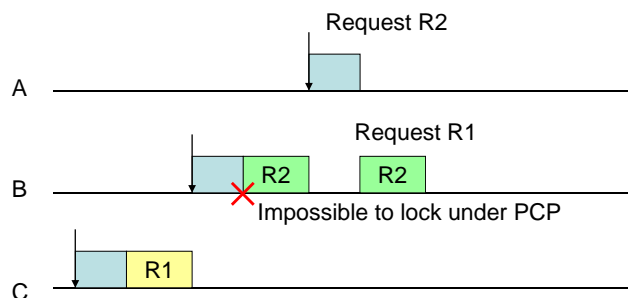


Why 1. Is True?



Why 2. Is True?


- Suppose three jobs J_l , J_m , and J_h are blocked transitively.



Computation of Blocking Time


- Theorem 2** makes it easy for us to compute an upper bound to the amount of time a job may be blocked due to resource conflicts.
 - We call this upper bound the **blocking time (due to resource conflicts)** of the job.

Example:




Although J_2 and J_3 do not require the resource *Black*, they can be priority-inheritance blocked by J_4 since J_4 can inherit priority π_1 . Hence, the blocking time $b_2(rc)$ and $b_3(rc)$ are also one.

A Slightly More Complicated Example

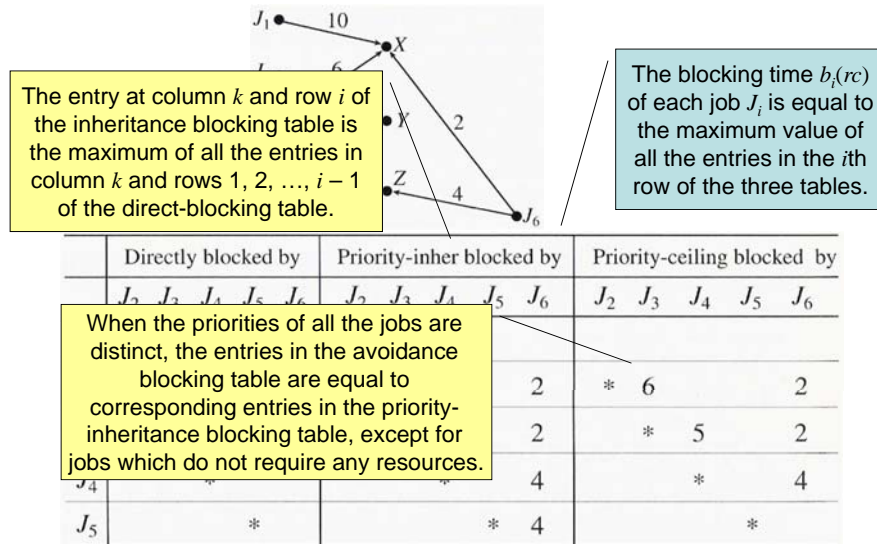


A Systematic Way to Compute Blocking



	Directly blocked by					Priority-inher blocked by					Priority-ceiling blocked by				
	J_2	J_3	J_4	J_5	J_6	J_2	J_3	J_4	J_5	J_6	J_2	J_3	J_4	J_5	J_6
J_1															
J_2		*				*	6			2	*	6			2
J_3			*				*	5		2		*	5		2
J_4				*				*		4			*		4
J_5					*				*	4				*	

A Systematic Way to Compute Blocking



Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

53

Outline

- Assumptions on Resources and Their Usage
- Effects of Resource Contention and Resource Access Control
- Nonpreemptive Critical Sections Protocol
- Basic Priority-Inheritance Protocol
- Basic Priority-Ceiling Protocol
- **Stack-Based, Priority-Ceiling (Ceiling-Priority) Protocol**
- Preemption-Ceiling Protocols
- Controlling Accesses to Multiple-Unit Resources

Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

54

Stack-Based, Priority-Ceiling (Ceiling-Priority) Protocol

- We will give two different definitions of a protocol
 - Simpler than the PCP
 - Has the same worst-case performance as the PCP
- The different definitions arise from two different motivations:
 - To provide stack-sharing capability (Stack-Based, Priority-Ceiling Protocol)
 - To simplify the priority-ceiling protocol (Ceiling-Priority Protocol)

Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

55

Motivation of Stack-Sharing Priority-Ceiling Protocol

- Sometimes, especially in systems where the number of jobs is large, it may be necessary for the jobs to share a common run-time stack, in order to **reduce overall memory demand**.

Copyright: All right reserved, Jen-Wei Hsieh, CSIE, NTUST

56

How Stack Sharing Operate

- Space in the (shared) stack is allocated to jobs contiguously in the last-in-first-out manner.
- When a job J executes, its stack space is on the top of the stack.
- The space is freed when the job completes.
- When J is preempted, the preempting job has the stack space above J 's.
- J can resume execution only after all the jobs holding stack space above its space complete, free their stack space, and leave J 's stack space on the top of the stack again.


Rules Defining Basic Stack-Based, Priority-Ceiling Protocol

1. **Update of the Current Ceiling:** Whenever all the resources are free, the ceiling of the system is Ω . The ceiling $\hat{\Pi}(t)$ is updated each time a resource is allocated or freed.
2. **Scheduling Rule:** After a job is released, it is blocked from starting execution until its assigned priority is higher than the current ceiling $\hat{\Pi}(t)$ of the system.
3. **Allocation Rule:** Whenever a job requests a resource, it is allocated the resource.

Stack-Based, Priority-Ceiling Protocol

- According to the scheduling rule, when a job begins to execute, all the resource it will ever need during its execution are free. (why?)
- ➡ No job is ever blocked once its execution begins.
- When a job J is preempted, all the resources the preempting job will require are free, ensuring that the preempting job can always complete so J can resume.
- ➡ Deadlock can never occur.

Example (vs. PCP)



Rules Defining the Ceiling-Priority Protocol

1. Scheduling Rule:

- a) Every job executes at its assigned priority when it does not hold any resource. Jobs of the same priority are scheduled on the FIFO basis.
- b) The priority of each job holding any resource is equal to the highest of the priority ceiling of all resources held by the job.

2. Allocation Rule: Whenever a job requests a resource, it is allocated the resource.

- Note that when jobs never self-suspend, the stack-based priority-ceiling and ceiling-priority protocols are the same.

Blocking Time and Context-Switch Overhead

- **Theorem 3.** The longest blocking time suffered by every job is the same for the stack-based and basic [priority-ceiling protocols](#).
- The context-switch overhead is smaller under the stack-based version.
 - Because no job is ever blocked once its execution starts, no job ever suffers more than two context switches.

Outline

- Assumptions on Resources and Their Usage
- Effects of Resource Contention and Resource Access Control
- Nonpreemptive Critical Sections Protocol
- Basic Priority-Inheritance Protocol
- Basic Priority-Ceiling Protocol
- Stack-Based, Priority-Ceiling (Ceiling-Priority) Protocol
- **Preemption-Ceiling Protocols**
- Controlling Accesses to Multiple-Unit Resources

Motivation

- In a dynamic-priority system, the priorities of the periodic tasks change with time while the resources required by each task remain constant.
- ➡ The priority ceilings of the resources may change with time.
- ➡ The [preemption-ceiling protocol](#) is based on the clever observation that the potentials of resource contentions in a dynamic-priority system do not change with time, just as in fixed-priority systems, and hence can be analyzed statically.

Observations (1/2)

- The fact that a job J_h has a higher priority than another job J_l and they both require some resource does not imply that J_l can directly block J_h . This blocking can occur only when it is possible for J_h to preempt J_l .

➡ When determining whether a free resource can be granted to a job, it is not necessary to be concerned with the resource requirements of all higher-priority jobs; only those that can preempt the job.

Observations (2/2)

- For some dynamic priority assignments, it is possible to determine a priori the possibility that jobs in each periodic task will preempt the jobs in other periodic tasks.

➡ In a deadline-driven system, no job in a periodic task with a smaller relative deadline is ever preempted by jobs in periodic tasks with identical or larger relative deadlines.


Preemption Levels (1/4)

- The possibility that a job J_i will preempt another job is captured by the parameter **preemption level** ψ_i of J_i .
 - The preemption levels of jobs are functions of their **priorities** and **release times**.
 - According to a **valid preemption-level assignment**, for every pair of jobs J_i and J_k , the preemption level ψ_i of J_i being equal to or higher than the preemption level ψ_k of J_k implies that it is never possible for J_k to preempt J_i .
- Validity Condition:** If π_i is higher than π_k and $r_i > r_k$, then ψ_i is higher than ψ_k .

Preemption Levels (2/4)


- The preemption levels of jobs that are not given by the above rule are valid as long as the linear order over all jobs defined by preemption-level assignment does not violate the validity condition.

Example:



Preemption Levels (3/4)

- Example:**




- Method 1: release time + priority (1, 1, 2, 3, 3)
- Method 2: release time (2, 1, 3, 5, 4)
- Method 3: relative deadline (EDF, periodic task)

Preemption Levels (4/4)

- A system of periodic tasks is a **fixed preemption-level system** if there is a valid assignment of preemption levels to all jobs such that the **jobs in every task have the same preemption level**.
 - All fixed-priority systems are also fixed preemption-level systems.
 - Periodic tasks scheduled on the LIFO basis have varying preemption levels. ([illustrate](#))

Example: Varying Preemption Levels

- Periodic tasks scheduled on the LIFO basis



- How about preemption levels?
 - (0, 3):
 - (3, 4.5):
 - (4.5, 5):
 - (5, 6):

Definitions of Protocols

- A **preemption-ceiling protocol** makes decisions on whether to grant a free resource to any job based on the preemption level of the job in a manner similar to the priority-ceiling protocol.
 - Assumptions:
 - The resource requirements of all jobs are known a priori.
 - There is only 1 unit of each resource
 - The **preemption ceiling** $\psi(R)$ of a resource R is the highest preemption level of all the jobs that require the resource.
 - The **(preemption) ceiling of the system** $\hat{\Psi}(t)$ at any time t is the highest preemption ceiling of all the resources that are in use at t .

Rules of Preemption-Ceiling Protocol (1/2)

- Basic priority-ceiling and preemption-ceiling protocols differ mainly in their [allocation rules](#).
- **Rules of Basic Preemption-Ceiling Protocol:**
 - 1 and 3. The [Scheduling Rule](#) and [Priority Inheritance Rule](#) are the same as the corresponding rules of the PCP.
 2. **Allocation Rule:** Whenever a job J requests resource R at time t , one of the following two conditions occurs:

Rules of Preemption-Ceiling Protocol (2/2)

- a) The resource R is held by another job. J 's request fails, and J becomes blocked.
- b) The resource R is free.
 - 1) If J 's preemption level $\psi(t)$ is higher than the current preemption ceiling $\hat{\Psi}(t)$ of the system, R is allocated to J .
 - 2) If J 's preemption level $\psi(t)$ is not higher than the ceiling $\hat{\Psi}(t)$ of the system, R is allocated to J only if J is the job holding the resource(s) whose preemption ceiling is equal to $\hat{\Psi}(t)$; otherwise, J 's request is denied, and J becomes blocked.

[back](#)

Stack-Based Protocol (SBP)

- The [stack-based preemption-ceiling protocol](#) is called [Stack-Based Protocol \(SBP\)](#) by Backer*.
- The difference is that priority levels/ceilings are replaced by preemption levels/ceilings.
 - Rules of [Stack-Based Priority-Ceiling Protocol](#)
- In addition, the stack-based preemption-ceiling protocol has an [inheritance rule](#).

Rules of Stack-Based Protocol (1/2)

0. **Update of the Current Ceiling:** Whenever all the resources are free, the preemption ceiling of the system is Ω . The preemption ceiling $\hat{\Psi}(t)$ is updated each time a resource is allocated or freed.
1. **Scheduling Rule:** After a job is released, it is blocked from starting execution until its preemption level is higher than the current ceiling $\hat{\Psi}(t)$ of the system and the preemption level of the executing job. At any time t , jobs that are not blocked are scheduled on the


*Baker, T.P., "A Stack-Based Resource Allocation Policy for Real-Time Processes," *IEEE Real-Time System Symposium*, 1991.

Rules of Stack-Based Protocol (2/2)

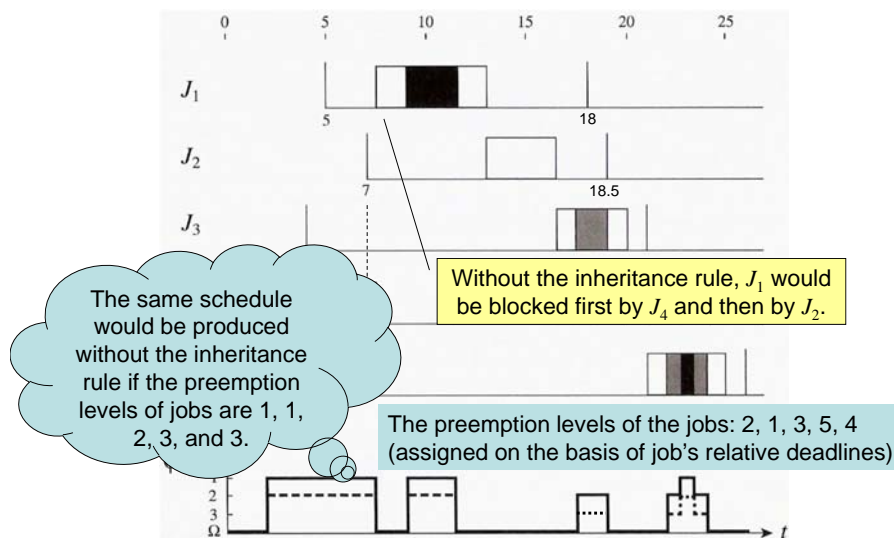
processor in a priority-driven, preemptive manner according to their assigned priorities.

2. **Allocation Rule:** Whenever a job J requests for a resource R , it is allocated resource.
3. **Priority-Inheritance Rule:** When some job is blocked from starting, the blocking job inherits the highest priority of all the blocked jobs.

Why We Need Priority-Inheritance Rule?



Why We Need Priority-Inheritance Rule?




Outline

- Assumptions on Resources and Their Usage
- Effects of Resource Contention and Resource Access Control
- Nonpreemptive Critical Sections Protocol
- Basic Priority-Inheritance Protocol
- Basic Priority-Ceiling Protocol
- Stack-Based, Priority-Ceiling (Ceiling-Priority) Protocol
- Preemption-Ceiling Protocols
- **Controlling Accesses to Multiple-Unit Resources**

Notations (1/2)

- We let $\Pi(R_i, k)$, for $k \leq v_i$, denote the priority ceiling of a resource R_i when k out of the v_i (≥ 1) units of R_i are free.
 - If one or more jobs in the system require more than k units of R_i , $\Pi(R_i, k)$ is the highest priority of all these jobs.
 - If no job requires more than k units of R_i , $\Pi(R_i, k)$ is equal to Ω , the nonexisting lowest priority.
 - The priority ceiling of a resource R_j that has only 1 unit is $\Pi(R_j, 0)$.
- Let $k_i(t)$ denote the number of units of R_i that are free at time t .

Example



Resources	$x(2)$	$y(3)$	$z(1)$
Units Required by			
J_1	1	0	0
J_2	0	2	0
J_3	2	3	1
J_4	1	0	1
J_5	0	1	0
$\Pi(*, 0)$			
$\Pi(*, 1)$			
$\Pi(*, 2)$			
$\Pi(*, 3)$			

Notations (2/2)

- The preemption ceiling $\psi(R_i, k)$ of the resource R_i when k units of R_i are free is the highest preemption level of all the jobs that require more than k units of R_i .
- The preemption ceiling of the system at time t is equal to the highest preemption ceiling of all the resources at the time.

Multiple-Unit Ceiling-Priority Protocol

- In essence, the scheduling and allocation rules remains unchanged except for the new definition of priority ceiling of resources.
- Scheduling Rule 1b)** needs to be rephrased: Upon acquiring a resource R and leaving $k \geq 0$ free units of R , a job executes at the higher of its priority and the priority ceiling $\Pi(R, k)$ of R .

Multiple-Unit Priority-(Preemption)-Ceiling Protocol (1/2)

- Similarly, the allocation rule of the [priority-ceiling](#) (or [preemption-ceiling](#)) protocol for multiple units of resources is a straightforward modification.

Whenever a job J requests k units of resource R at time t , one of the following two condition occurs:

- Less than k units of R are free. J 's request fails and J becomes directly blocked.
- k or more units of R are free.
 - If J 's priority $\pi(t)$ [preemption level $\psi(t)$] is higher than the current priority ceiling $\hat{\Pi}(t)$ [preemption ceiling $\hat{\Psi}(t)$] of the system at the time, k units of R are allocated to J until it releases them.

Multiple-Unit Priority-(Preemption)-Ceiling Protocol (2/2)

- If J 's priority $\pi(t)$ [preemption level $\psi(t)$] is not higher than the system ceiling $\hat{\Pi}(t)$ [preemption ceiling $\hat{\Psi}(t)$], k units of R are allocated to J only if J holds the resource(s) whose priority ceiling (preemption ceiling) is equal to $\hat{\Pi}(t)$ [$\hat{\Psi}(t)$]; otherwise, J 's request is denied, and J becomes blocked.

You can see that this rule is essentially the same as the allocation rule of the basic version. The only change is in the wording to accommodate multiple-unit requests.

Priority-Inheritance Rule (1/2)

- In the case where there is only 1 unit of each resource, we have shown that when a job J is blocked, only one lower-priority job is responsible for this blocking and this lower-priority job inherits J 's priority.

➡ In a system containing multiple resource units, **which job shall inherits the priority?**

- Example:** 3 units of resource R , 4 jobs each requiring 1 unit of R . J_1 request a unit of R , all 3 units are held by the other 3 jobs.

Priority-Inheritance Rule (2/2)

- Priority-Inheritance Rule:** When the requesting job J becomes blocked at t , the job with the highest priority among all the jobs holding the resource R that has the highest priority ceiling among all the resources inherits J 's priority until it releases its unit of R .

➡ With this rule, each job is blocked at most once for the duration of one critical section.

- Example:**

	no. of units	units required					$\Pi(*, k), k =$					
		J_1	J_2	J_3	J_4	J_5	0	1	2	3	4	5
Black	5	2	4	0	1	1	1	1	2	2	Ω	Ω
Shaded	1	1	1	0	0	1	1	Ω	Ω	Ω	Ω	Ω

